

- Naam: Mikrocentrum_gasten
- Wachtwoord: mikrowifi1115!

<https://tinyurl.com/AIVEslides>

Simcha van Helvoort
Dag 1

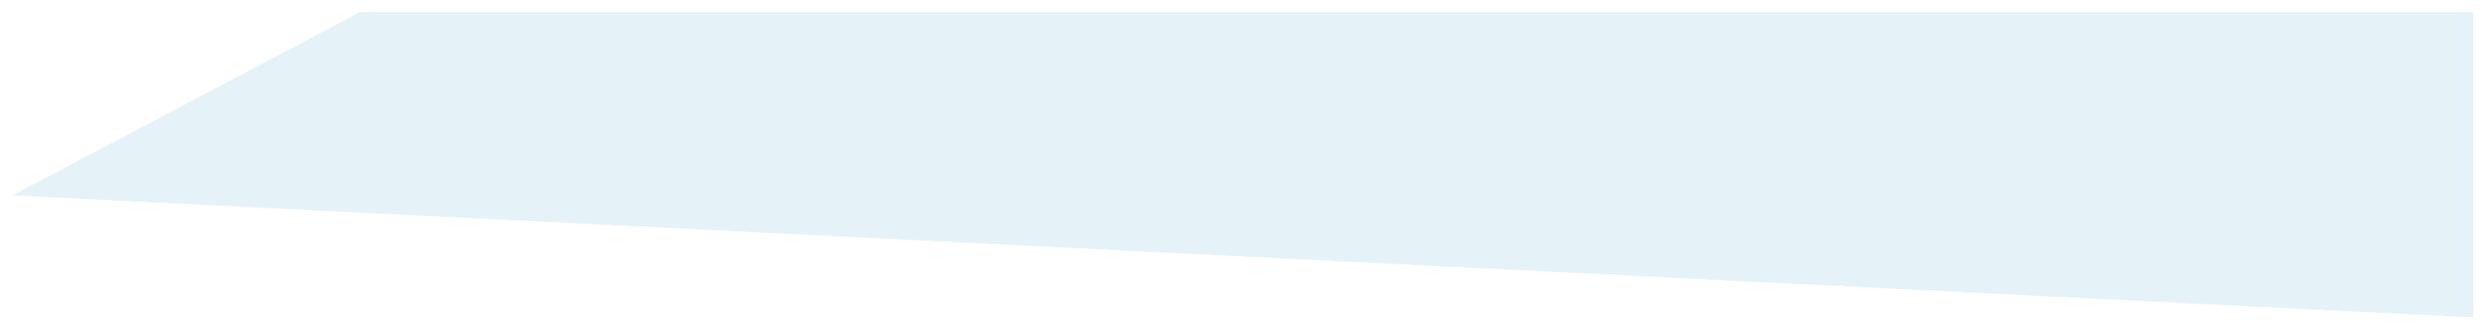


**Artificial Intelligence voor
Engineers**

WIE BEN IK?

WIE ZIJN JULLIE?

- Wie ben je en wat is je achtergrond?
- Waar werk je en wat is je functie/dagelijkse bezigheden?
- Wat hoop je te leren?
- Wat is je ervaring met data en AI?



WAT WILLEN WE BEREIKEN?

- Kun jij zelf een data gedreven AI model maken
- Weten bij welke problemen je AI zou kunnen gebruiken
- Kan je van ruwe data informatie halen
- Kan je een AI model evalueren

Leerdoelen

- Python en libraries
- Crash course basis statistiek
- Wat is het verschil tussen AI en ML?
- Hoe werkt ML (globaal)
- Wat is het verschil tussen supervised en unsupervised learning?
- 3 verschillende ML technieken
- Hoe evalueer je een ML model?
- Hoe trek je conclusies uit een ML model?

Terminologie

- Begrippen en concepten zullen *schuin* gedrukt zijn
- Vaak direct uitgelegd
- Anders komt de uitleg later in de cursus
of kan verder over gelezen worden in leesmateriaal

Waarmee gaan we werken

Programmeertaal: Python

Hoe: Jupyter Notebook op je [eigen machine](#) via Anaconda

Bestanden: <https://github.com/SimchaGD/AIVE>

De core ligt op deze packages:

- SciKit Learn - <https://scikit-learn.org/>
- Lifelines - <https://lifelines.readthedocs.io/>
- Numpy - <https://numpy.org/>
- Pandas - <https://pandas.pydata.org/>
- Matplotlib - <https://matplotlib.org/>
- Scipy - <https://www.scipy.org/>

- Hoe werkt het?
(demo)
- Voordelen

The screenshot shows a Jupyter Notebook interface with the following content:

```
jupyter Untitled Last Checkpoint: 05/31/2021 (unsaved changes)
```

File Edit View Insert Cell Kernel Help

Trusted Python 3 Logout

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
```

```
In [2]: filename = '../3dprinter/data.csv'
df = pd.read_csv(filename)
```

```
In [3]: df.head()
```

```
Out[3]:
```

	layer_height	wall_thickness	infill_density	infill_pattern	nozzle_temperature	bed_temperature	print_speed	material	fan_speed	roughness	tension_strenght
0	0.02	8	90	grid	220	60	40	abs	0	25	18
1	0.02	7	90	honeycomb	225	65	40	abs	25	32	16
2	0.02	1	80	grid	230	70	40	abs	50	40	8
3	0.02	4	70	honeycomb	240	75	40	abs	75	68	10
4	0.02	6	90	grid	250	80	40	abs	100	92	5

```
In [4]: from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score
```

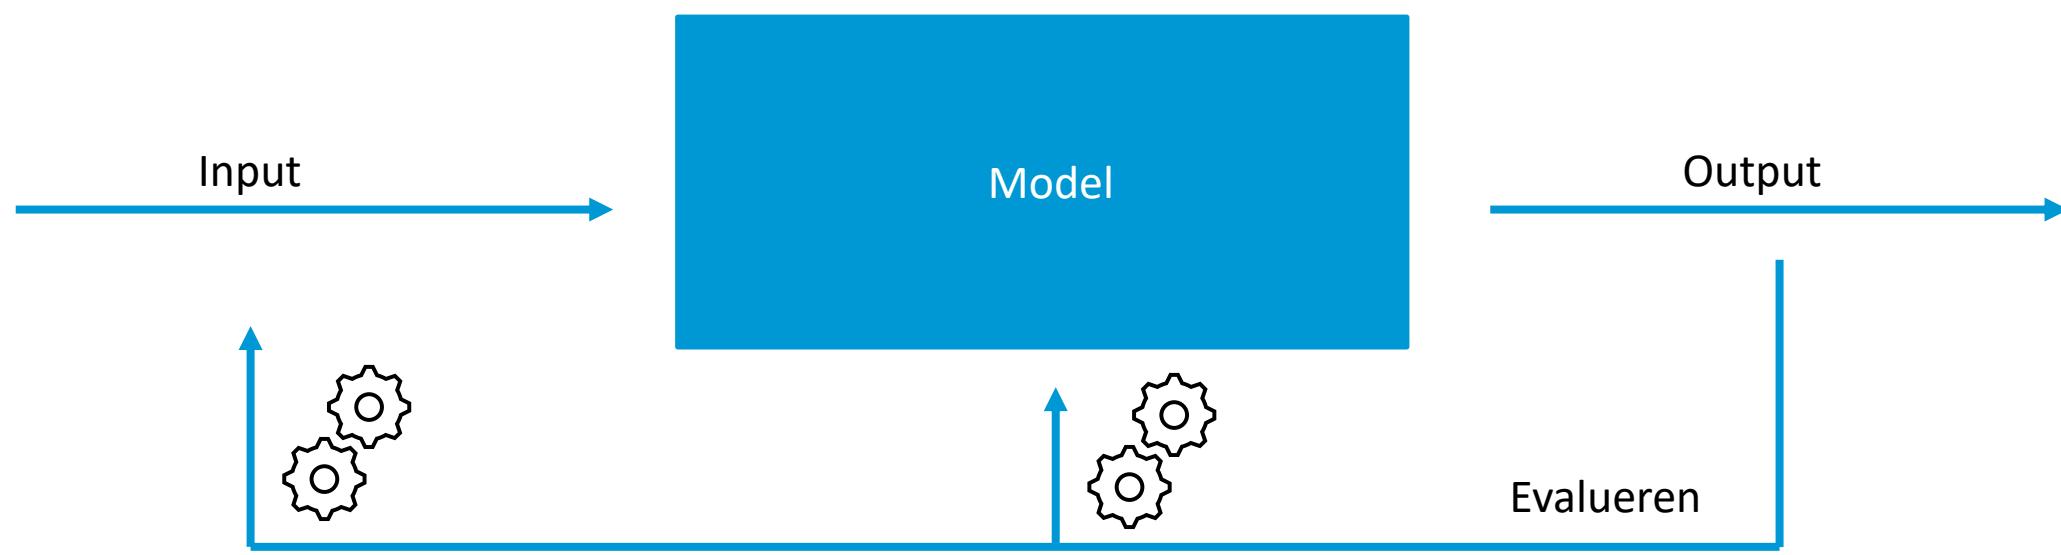
```
In [5]: # Verander alle categorien naar cijfers. We zetten de infill_pattern "grid" op 1 en "honeycomb" op 0.
df.infill_pattern = df.infill_pattern.apply(lambda x: 1 if x == "grid" else 0)

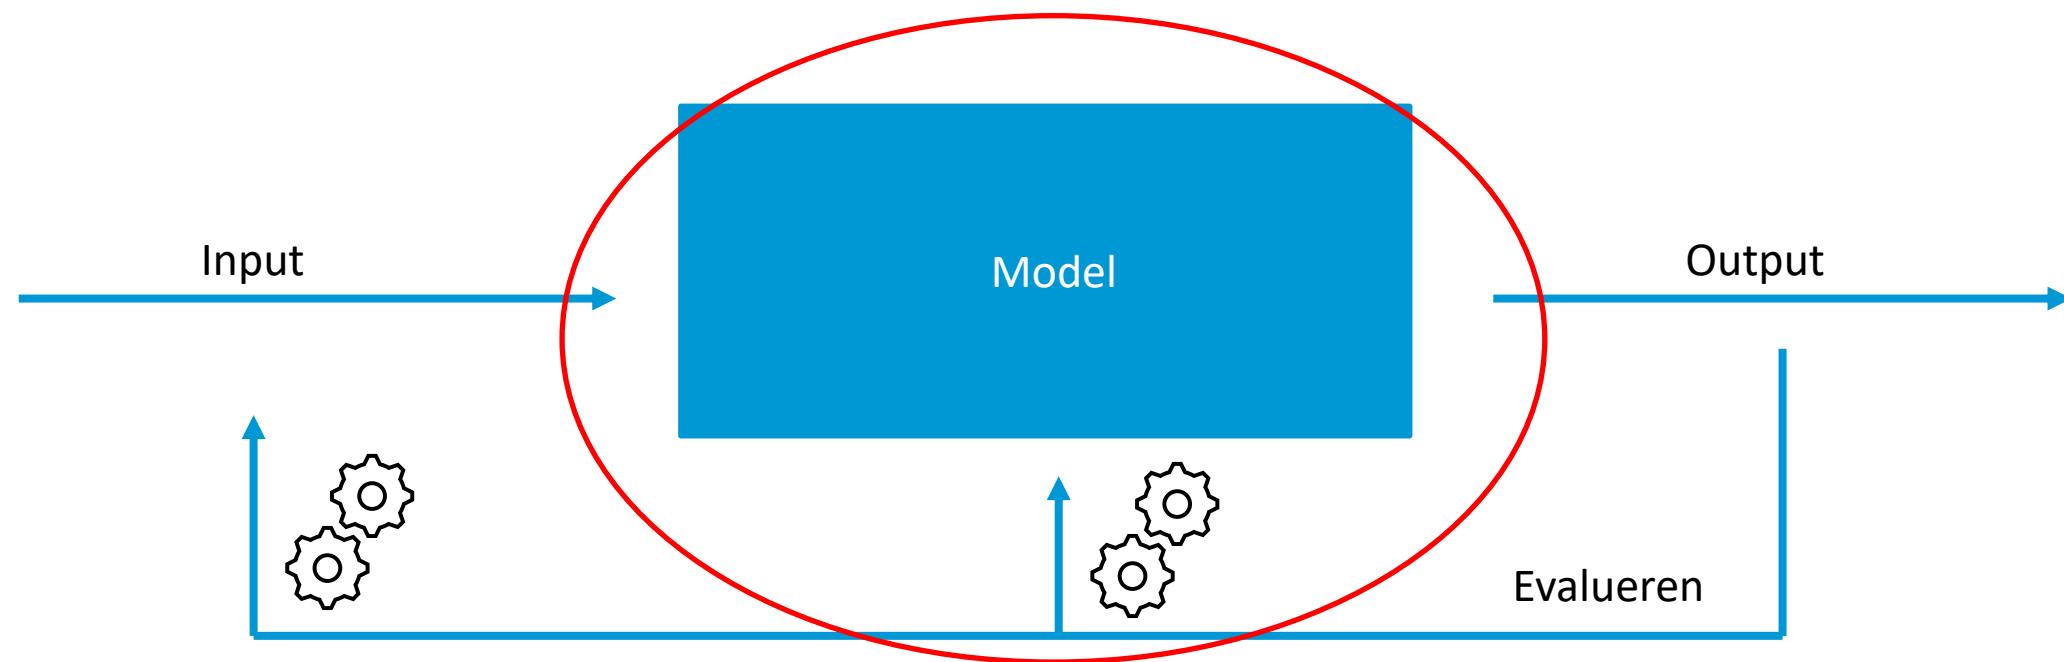
# We zetten de material "abs" op 1 en "pla" op 0.
df.material = df.material.apply(lambda x: 1 if x == "abs" else 0)

df.head() # Hier zien we dat de infill pattern 1 en 0 zijn geworden. En de kolom material staat er niet meer in
```

Pipeline

Waar gaan we naar toe werken





Maar eerst:

BASIS STATISTIEK

2 soorten data

Categorisch

- Benaderingen
- Groeperingen
- Beschrijvingen

Numeriek

- Exact
- Vaste ordening

Data type

Categorisch

Nominaal

Categorieën zonder volgorde

- Man/Vrouw
- Hond/Kat
- Groep A/
Group B/
Groep C

Ordinaal

Categorieën met volgorde

- Eens/Middel/Oneens
- Basisschool/
Middelbare school/
Universiteit

Numeriek

Interval

Numeriek zonder "echt nulpunt". Gelijke afstand tussen waarden

- Temperatuur °C:
-10/-5/0/5/10/20
- IQ
- Tijd op een klok

Ratio

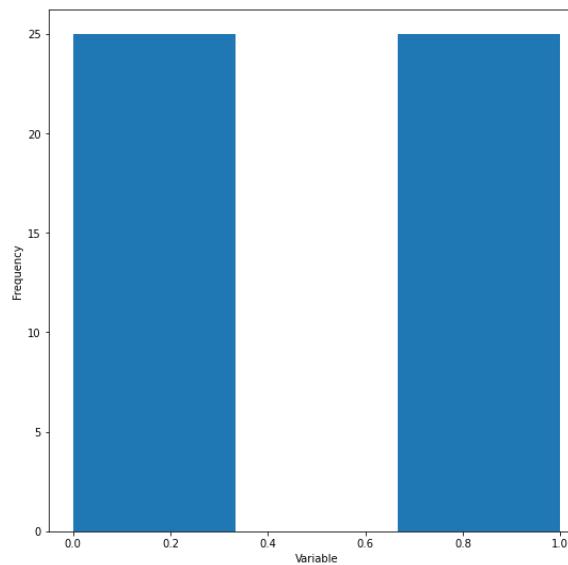
Numeriek met nulpunt.
Gelijke afstand

- Meter
- Gram
- Aantal presidenten van een land

Data types plotten

Categorisch

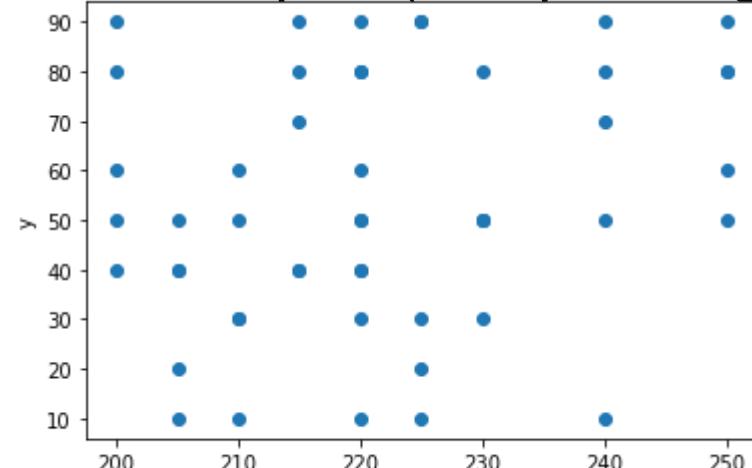
- Histogram (frequentie)



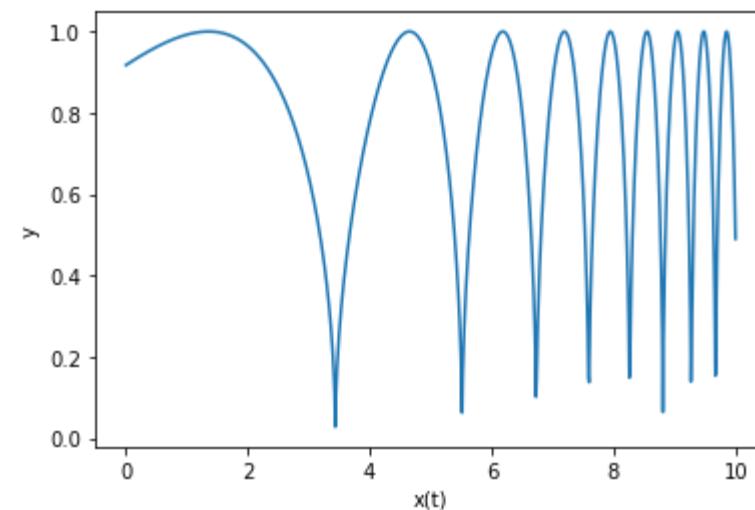
Calling Bullshit – Hoe te misleiden met grafieken en data

Numeriek

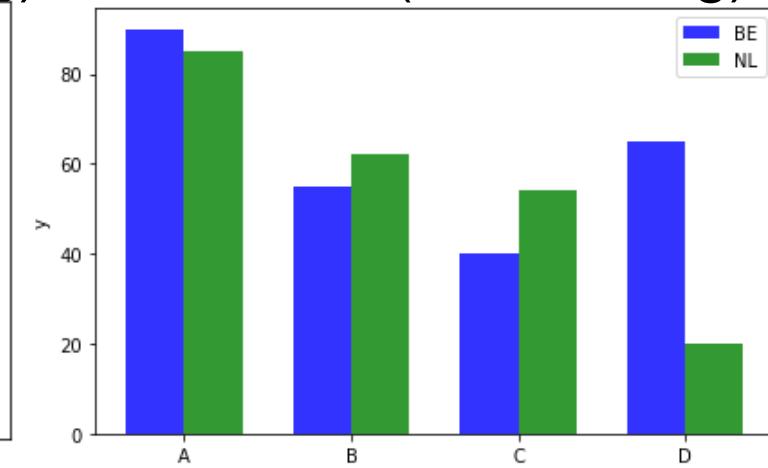
- Scatterplot (verspreiding)



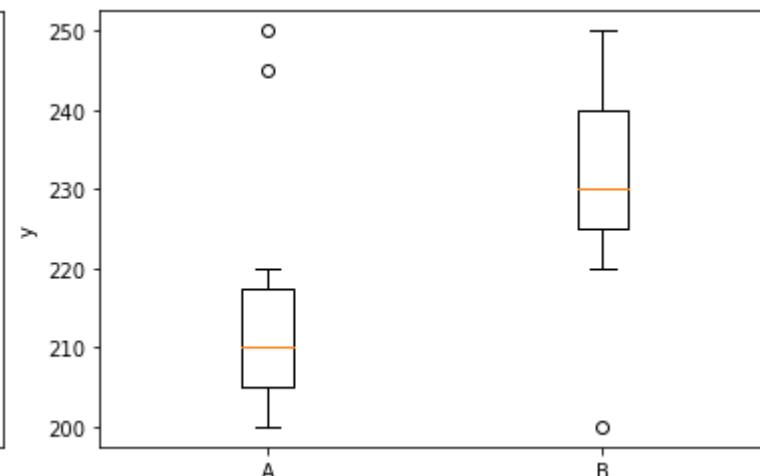
- Lineplot (verandering)



- Barchart (verhouding)



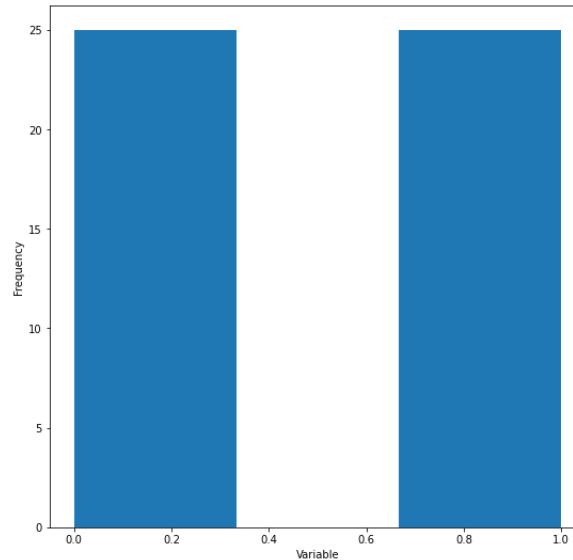
- Boxplot (verdeling)



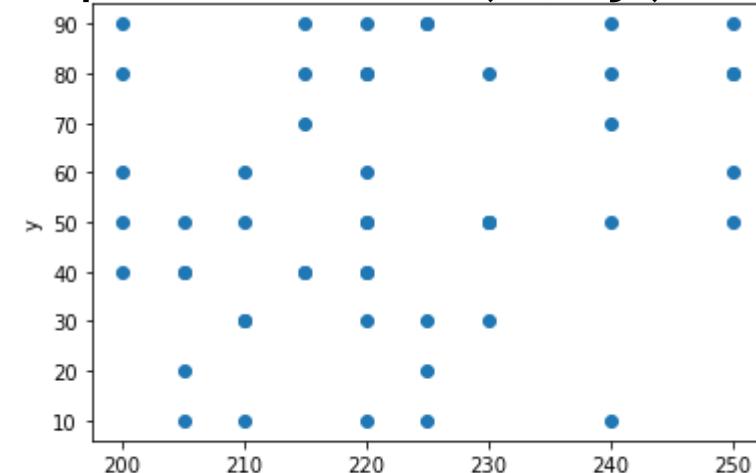
```
import matplotlib.pyplot as plt
```

Data types plotten

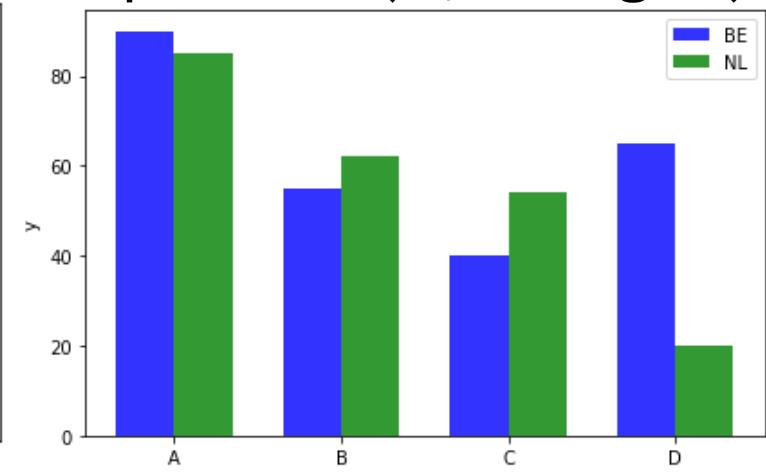
- `plt.hist(x)`



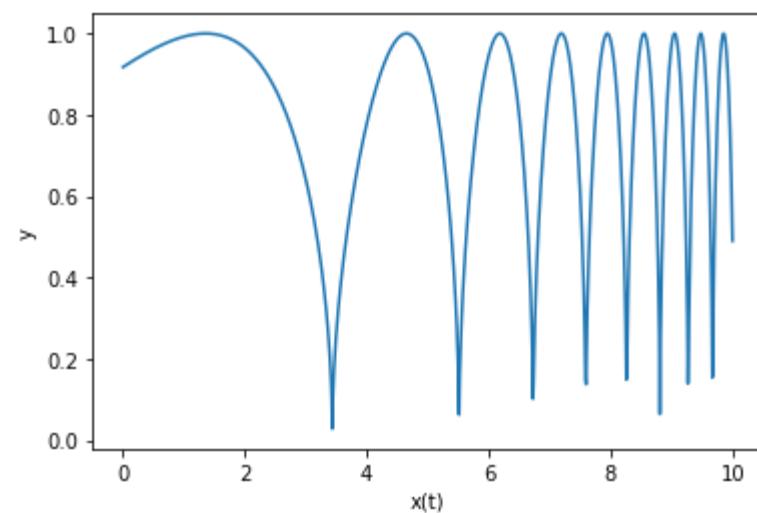
- `plt.scatter(x, y)`



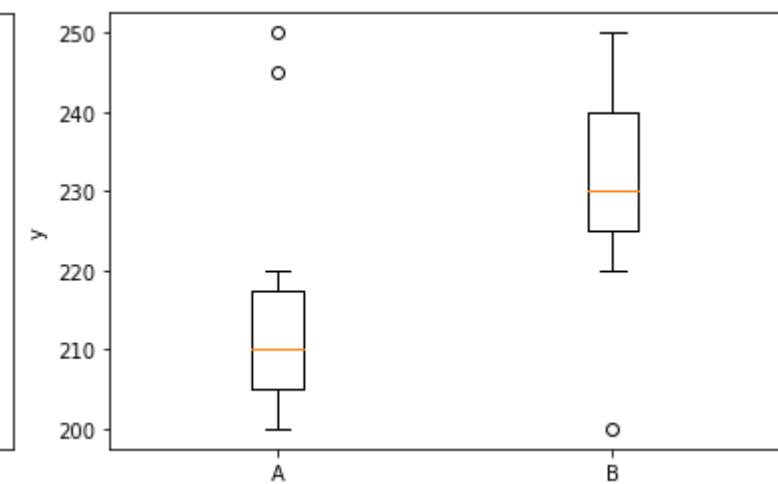
- `plt.bar(x, height)`



- `plt.plot(x, y)`

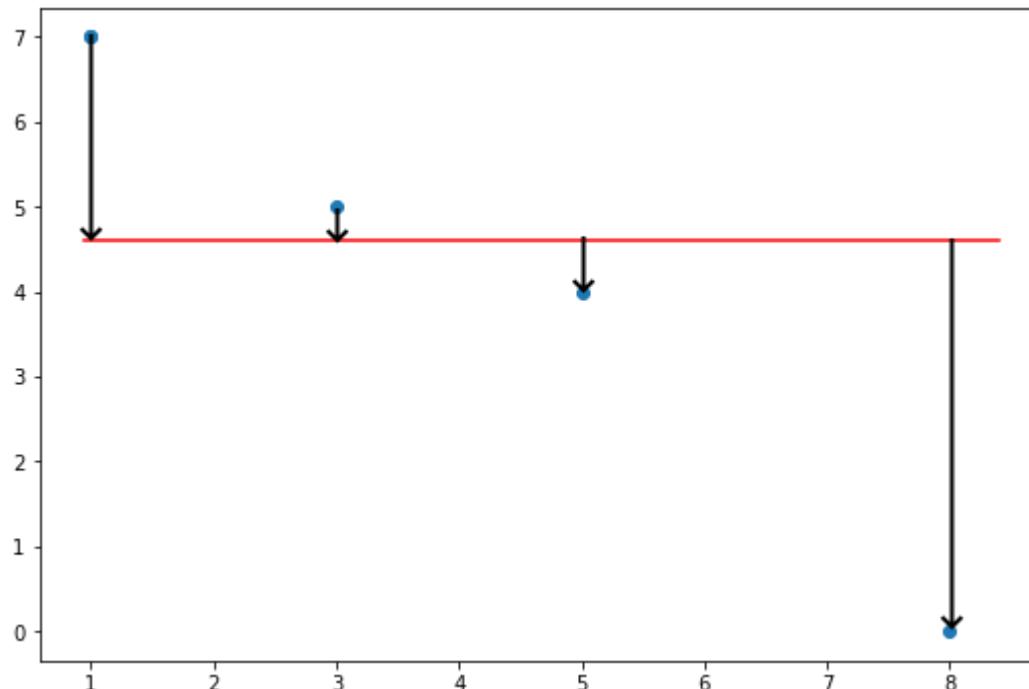


- `plt.boxplot(x)`



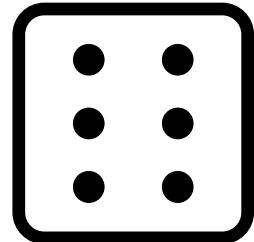
Calling Bullshit – Hoe te misleiden met grafieken en data

- **Gemiddelde** $\mu = \frac{1}{n} \sum_{i=1}^n x_i$
- **Variantie**
Het kwadraat van de gemiddelde afstand van het gemiddelde
- **Standaarddeviatie**
De wortel van de variantie
- **Modus**
De waarde die het meest voorkomt
- **Mediaan**
De waarde die in het midden ligt
- **25%, 50%, 75% percentiel**
Onder waarde X heb je Y% van alle data

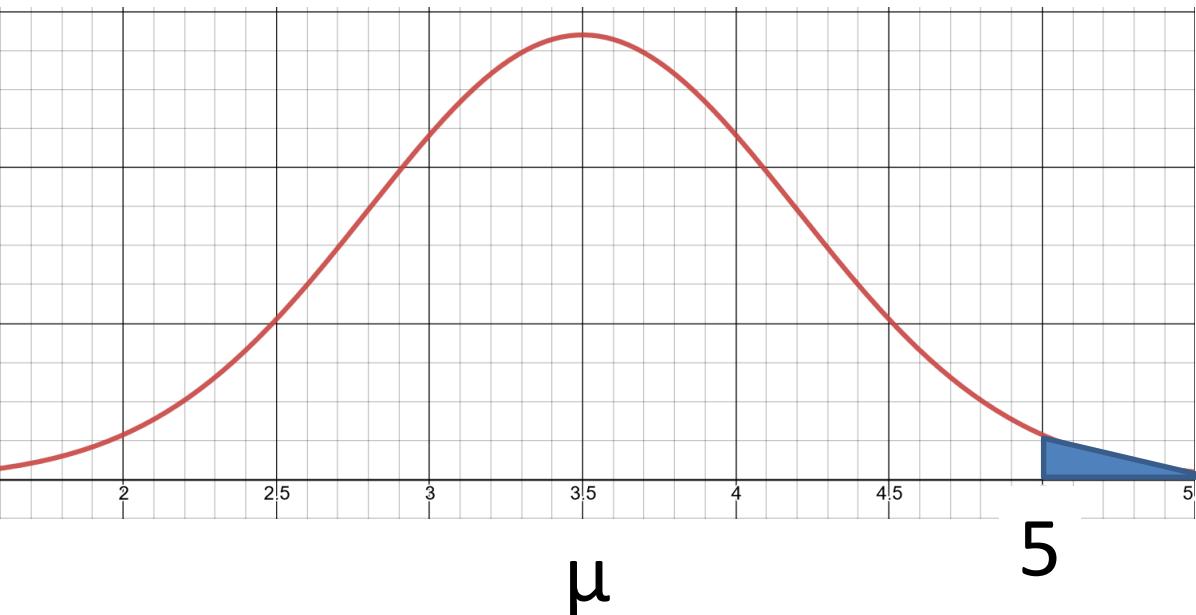


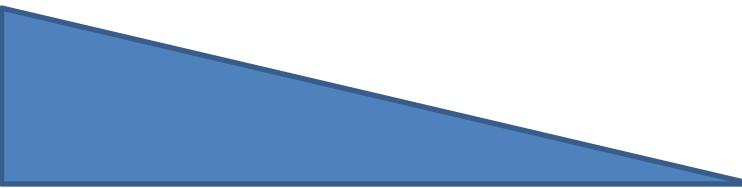
Hypothese toetsen

- Statistische significantie:
De kans dat iets waar is, is groot
OF
De kans dat iets niet waar is, door
toeval, is klein
- Hypothese toetsen
Ik gooi 10 keer. Het gemiddelde van
mijn worpen is 5. Is de dobbelsteen
frauduleus?



kans



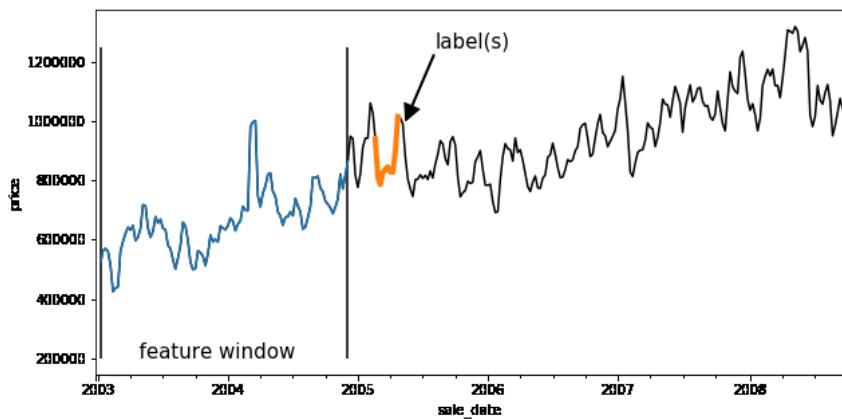
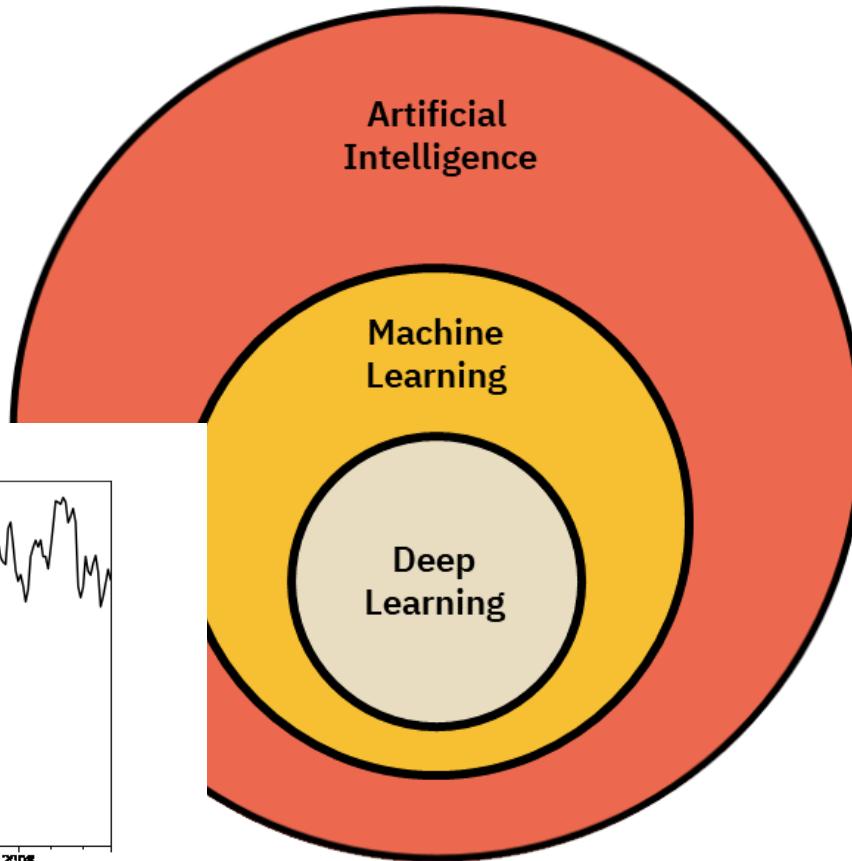
OPP() = Probability (p-waarde)

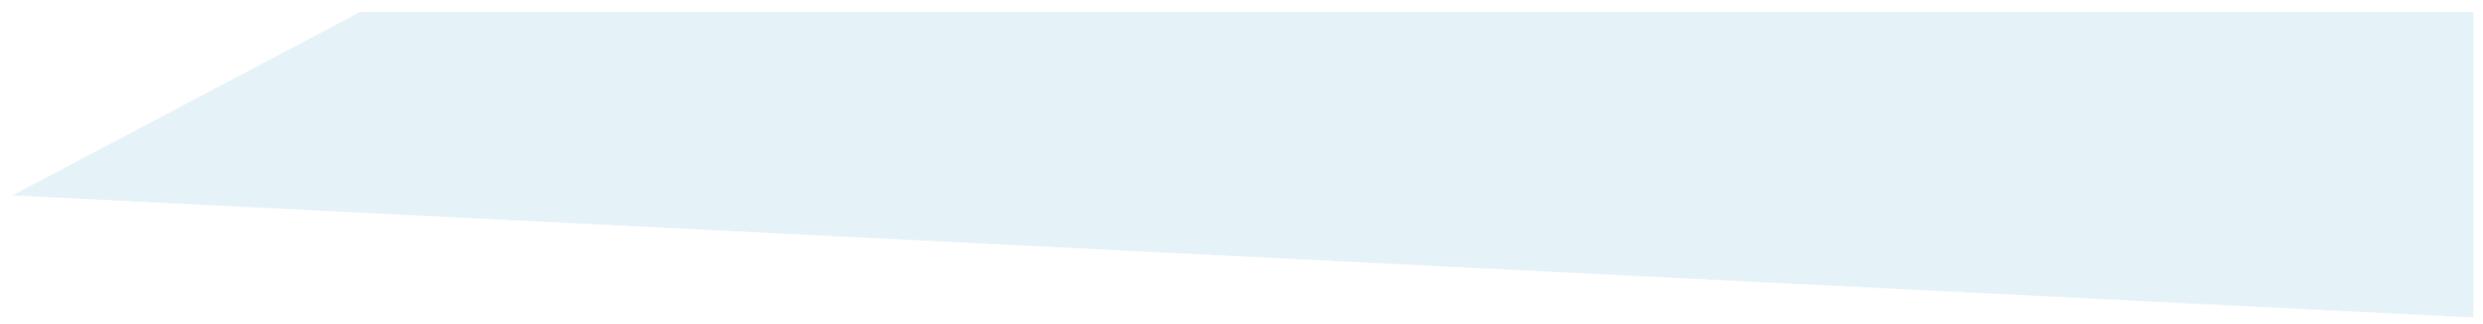
$p \leq 0.05 \rightarrow$ statistisch significant
(geen normale dobbelsteen)
 $p \geq 0.05 \rightarrow$ onvoldoende bewijs

Artificial Intelligence
Namaken van intelligentie

Machine Learning
Namaken hoe we leren

Het verschil tussen AI en ML





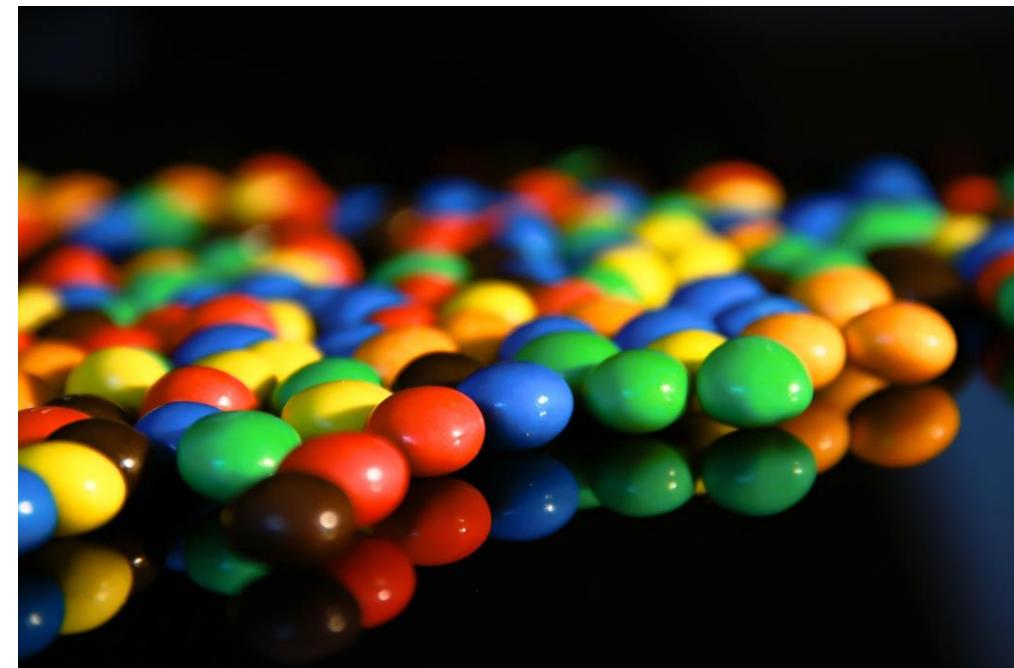
**JE ZULT ME VEEL TERMEN HOREN
GEBRUIKEN DIE EIGENLIJK VAN
TOEPASSING ZIJN OP EEN MENS**

Twee soorten machine learning

Supervised learning



Unsupervised learning



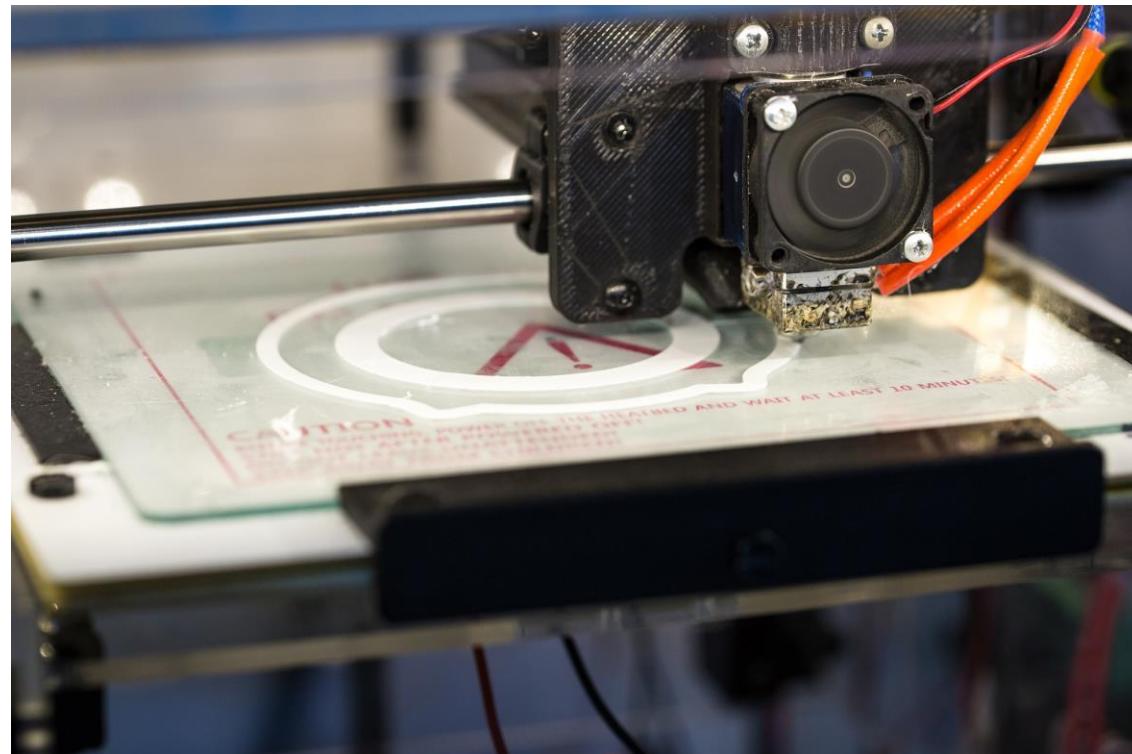
3 verschillende methoden

3 verschillende *classificatie* methoden¹:

- Decision Tree/Random Forest
- Support Vector Machine
- Neural Networks

1: Er zijn meer methoden, maar deze komen het meest voor.

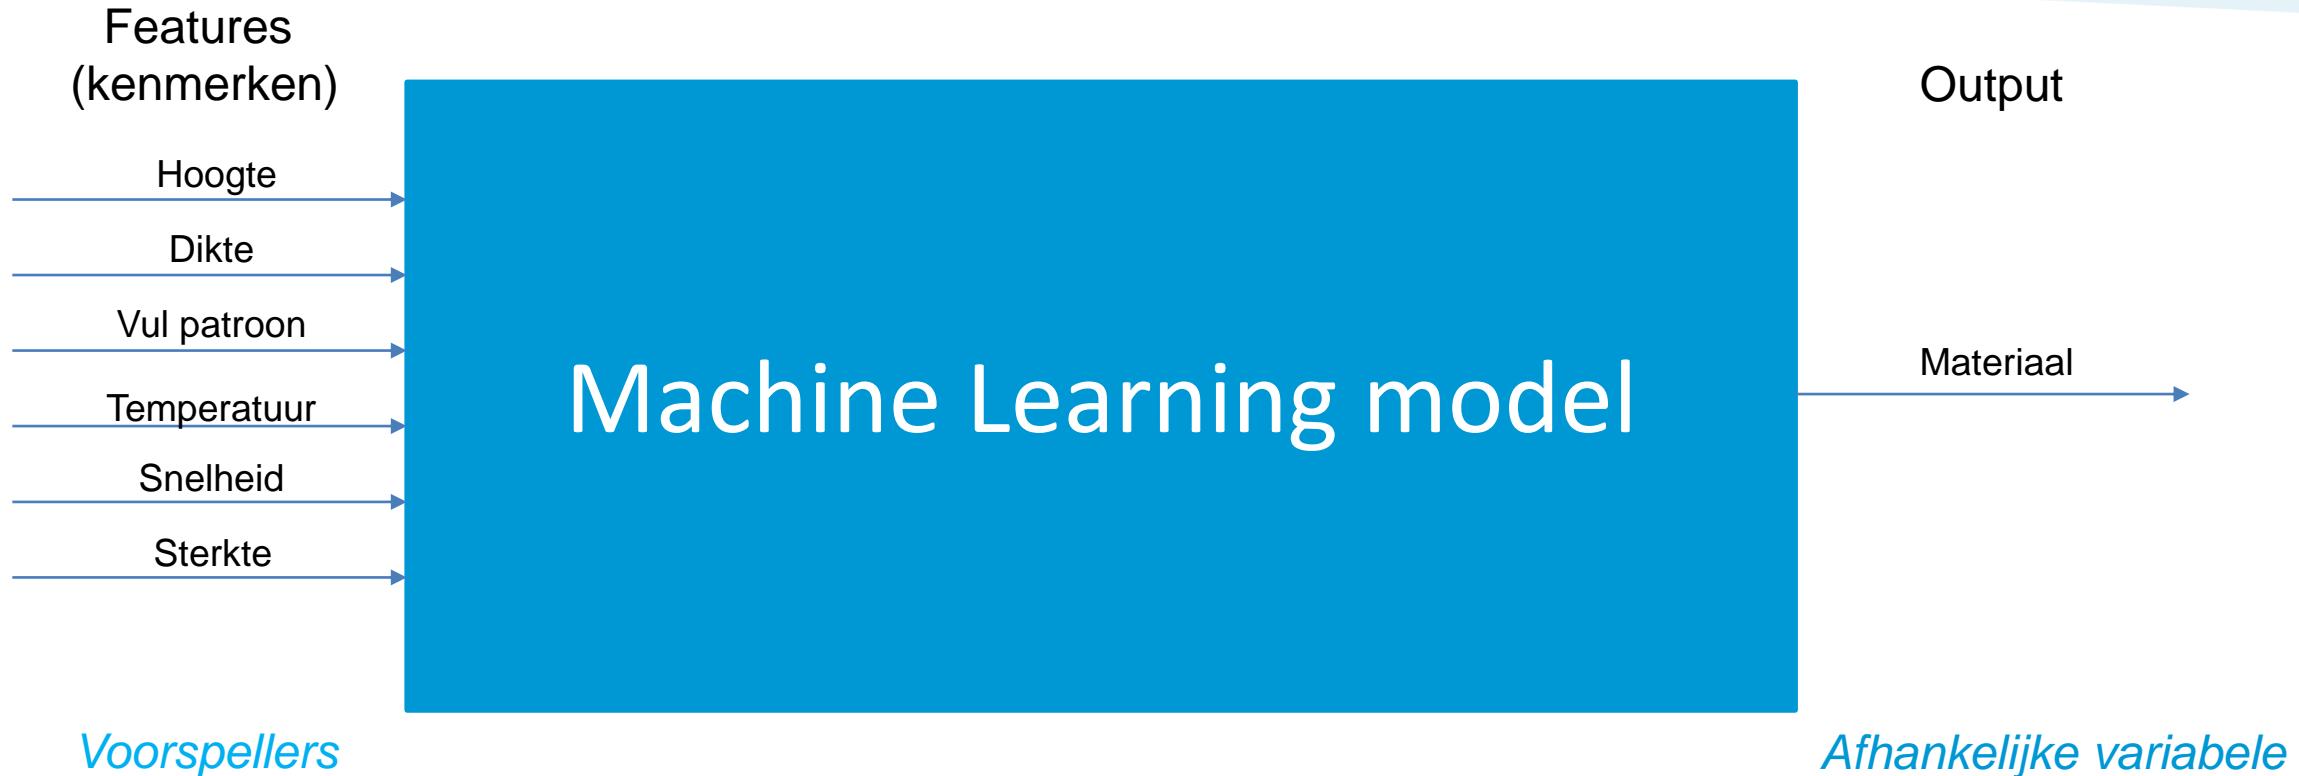
Voorbeeld casus



3D printer:

- https://www.kaggle.com/afumetto/3d_printer (staat ook op github)
- Eigenschappen van printopdracht:
 - Hoogte
 - Temperatuur
 - Snelheid
 - Etc.
- Kunnen we met Machine Learning raden met welk materiaal geprint is?

Het idee achter elk machine learning model



$$Y = f(X)$$

3D printer dataset

	layer_height	wall_thickness	infill_density	infill_pattern	nozzle_temperature	bed_temperature	print_speed	fan_speed	roughness	tension_strenght	elongation		
0	0.02	8	90	1	220	60	40	0	25	18	1.2	0	1
1	0.02	7	90	0	225	65	40	25	32	16	1.4	1	1
2	0.02	1	80	1	230	70	40	50	40	8	0.8	2	1
3	0.02	4	70	0	240	75	40	75	68	10	0.5	3	1
4	0.02	6	90	1	250	80	40	100	92	5	0.7	4	1

Name: material, dtype: int64

X

Kunnen we aan de hand van de laag hoogte, print dikte etc. voorspellen met welk materiaal is geprint?

Y

Materiaal
1 = ABS
0 = PLA

- <https://github.com/SimchaGD/AIVE>
- Download als ZIP
- Zet in map documenten (of andere map) en pak uit

The screenshot shows a GitHub repository page for "SimchaGD / AIVE" and a file manager interface.

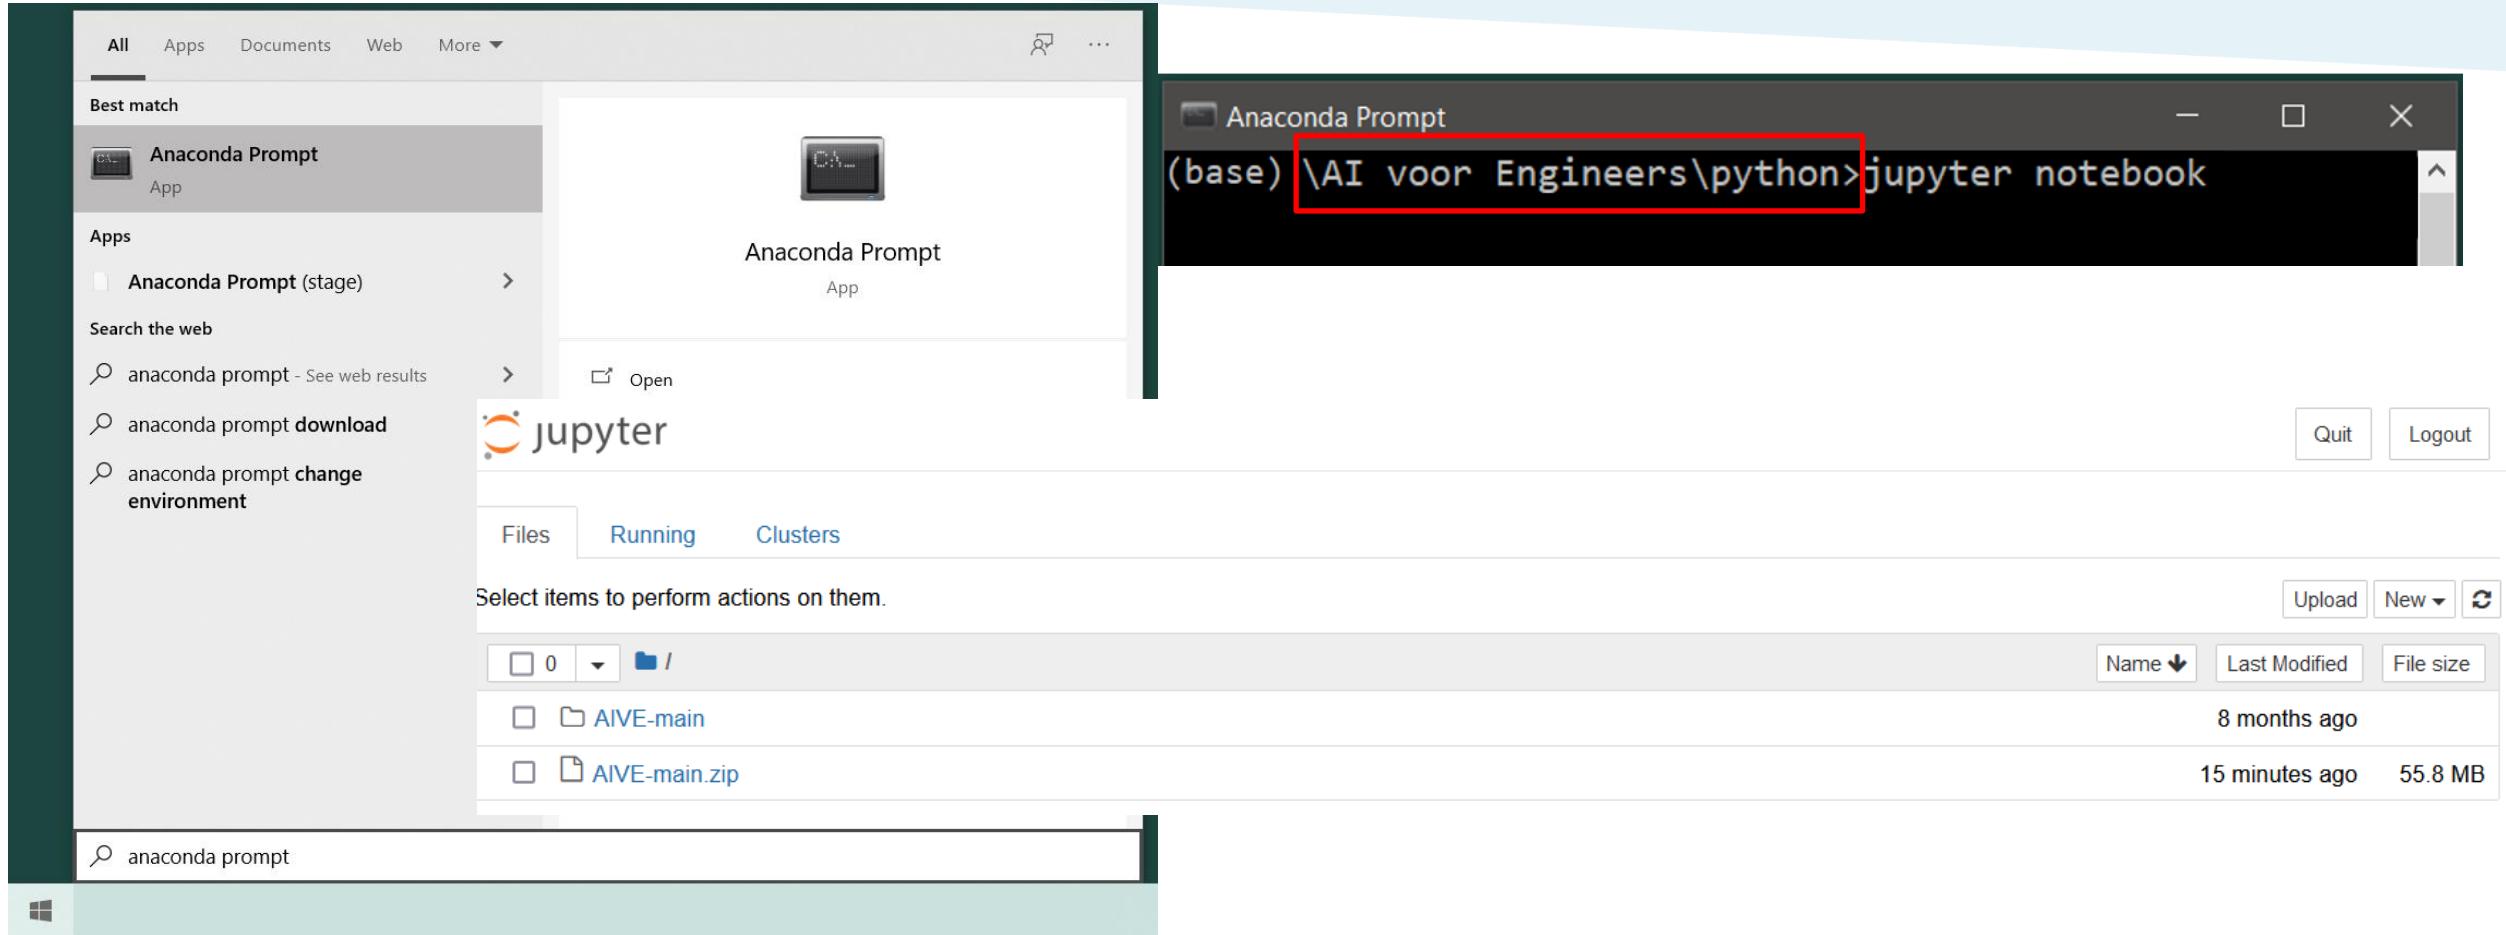
GitHub Repository Page:

- Header:** SimchaGD / AIVE, Projects, Wiki, Security, Insights, Settings.
- Code Button:** A green button labeled "Code" with a downward arrow, circled with orange number 1.
- Clone Section:** Clone options via HTTPS, SSH, GitHub CLI, or SVN.
- Download ZIP:** A red box highlights the "Download ZIP" button, circled with orange number 2.
- File List:**
 - SimchaGD Update README.md (Initial commit)
 - 3dprinter
 - dag 1
 - dag 2
 - dag 3
 - Update README.md

File Manager Interface:

- Toolbar:** Home, Share, View, Compressed Folder Tools (highlighted with a pink border).
- Folders:** mega, Prive, datasets, Tau Omega projecten, datasets, Al voor Engineers, Bisc.
- Extract To:** A dropdown menu showing "Extract all" and "Extract to <> AI voor Engineers > python". The dropdown menu is highlighted with a red box.
- Search:** Search bar with placeholder "Search python".
- Table:** A list of files and folders.

Name	Status	Date modified	Type	Size
AIVE-main.zip	⟳	8-7-2021 08:54	Compressed (zipp...)	54,493 KB



Hoe laden we data in?

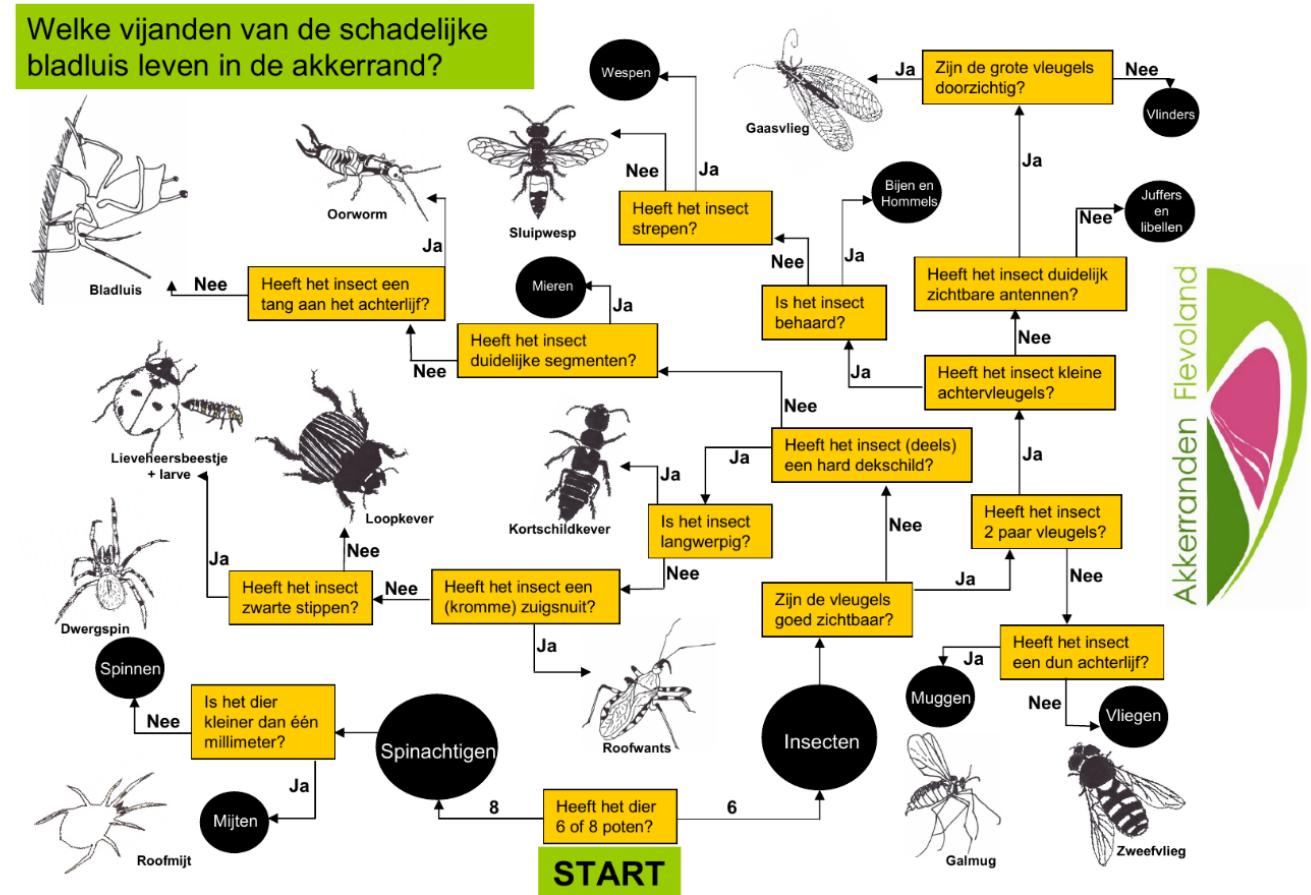
```
import pandas as pd  
filename = "3dprinter/data.csv"  
df = pd.read_csv(filename, sep = ';')  
df.head()
```

	layer_height	wall_thickness	infill_density	infill_pattern	nozzle_temperature	bed_temperature	print_speed	fan_speed	roughness	tension_strenght	elongation
0	0.02	8	90	1	220	60	40	0	25	18	1.2
1	0.02	7	90	0	225	65	40	25	32	16	1.4
2	0.02	1	80	1	230	70	40	50	40	8	0.8
3	0.02	4	70	0	240	75	40	75	68	10	0.5
4	0.02	6	90	1	250	80	40	100	92	5	0.7

0 = grid

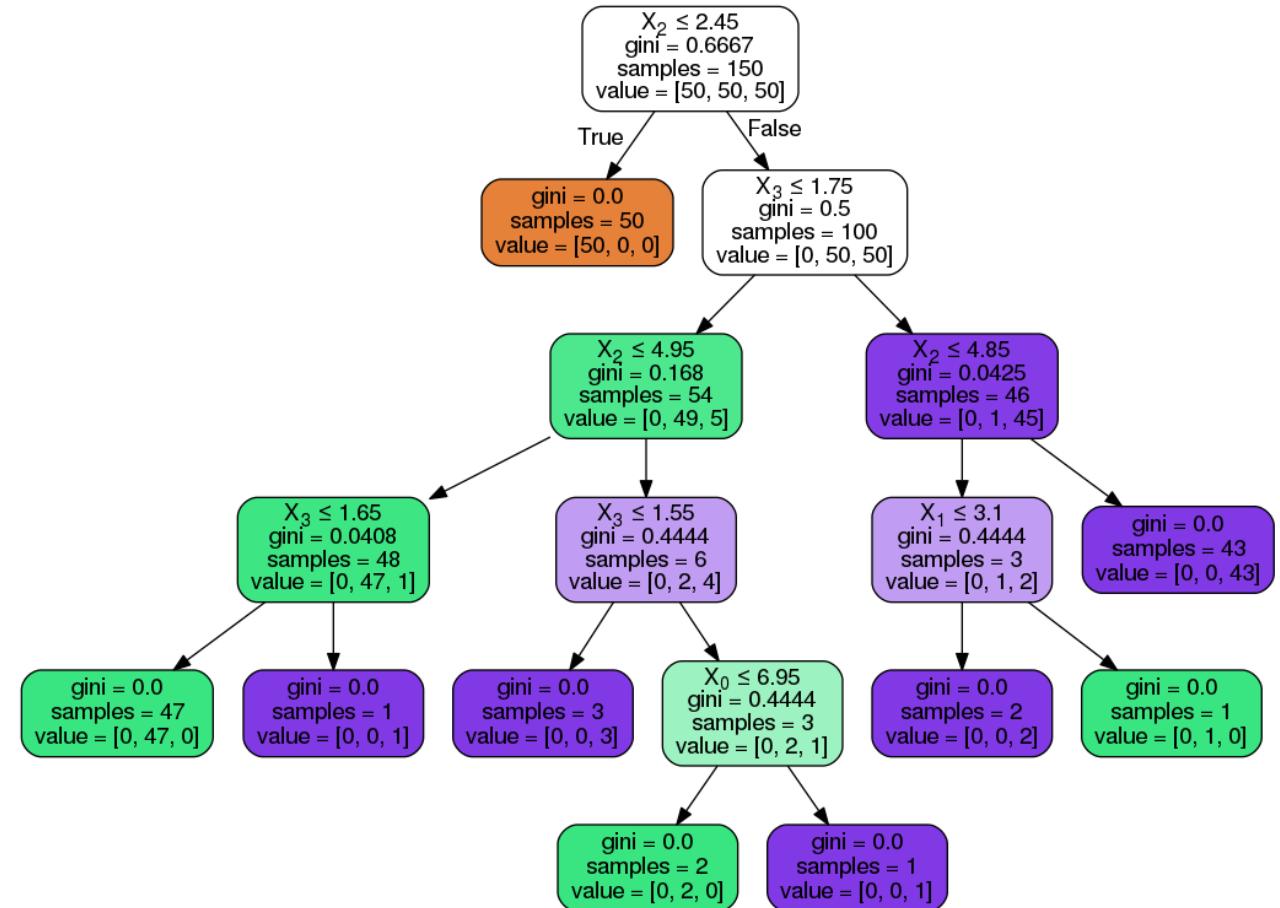
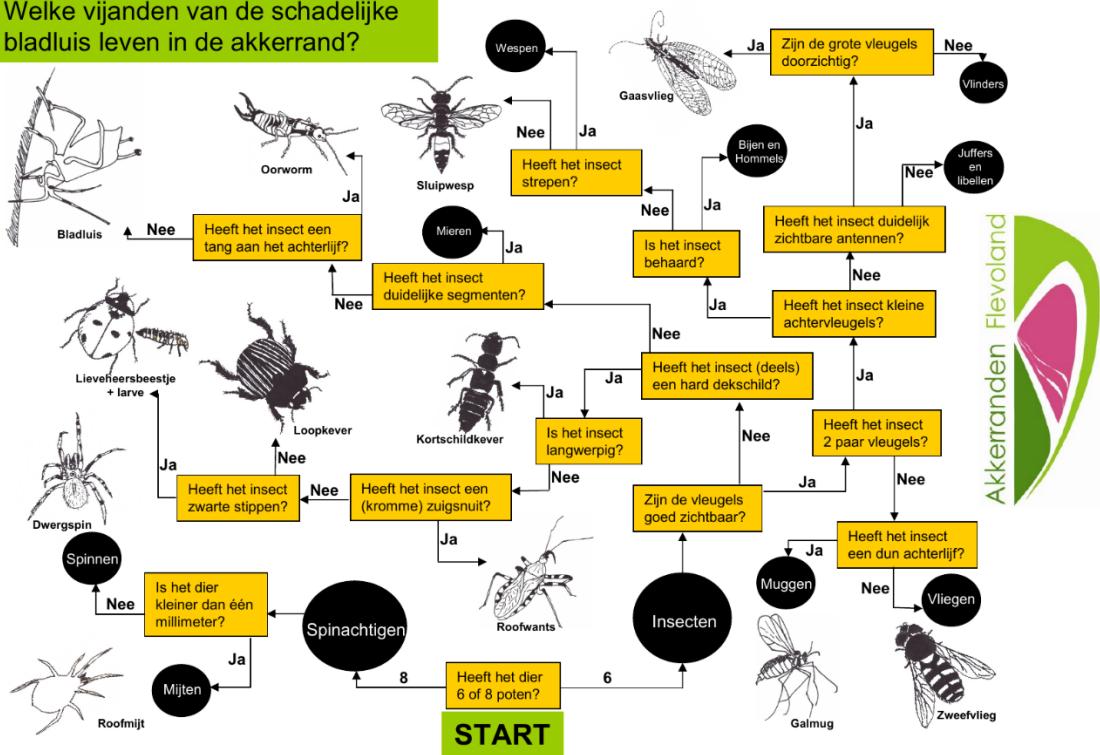
1 = honeycomb/hexagon

Model 1:
DECISION TREES

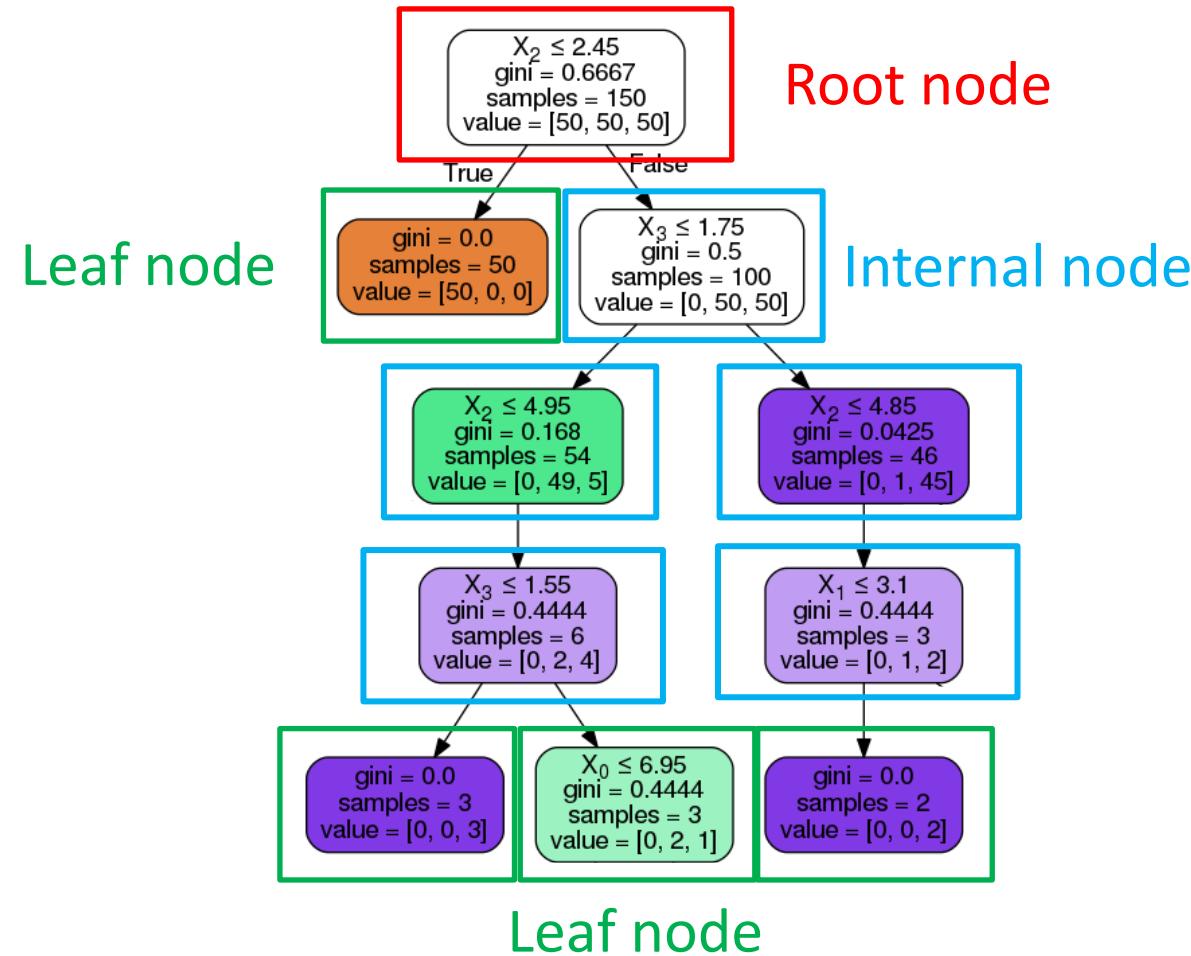


Decision Trees

Welke vijanden van de schadelijke bladluis leven in de akkerrand?



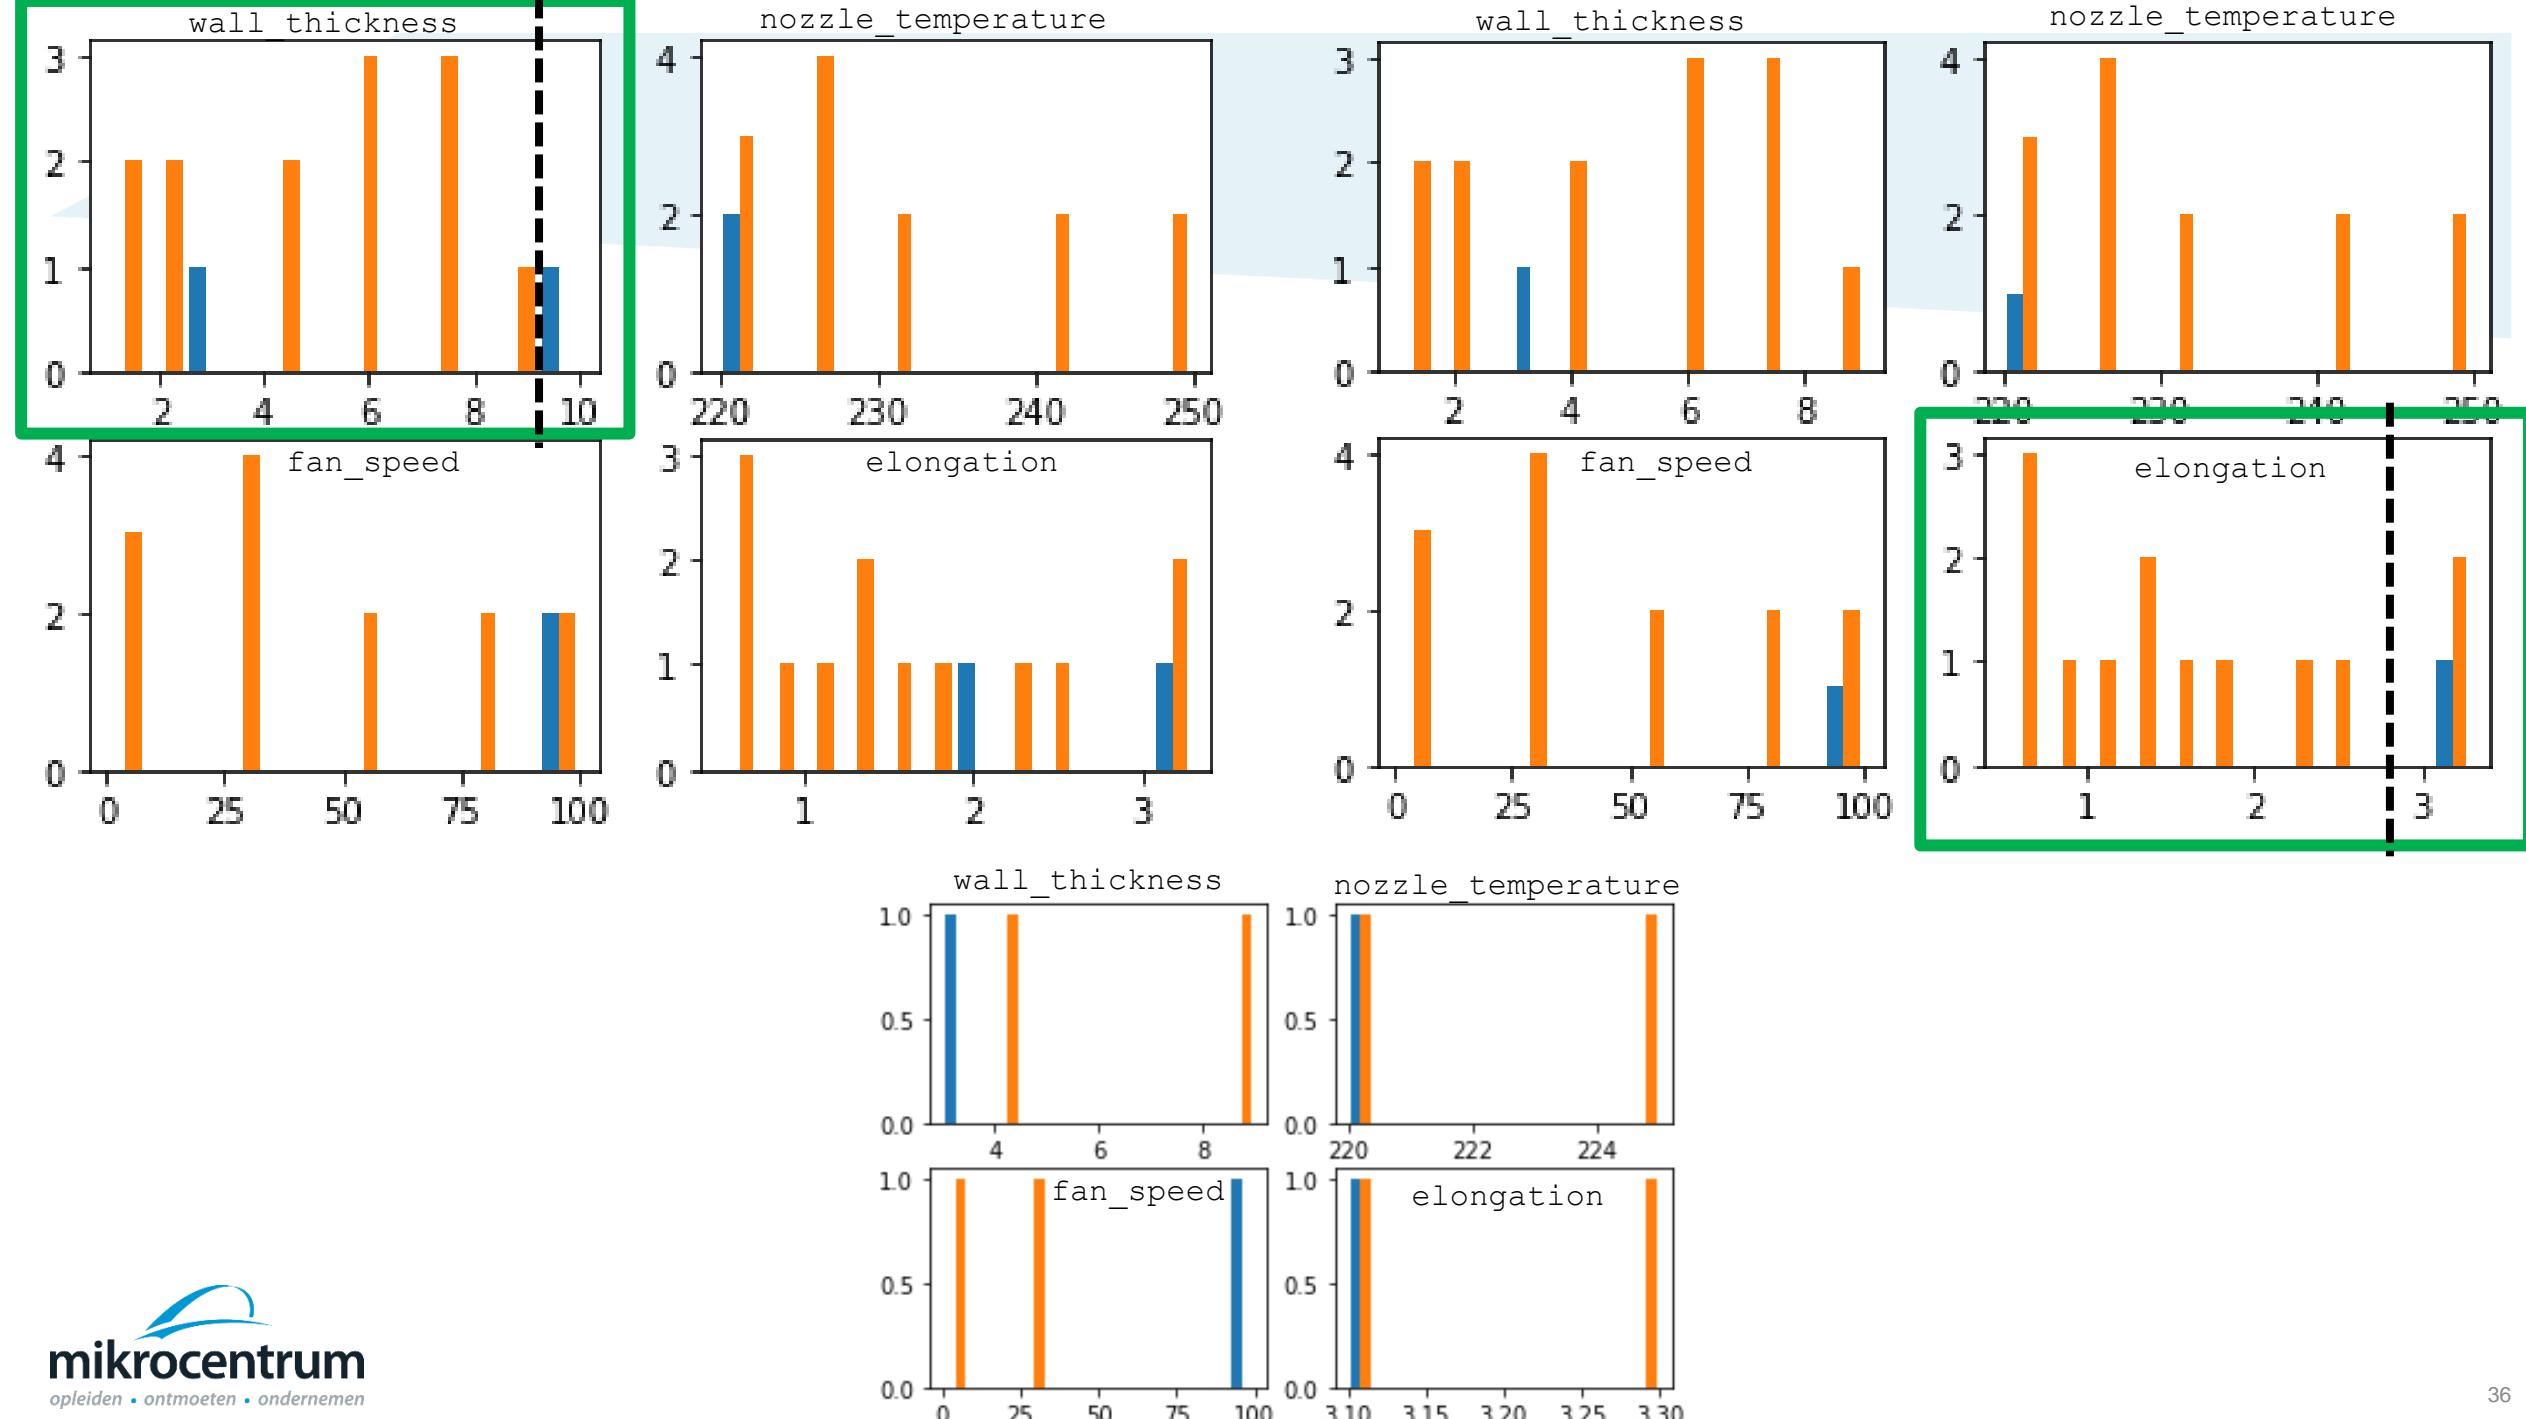
Terminologie

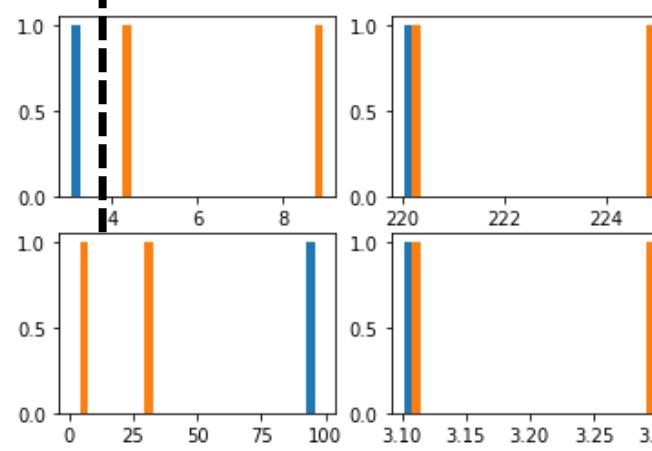
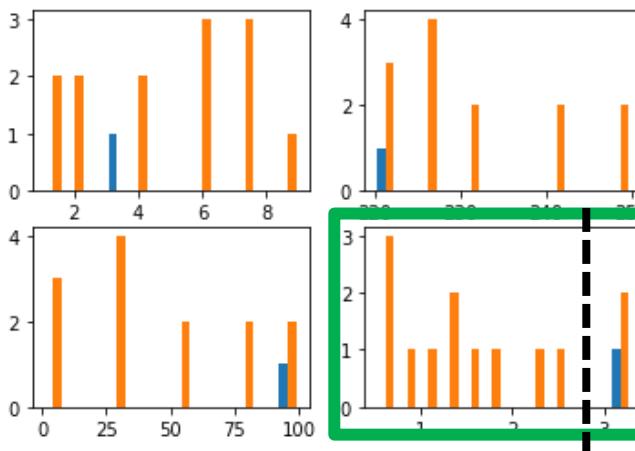
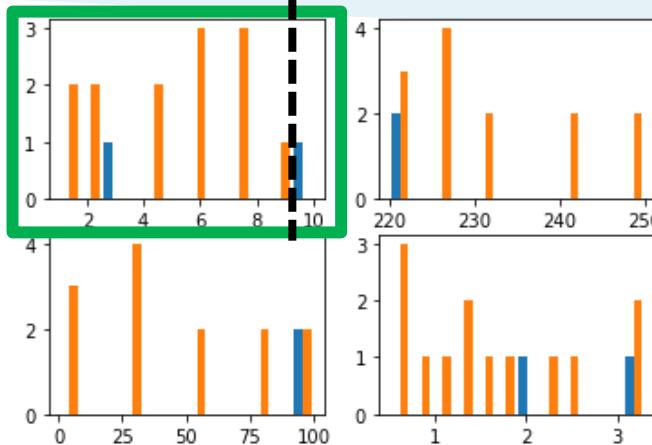
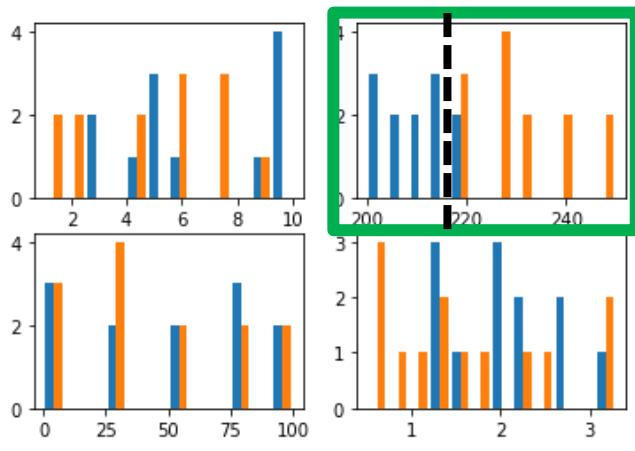


Met de hand

PLA

ABS





$X[1] \leq 217.5$
gini = 0.499
samples = 25
value = [12, 13]

gini = 0.0
samples = 10
value = [10, 0]

$X[0] \leq 9.5$
gini = 0.231
samples = 15
value = [2, 13]

$X[3] \leq 2.75$
gini = 0.133
samples = 14
value = [1, 13]

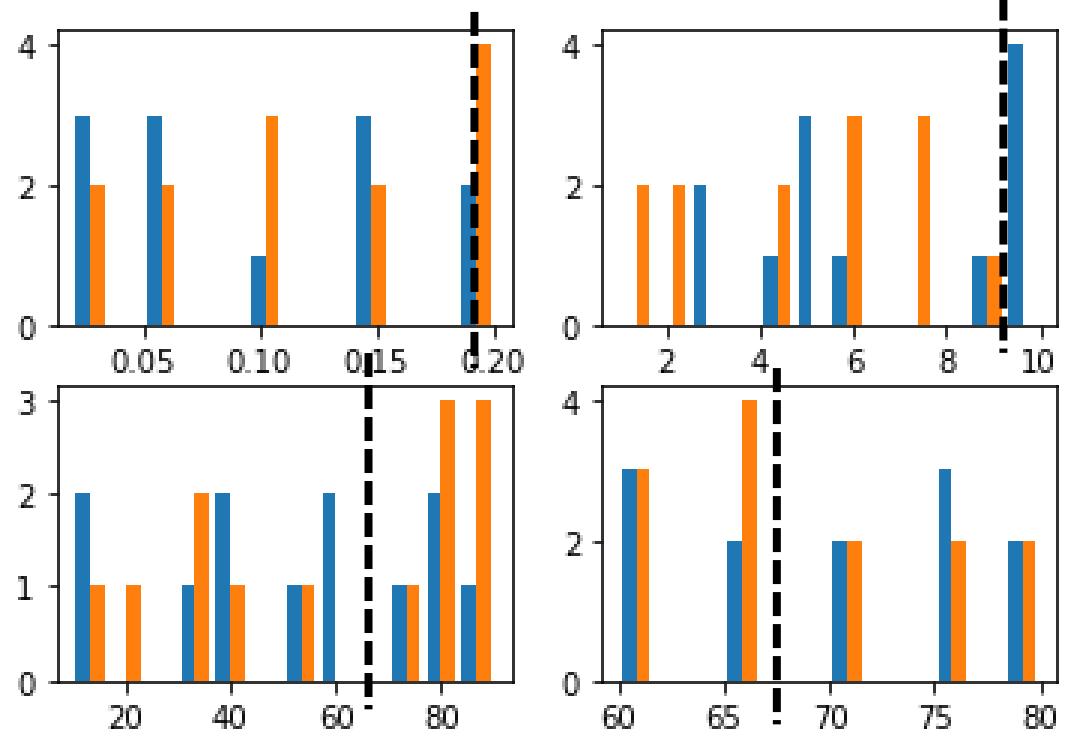
gini = 0.0
samples = 1
value = [1, 0]

gini = 0.0
samples = 11
value = [0, 11]

$X[0] \leq 3.5$
gini = 0.444
samples = 3
value = [1, 2]

gini = 0.0
samples = 1
value = [1, 0]

gini = 0.0
samples = 2
value = [0, 2]

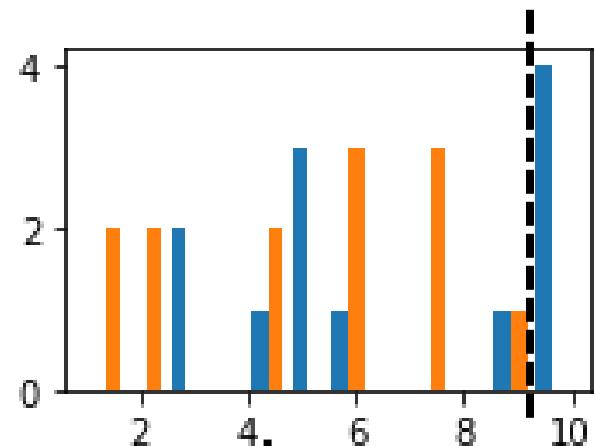
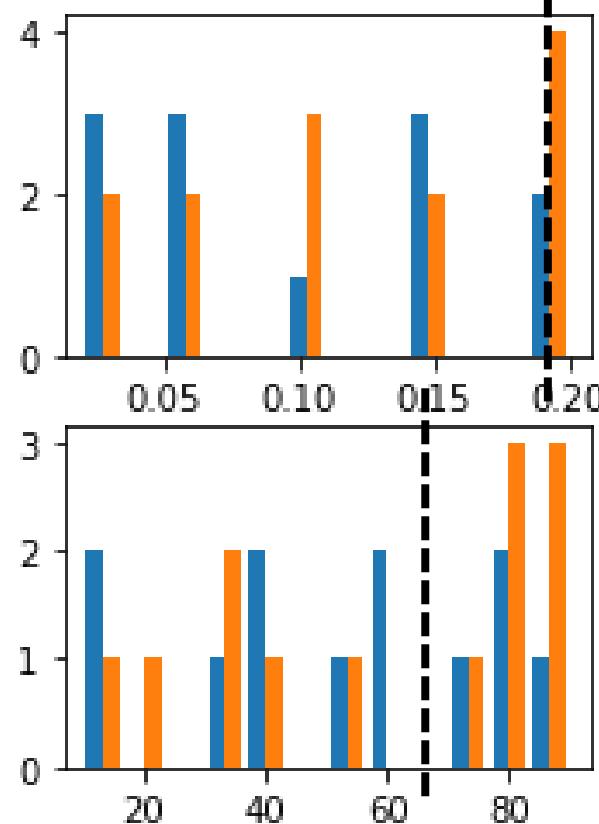


?

$$\text{Gini} = 1 - \sum_{i=1}^n p_i^2$$

> 0.19	
ABS	PLA
12	9

ABS		ABS	PLA
0	4		

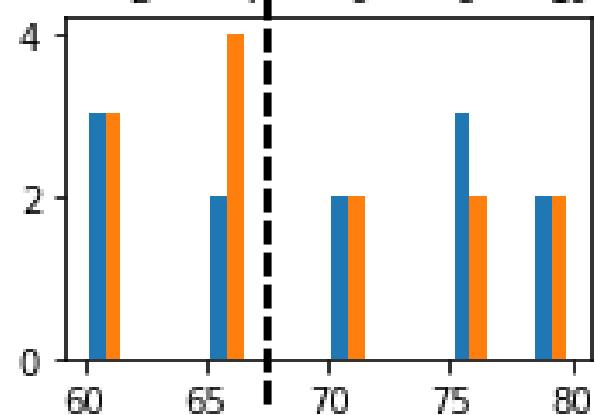
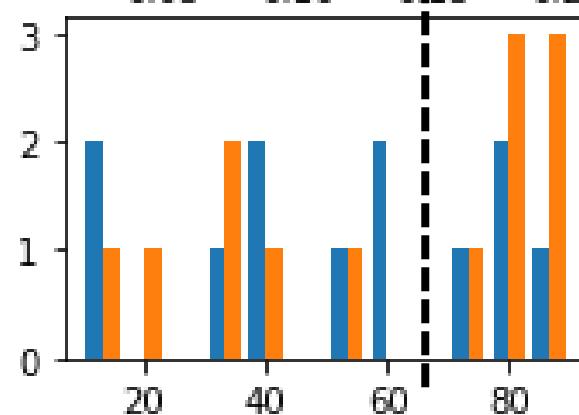


> 9.5	
ABS	PLA
8	13

ABS		ABS	PLA
4	0		

> 65	
ABS	PLA
8	8

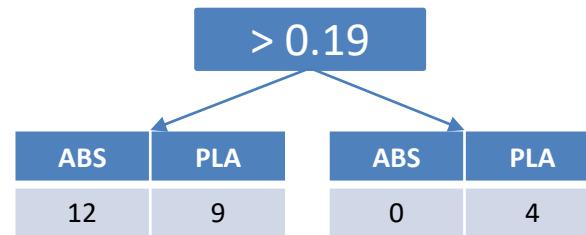
ABS		ABS	PLA
4	7		



> 67.5	
ABS	PLA
5	7

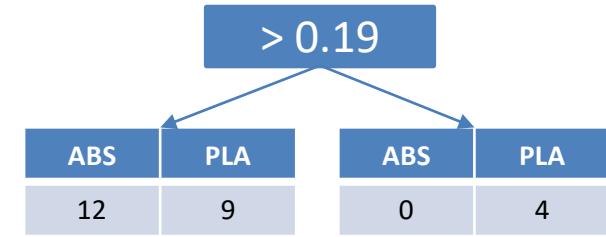
ABS		ABS	PLA
7	6		

$$\text{Gini} = 1 - \sum_{i=1}^n p_i^2 = 1 - P(\text{ABS})^2 - P(\text{PLA})^2$$

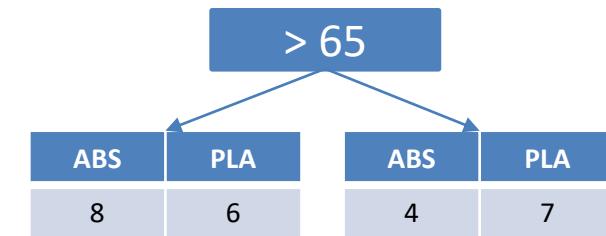
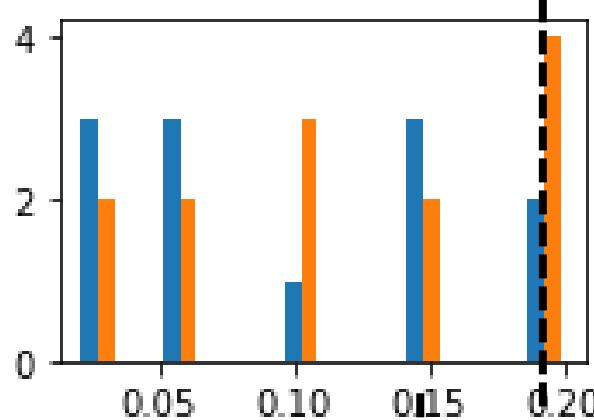


$$1 - \left(\frac{12}{12+9} \right)^2 - \left(\frac{9}{12+9} \right)^2 = 0.490 \quad 1 - \left(\frac{4}{4+0} \right)^2 - \left(\frac{0}{4+0} \right)^2 = 0$$

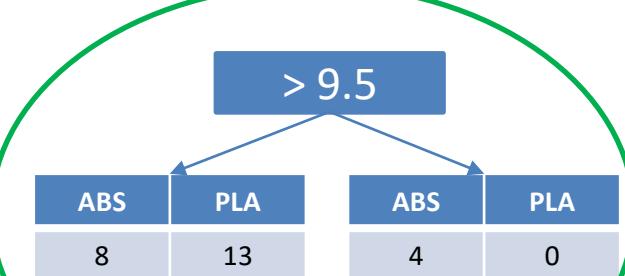
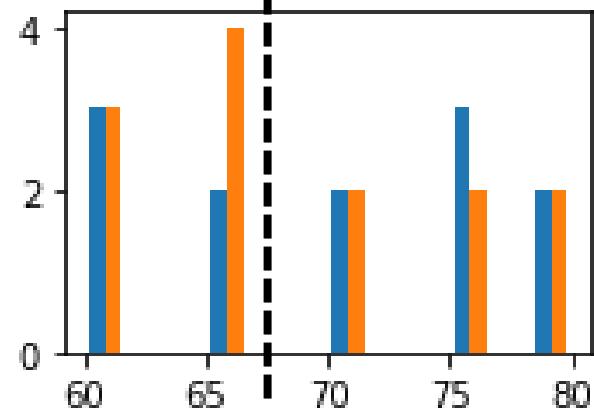
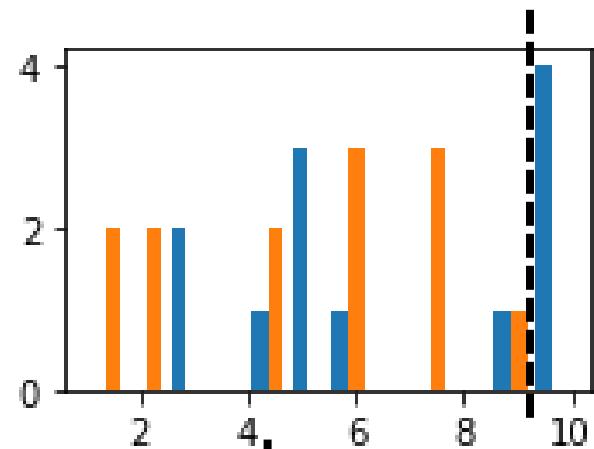
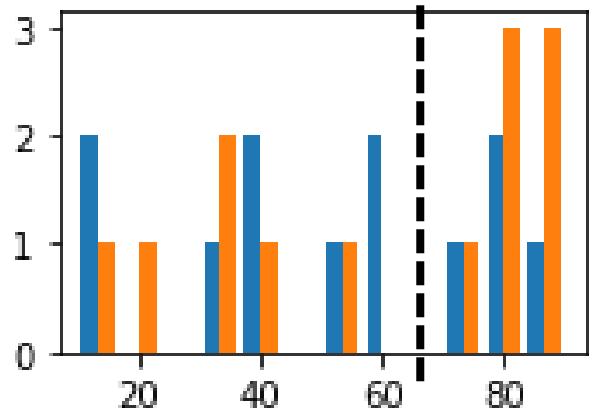
$$\left(\frac{21}{25} \right) \times 0.490 + \left(\frac{4}{25} \right) \times 0 = 0.412$$



Gini = 0.412



Gini = 0.478



Gini = 0.396



Gini = 0.492

In: `from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score`

```
clf = DecisionTreeClassifier()  
y_pred = clf.fit(X, Y).predict(X)  
print(accuracy_score(y_pred, Y))
```

Out: 1.0

Laad de 3D-printer data in met Pandas

Splits de voorspellers van de afhankelende variabelen.

Maak een Decision Tree met behulp van de SciKit-Learn library

~10 minuten

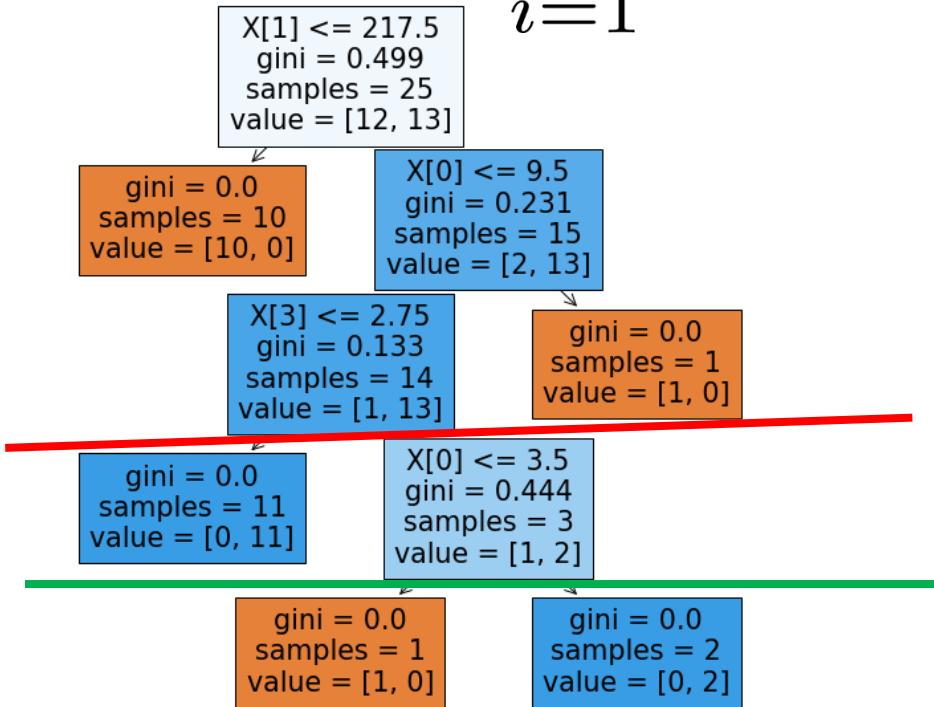
Hyperparameters

- criterion = ("gini", "entropy")

$$\text{Gini} = 1 - \sum_{i=1}^n p_i^2$$

- max_depth = None
- **max_depth = 2**
- min_samples_split = 2
- **min_samples_split = 6**

$$\text{Entropy} = - \sum_{i=1}^n p_i \log(p_i)$$



```
In: dt_clf = DecisionTreeClassifier(max_depth=4)
y_pred = dt_clf.fit(x_train, y_train).predict(x_test)
print(accuracy_score(y_pred, y_test))
```

Out: 0.8

Train test split

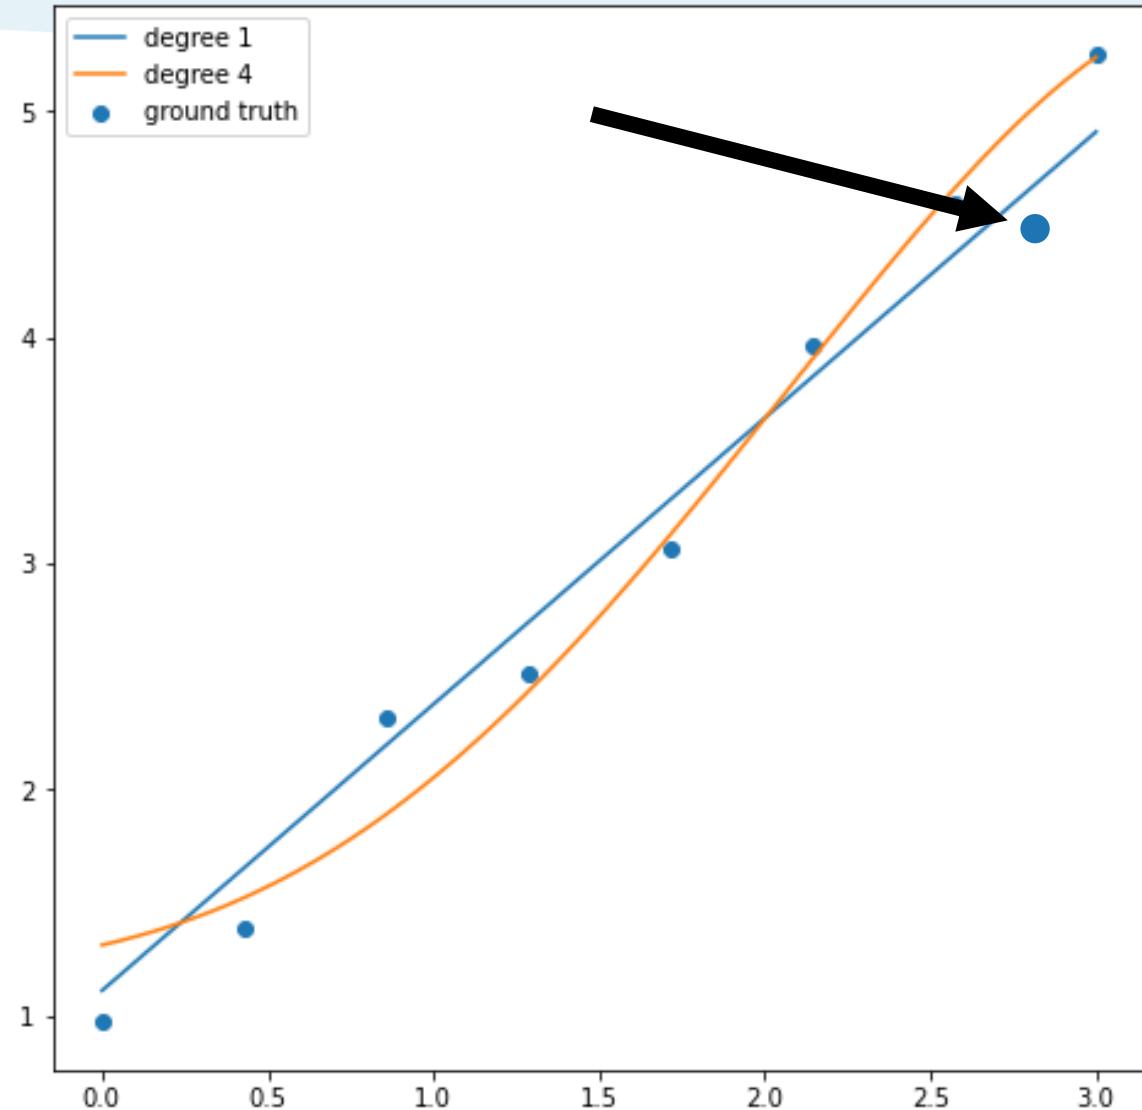


Overfitting:

Creating over-complex models
that do not generalize well to new data

Het creëren van over-complexe modellen
die niet generaliseerbaar zijn op nieuwe data

https://en.wikipedia.org/wiki/Gambler%27s_fallacy



Train test split



Voorbeeld

Hele dataset

1
2
3
4
5
6
7
8
9
10

Train test split
Verhouding 70/30

Trainset

3
5
2
8
6
10
9

Test

4
1
7

Train test split

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(X, Y,
                                                    test_size = 0.5,
                                                    random_state = 2)

clf = DecisionTreeClassifier()
y_pred = clf.fit(x_train, y_train).predict(x_test)
print(accuracy_score(y_test, y_pred))
```

Voordelen

- Makkelijk/Simpel
- Goede resultaten met relatief weinig data
- Snel
- Goed met *numerieke en categorische data*
- *Glass box*

Nadelen

- *Niet robuust*
- Erg gevoelig voor overfitting
- Optimale oplossing niet altijd gegarandeerd
- Kan *biased* uitkomsten produceren

Oefening met een dataset

Maak een *decision tree classifier* die voorspelt met welk materiaal is geprint

Split je dataset in een *train-* en *testset*

Bereken de *accuracy*

~ 10 minuten

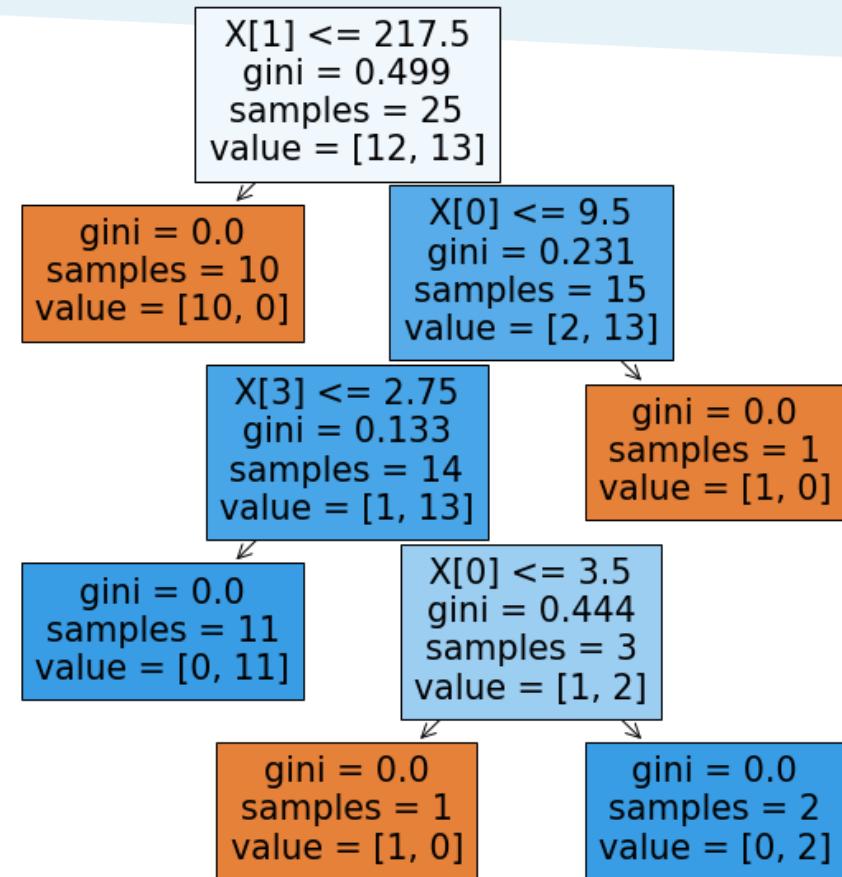
Model opslaan en inladen

```
import pickle
s = pickle.dumps(clf)
with open("decision_tree.sav", "wb") as f:
    f.write(s)
```

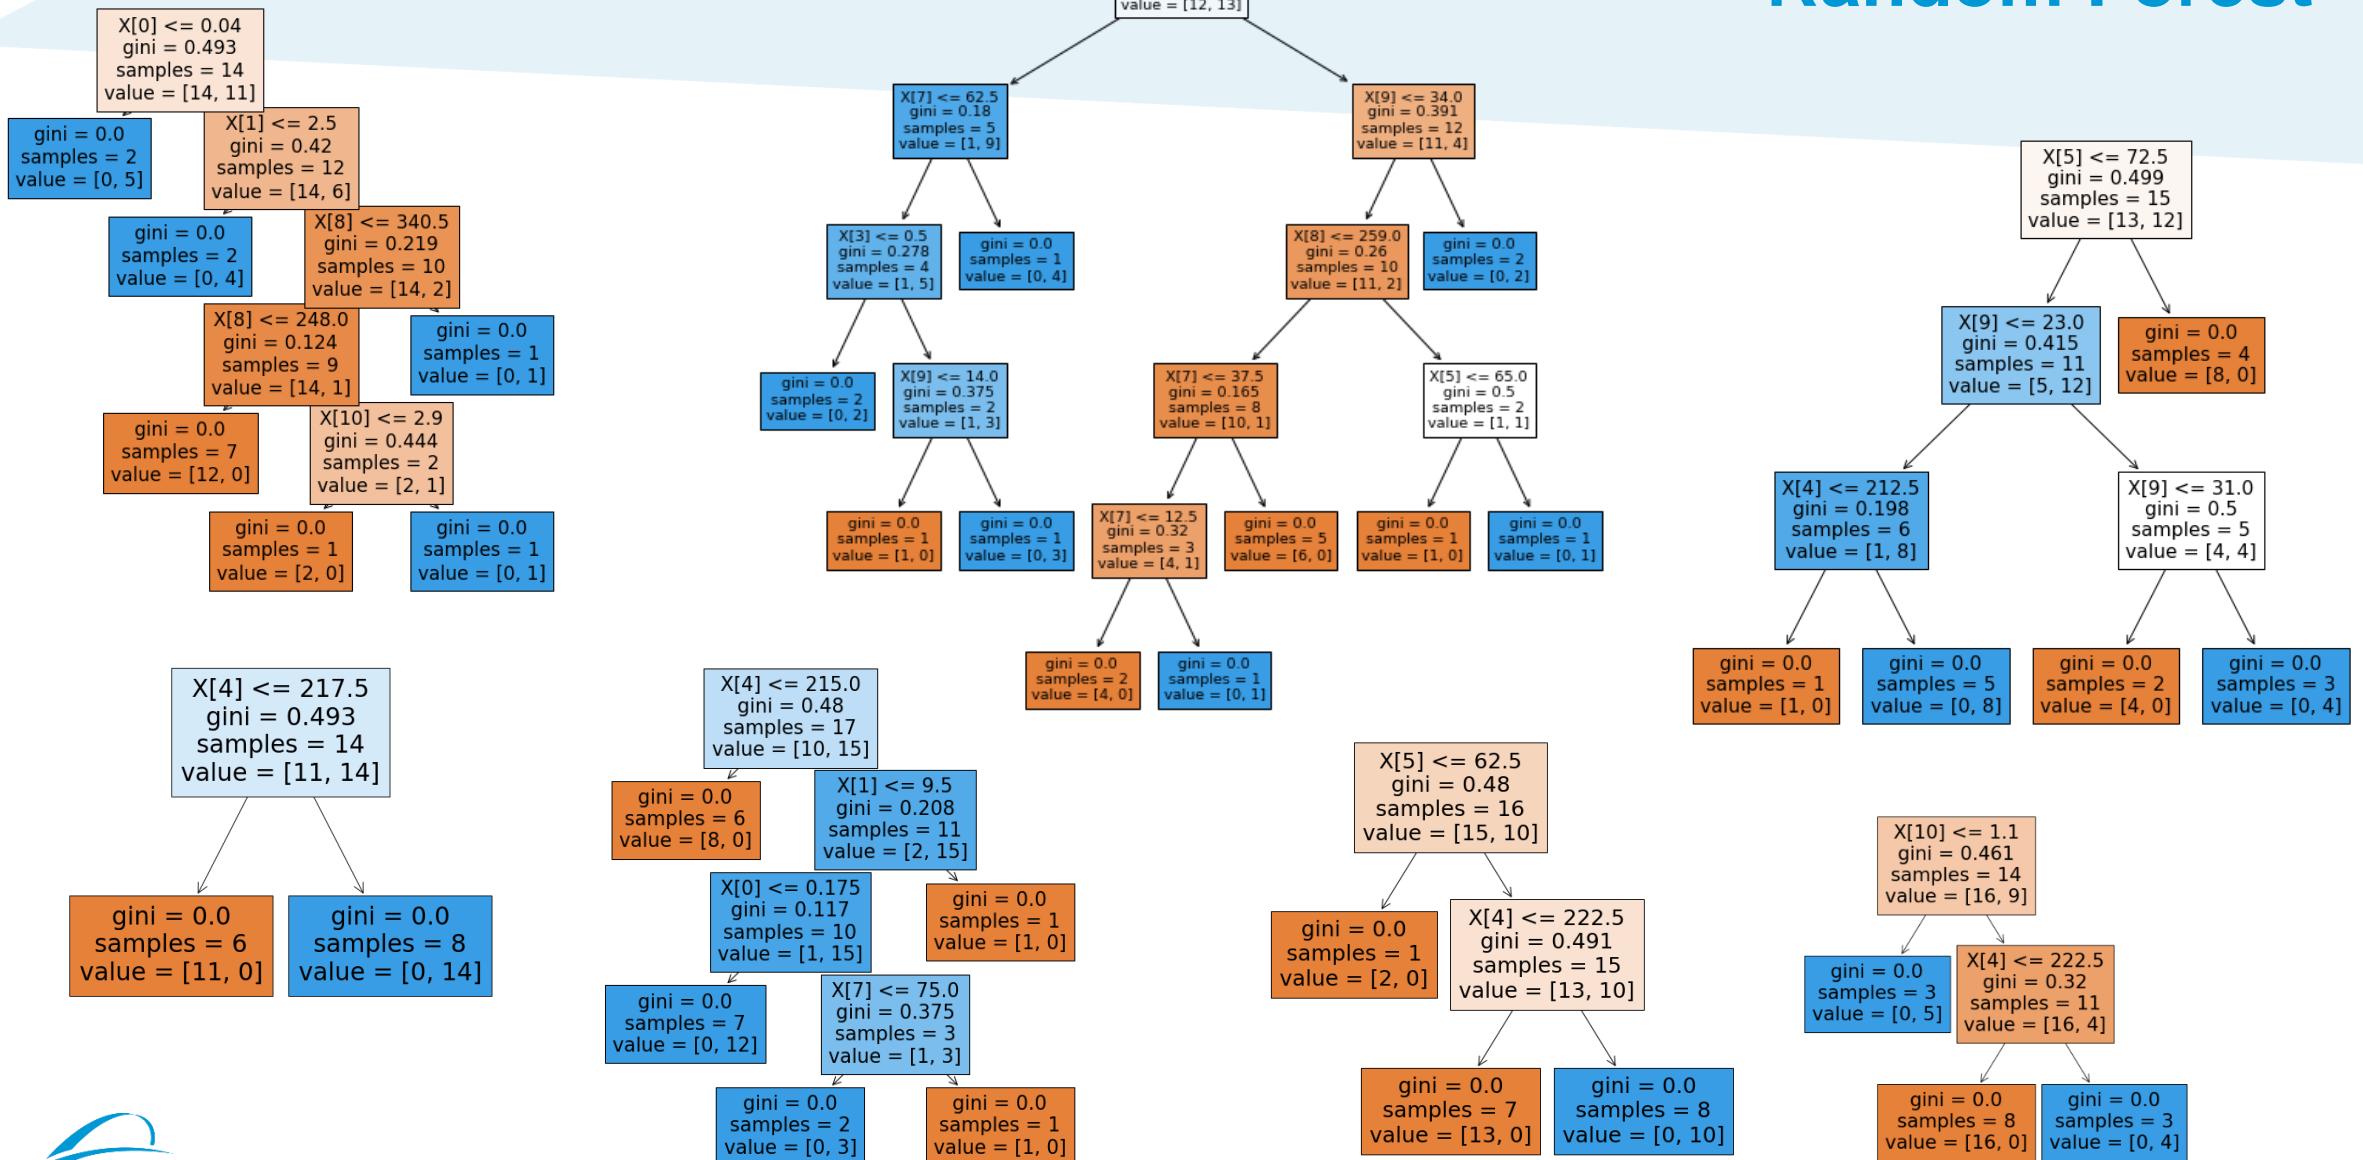
```
import pickle
with open("decision_tree.sav", "rb") as f:
    clf = pickle.load(f)
```

Model 1b:

RANDOM FOREST



Random Forest



Stappenplan

1. Maak een 'bootstrapped' dataset
2. Maak een decision tree van deze dataset (met een paar kolommen)
3. Herhaal 1 en 2
4. Alle decision trees bij elkaar representeren de Random Forest

Stap 1. bootstrapped dataset

	layer_height	wall_thickness	infill_density	infill_pattern	nozzle_temperature	bed_temperature	print_speed	material	fan_speed	roughness	tension_strenght
0	0.02	8	90	0	220	60	40	0	0	25	18
1	0.02	7	90	1	225	65	40	0	25	32	16
2	0.02	1	80	0	230	70	40	0	50	40	8
3	0.02	4	70	1	240	75	40	0	75	68	10
4	0.02	6	90	0	250	80	40	0	100	92	5
.											
.											
.											
45	0.2	5	60	1	200	60	40	1	0	321	28
46	0.2	4	20	0	205	65	40	1	25	265	14
47	0.2	5	60	1	210	70	40	1	50	278	30
48	0.2	7	40	0	215	75	40	1	75	244	29
49	0.2	3	60	1	220	80	40	1	100	220	27



	layer_height	wall_thickness	infill_density	infill_pattern	nozzle_temperature	bed_temperature	print_speed	fan_speed	roughness	tension_strenght	elongatik
48	0.20	7	40	0	215	75	40	75	244	29	3
29	0.10	1	30	1	220	80	120	100	121	14	1
28	0.10	4	90	0	215	75	120	75	138	34	2
1	0.02	7	90	1	225	65	40	25	32	16	1
1	0.02	7	90	1	225	65	40	25	32	16	1
22	0.10	1	50	0	230	70	120	50	265	10	0
45	0.20	5	60	1	200	60	40	0	321	28	2
36	0.15	4	50	0	205	65	60	25	212	24	1
44	0.20	3	80	0	250	80	40	100	357	21	1
20	0.10	1	40	0	220	60	120	0	120	16	1

bootstrapped

Stap 2. decision tree maken

Bij elke stap selecteren we X random kolommen waarmee we een decision tree gaan vullen

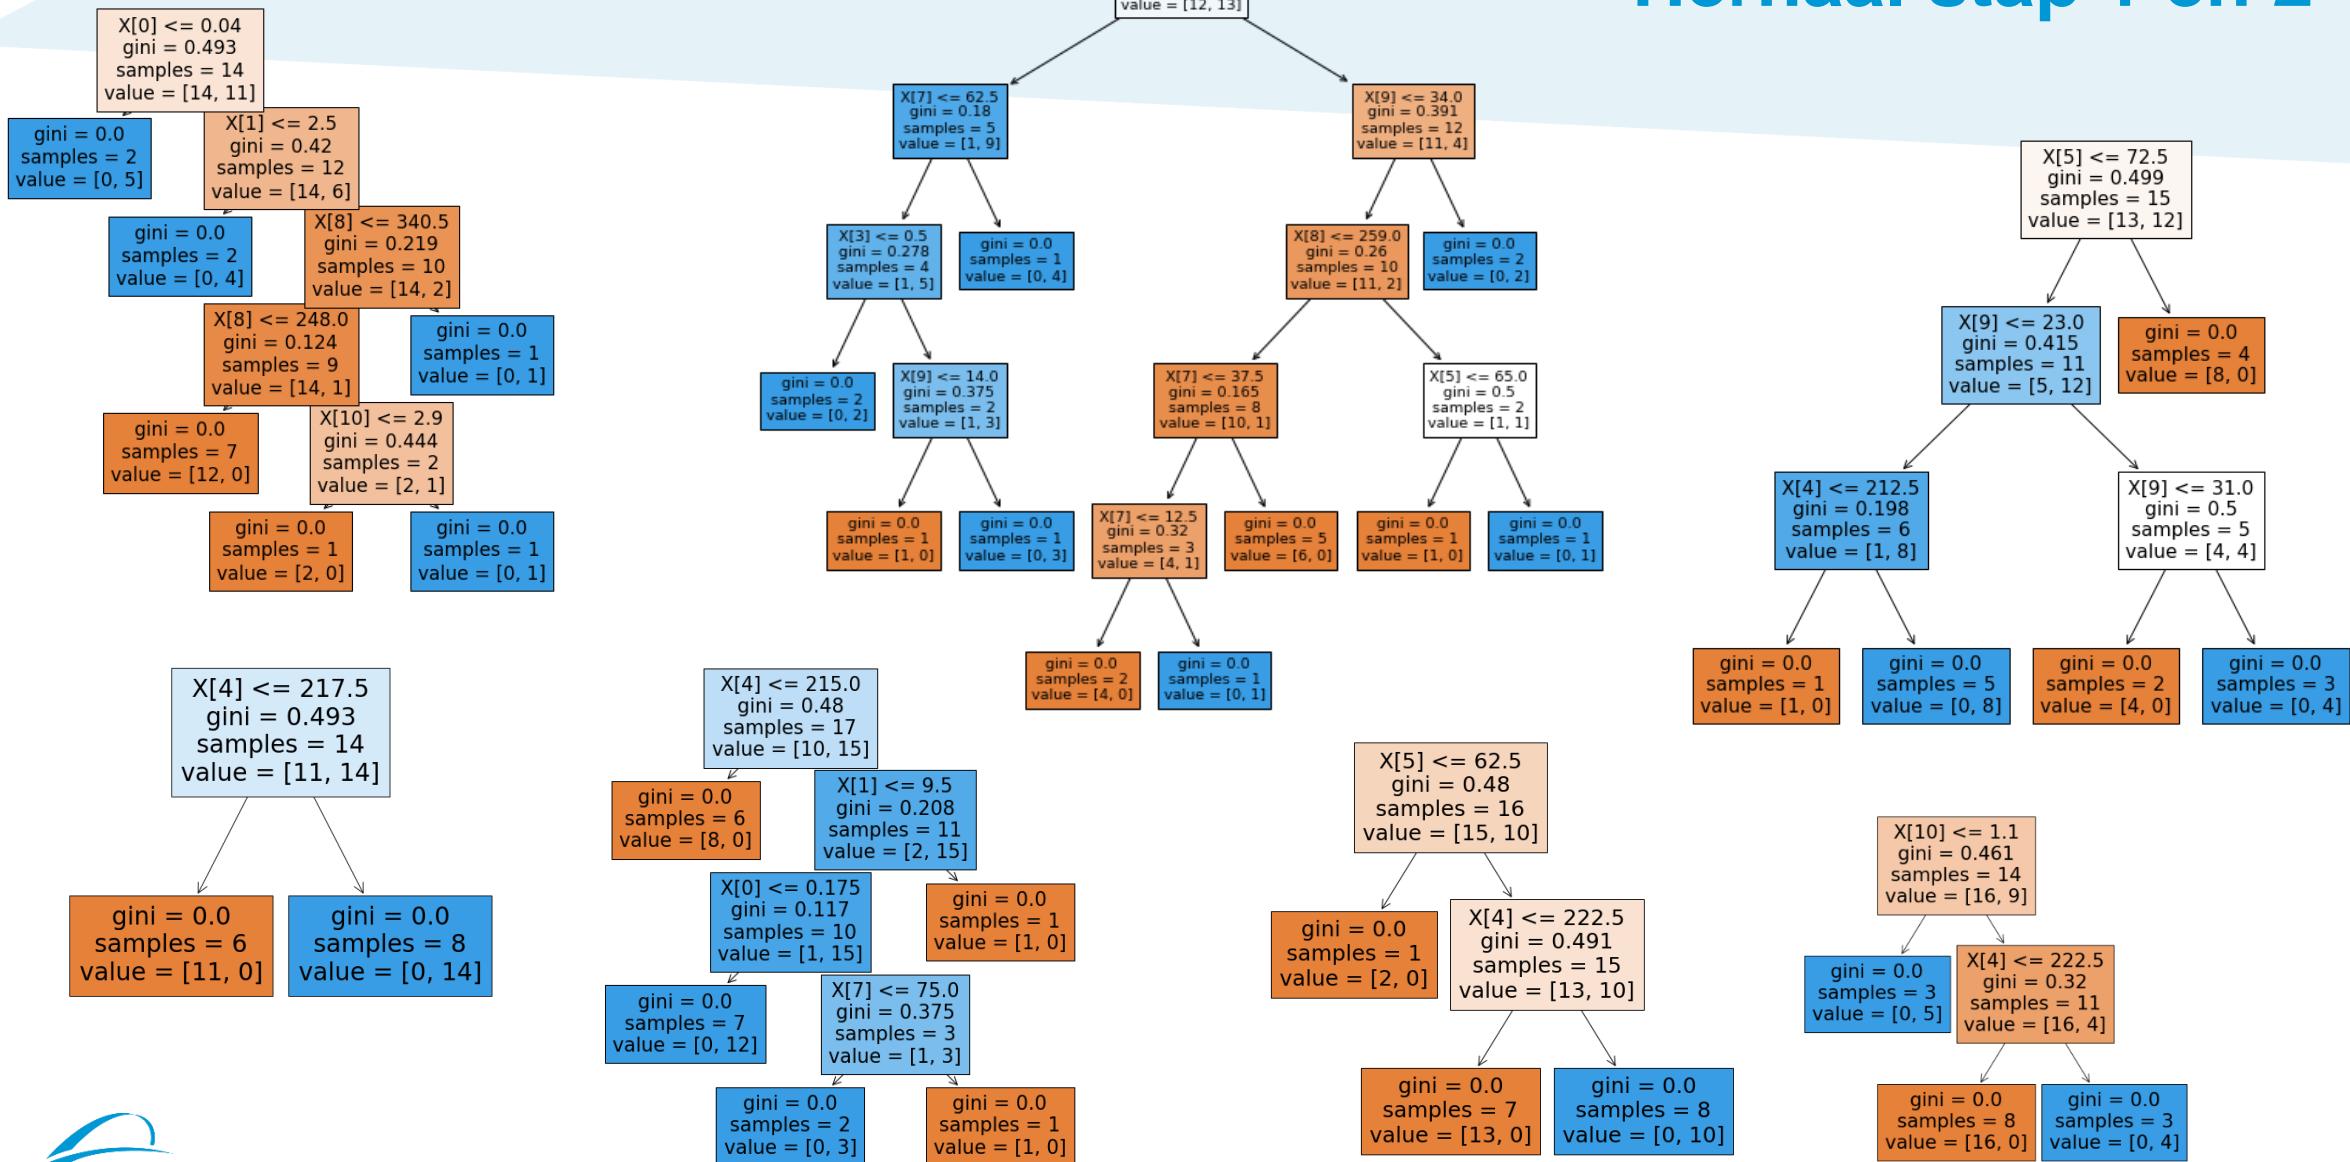
Stel, we selecteren elke keer 2 kolommen

	layer_height	wall_thickness	infill_density	infill_pattern	nozzle_temperature	bed_temperature	print_speed	fan_speed	roughness	tension_strenght	elongation
48	0.20	7	40	0	215	75	40	75	244	29	3
29	0.10	1	30	1	220	80	120	100	121	14	1
28	0.10	4	90	0	215	75	120	75	138	34	2
1	0.02	7	90	1	225	65	40	25	32	16	1
1	0.02	7	90	1	225	65	40	25	32	16	1
22	0.10	1	50	0	230	70	120	50	265	10	0
45	0.20	5	60	1	200	60	40	0	321	28	2
36	0.15	4	50	0	205	65	60	25	212	24	1
44	0.20	3	80	0	250	80	40	100	357	21	1
20	0.10	1	40	0	220	60	120	0	120	16	1

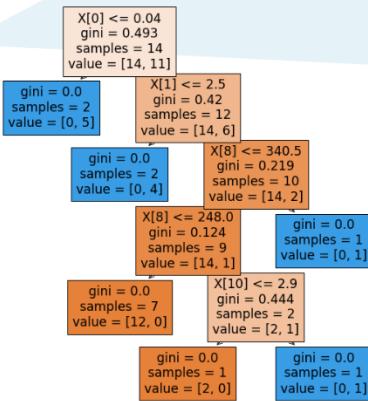
X[0] <= 0.04
gini = 0.493
samples = 14
value = [14, 11]

We zorgen hiermee voor grote variatie tussen alle decision trees. Dit maakt de classifier robuster.

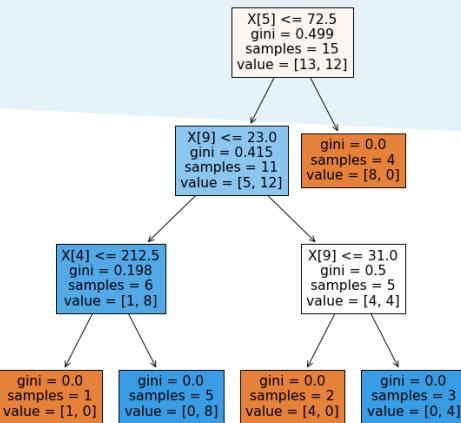
Herhaal stap 1 en 2



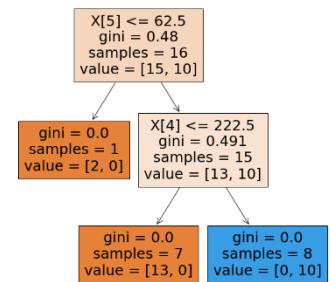
Classificeren



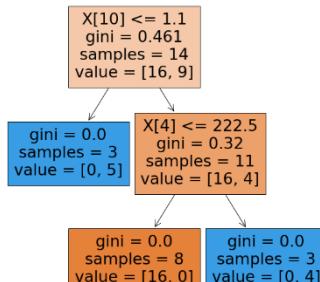
$$\mathbf{P} = [0.82 \quad 0.18]$$



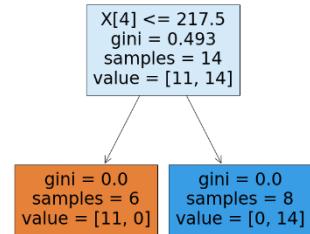
$$\mathbf{P} = [0.99 \quad 0.01]$$



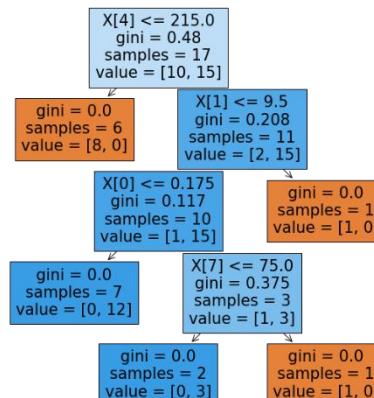
$$\mathbf{P} = [0.72 \quad 0.28]$$



$$\mathbf{P} = [0.84 \quad 0.16]$$



$$\mathbf{P} = [0.65 \quad 0.35]$$



$$\mathbf{P} = [0.93 \quad 0.07]$$

Aggereren van de uitkomsten

$$\begin{aligned} \mathbf{P} &= [0.82 \quad 0.18] \\ \mathbf{P} &= [0.72 \quad 0.28] \\ \mathbf{P} &= [0.65 \quad 0.35] \\ \mathbf{P} &= [0.99 \quad 0.01] \\ \mathbf{P} &= [0.84 \quad 0.16] \\ \mathbf{P} &= [0.93 \quad 0.07] \end{aligned}$$

Aggereren:

- Gemiddelde
- Meeste 'stemmen'

$$\mathbf{P} = [0.825 \quad 0.175]$$

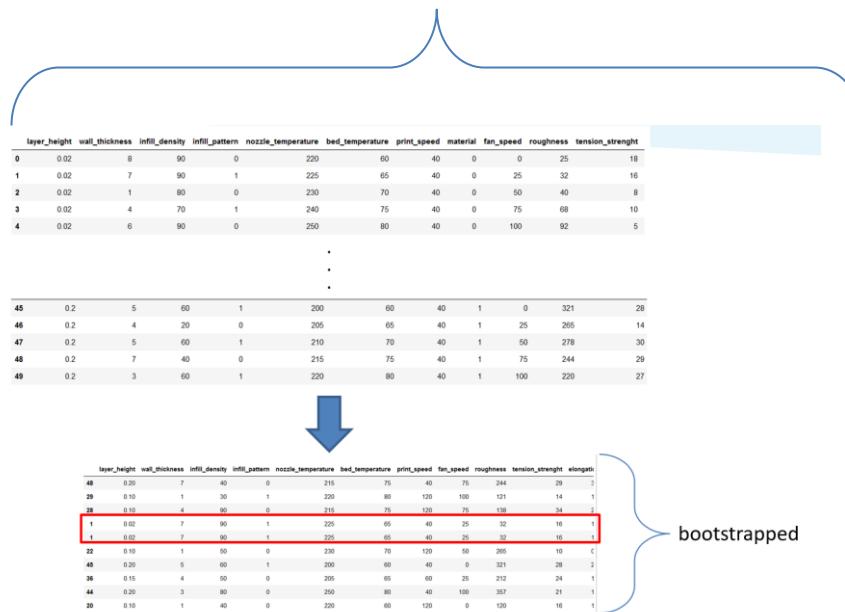
Wordt ook wel eens Bagging trees genoemd

Bootstrapping

+

Aggeregren

= Bagging



$$\begin{aligned} \mathbf{P} &= [0.82 \quad 0.18] \\ \mathbf{P} &= [0.72 \quad 0.28] \\ \mathbf{P} &= [0.65 \quad 0.35] \\ \mathbf{P} &= [0.99 \quad 0.01] \\ \mathbf{P} &= [0.84 \quad 0.16] \\ \mathbf{P} &= [0.93 \quad 0.07] \end{aligned} \quad \left. \begin{array}{l} \\ \\ \\ \\ \\ \end{array} \right\} \mathbf{P} = [0.825 \quad 0.175]$$

Voordelen

- Makkelijk/Simpel
- Goede resultaten met relatief weinig data
- Betrouwbaar/Robuust
- Minder gevoelig voor overfitting (kan nog steeds gebeuren)
- Goed met categorische en numerieke data

Nadelen

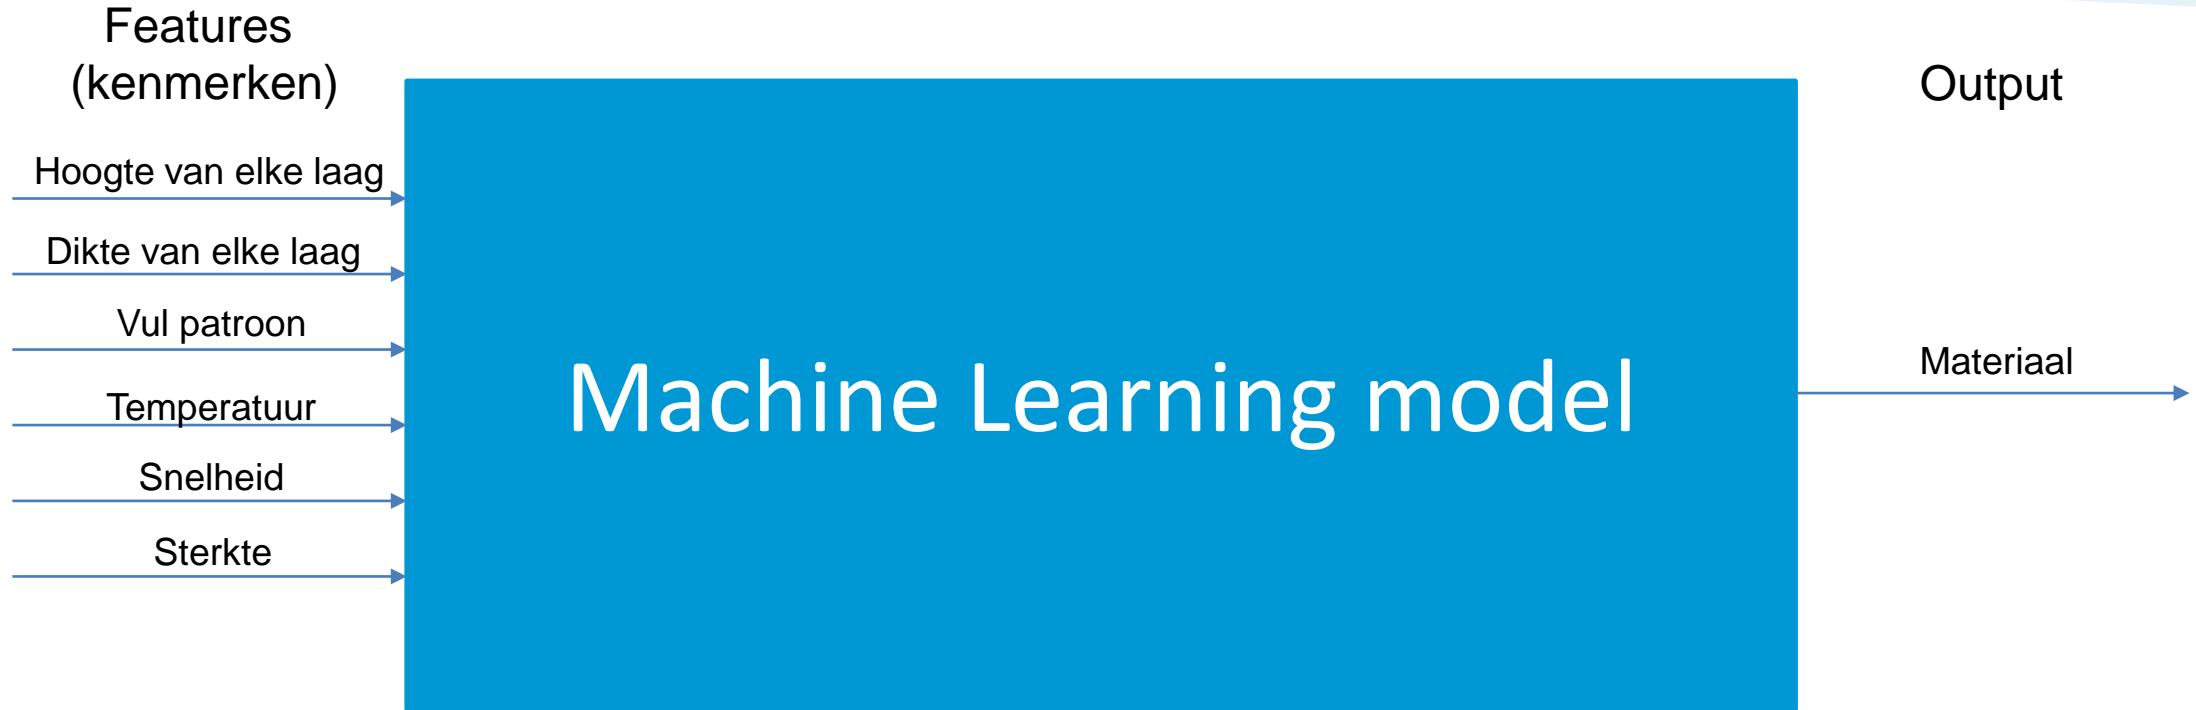
- Langzamer dan DT's
- *Black box*
- Optimale oplossing niet altijd gegarandeerd

Maak een *Random Forest* classifier die voorspelt met welk materiaal is geprint

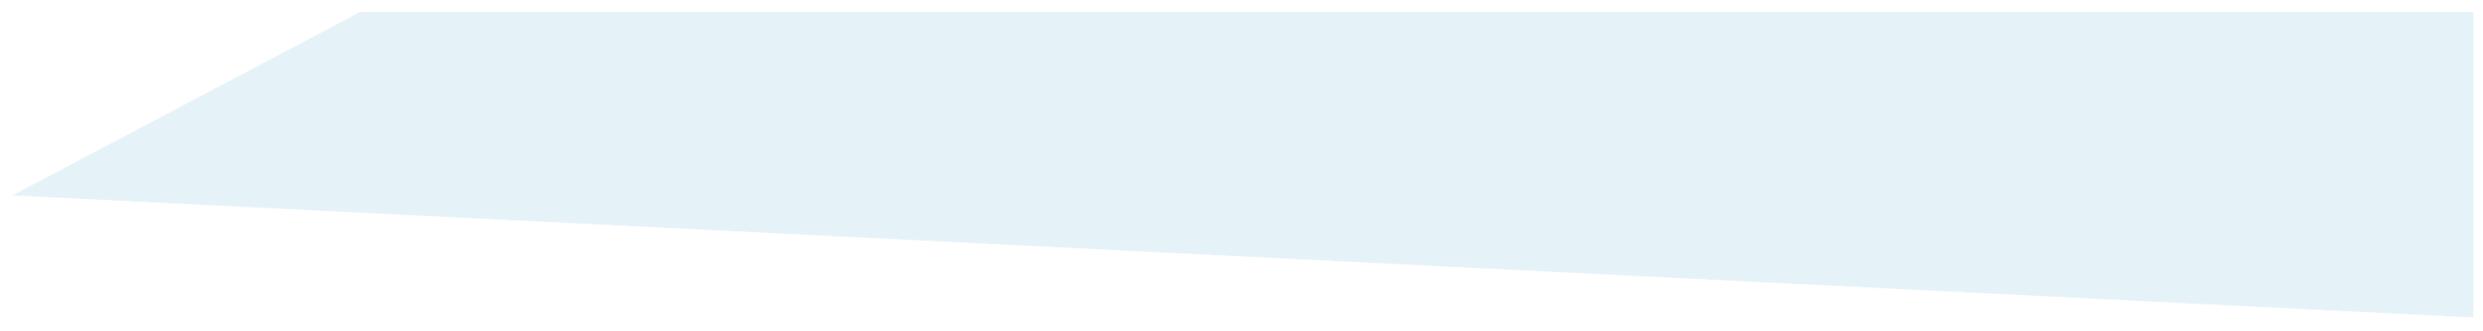
```
from sklearn.ensemble import RandomForestClassifier
```

Welke werkt beter? Random Forest of Decision Tree?

~ 5 minuten

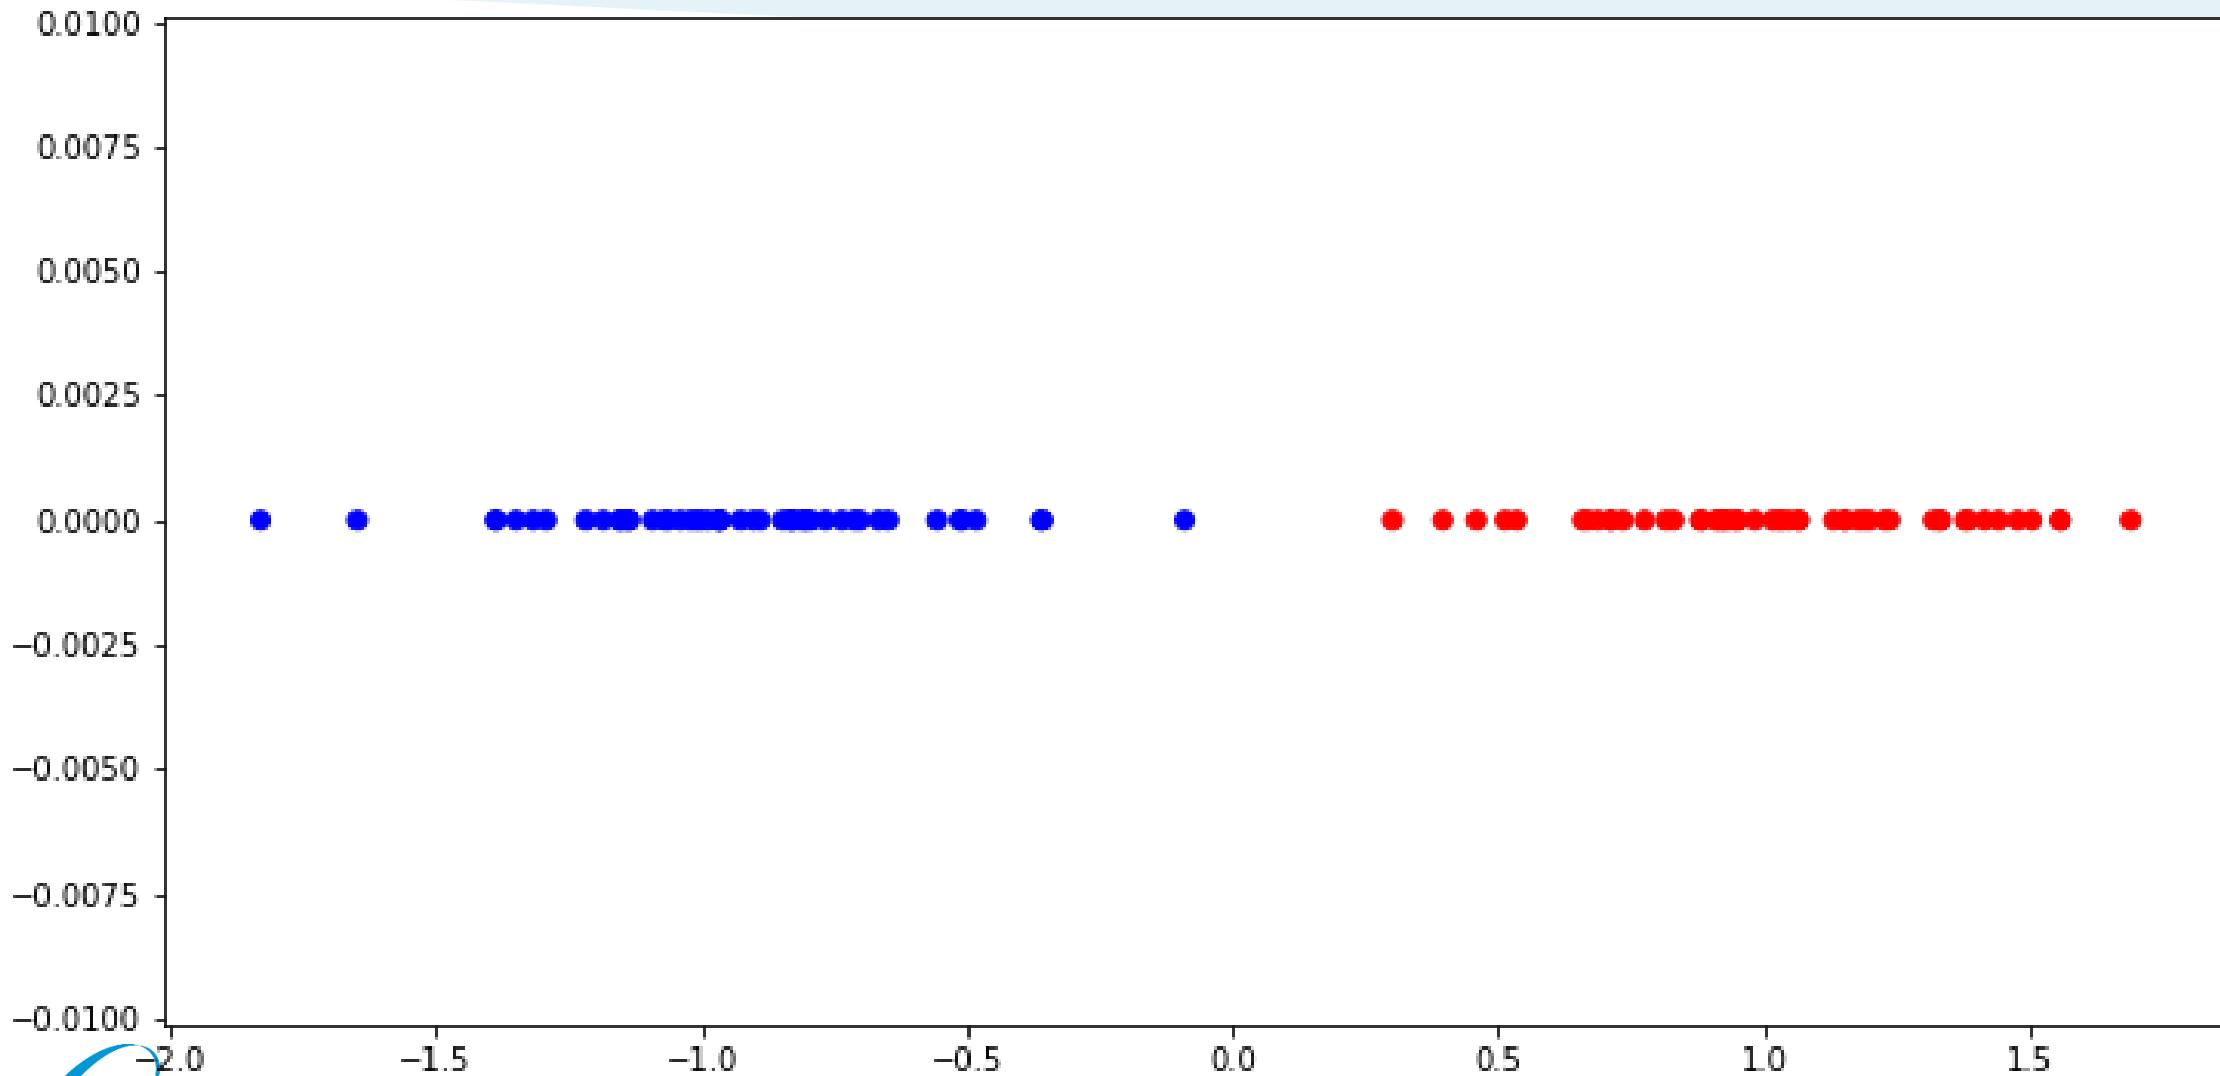


$$Y = f(X)$$

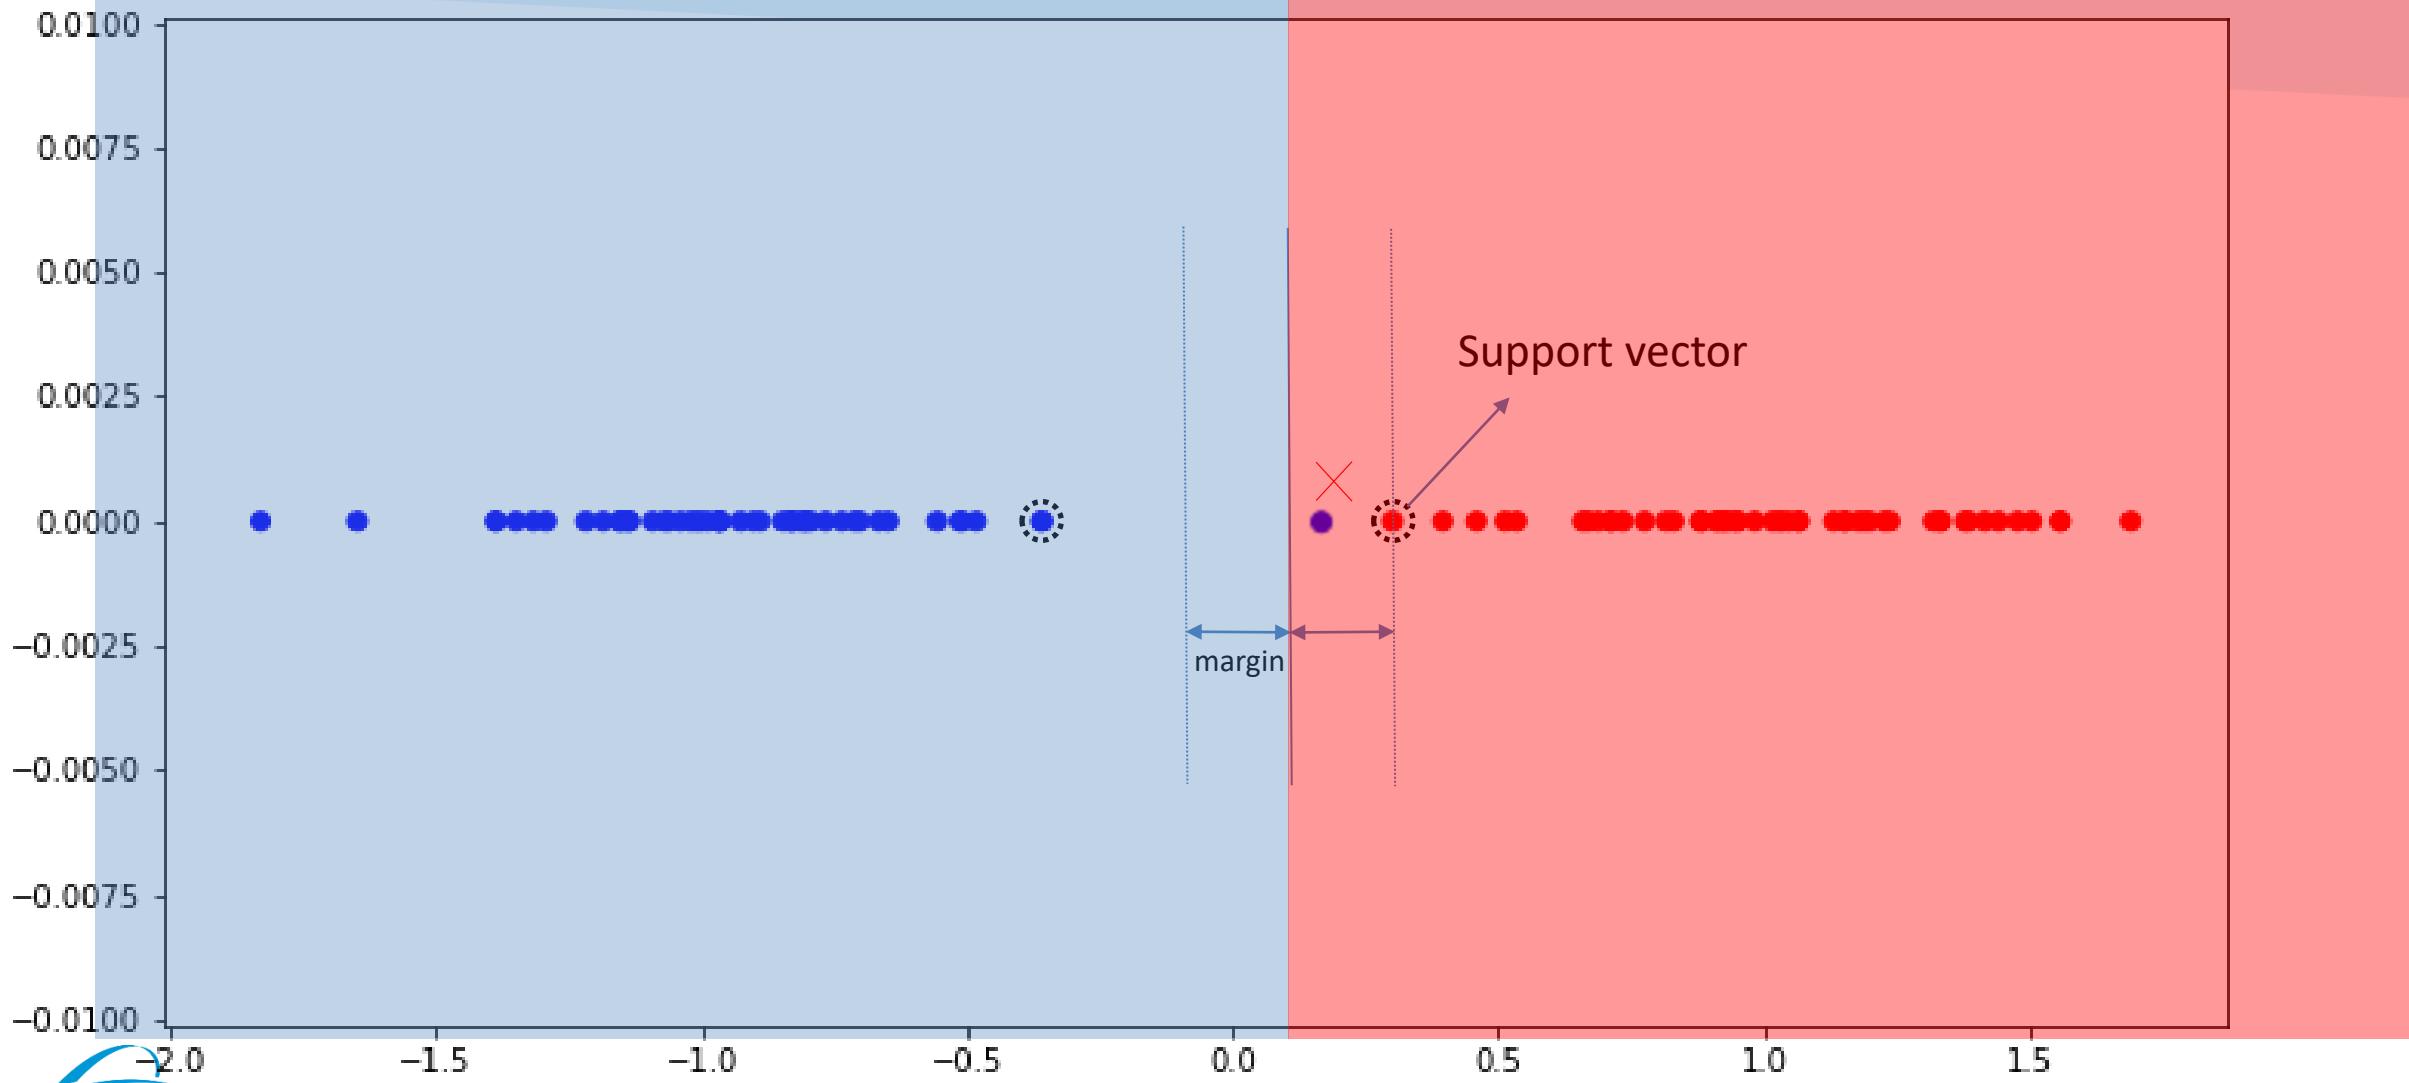


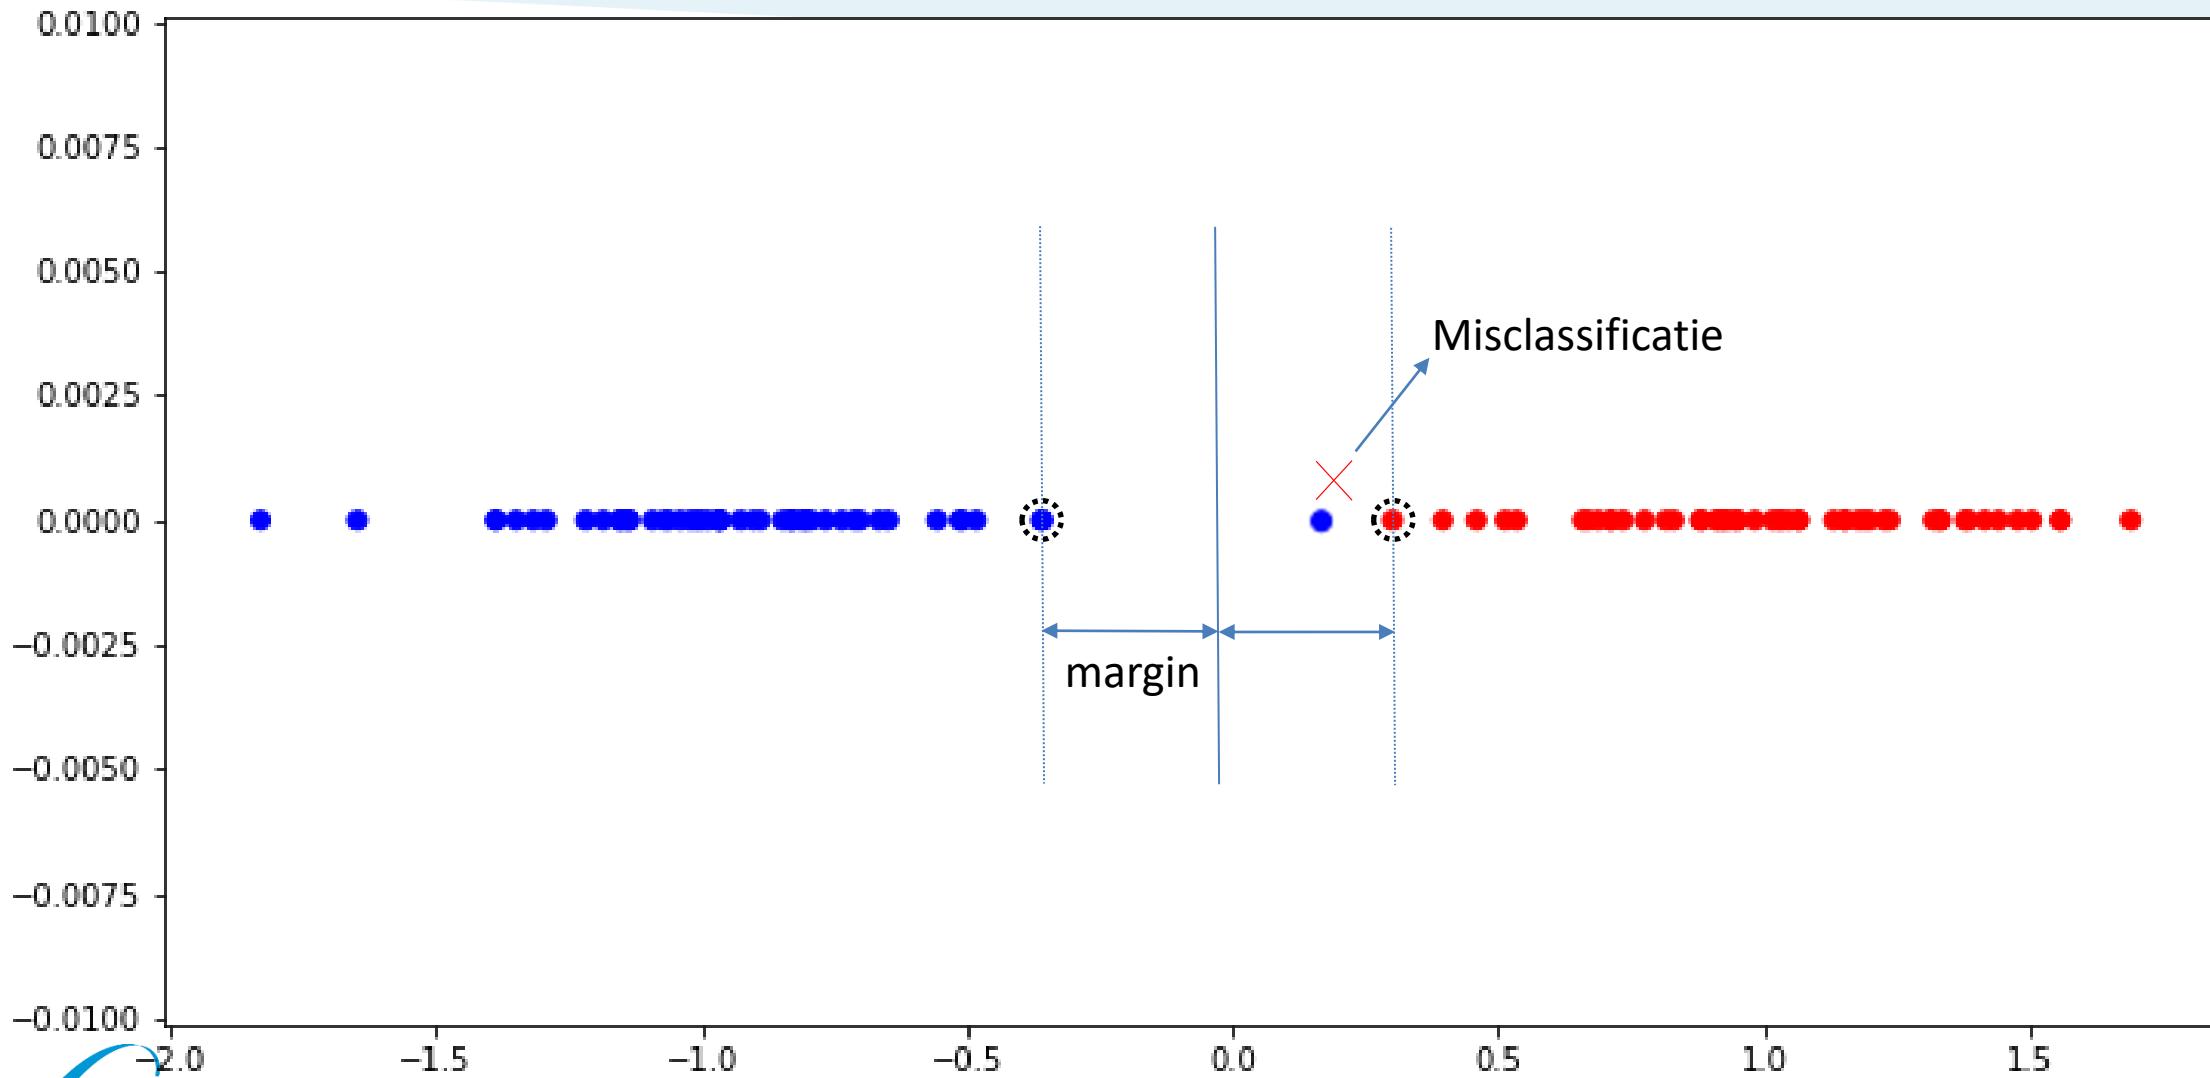
Model 2:

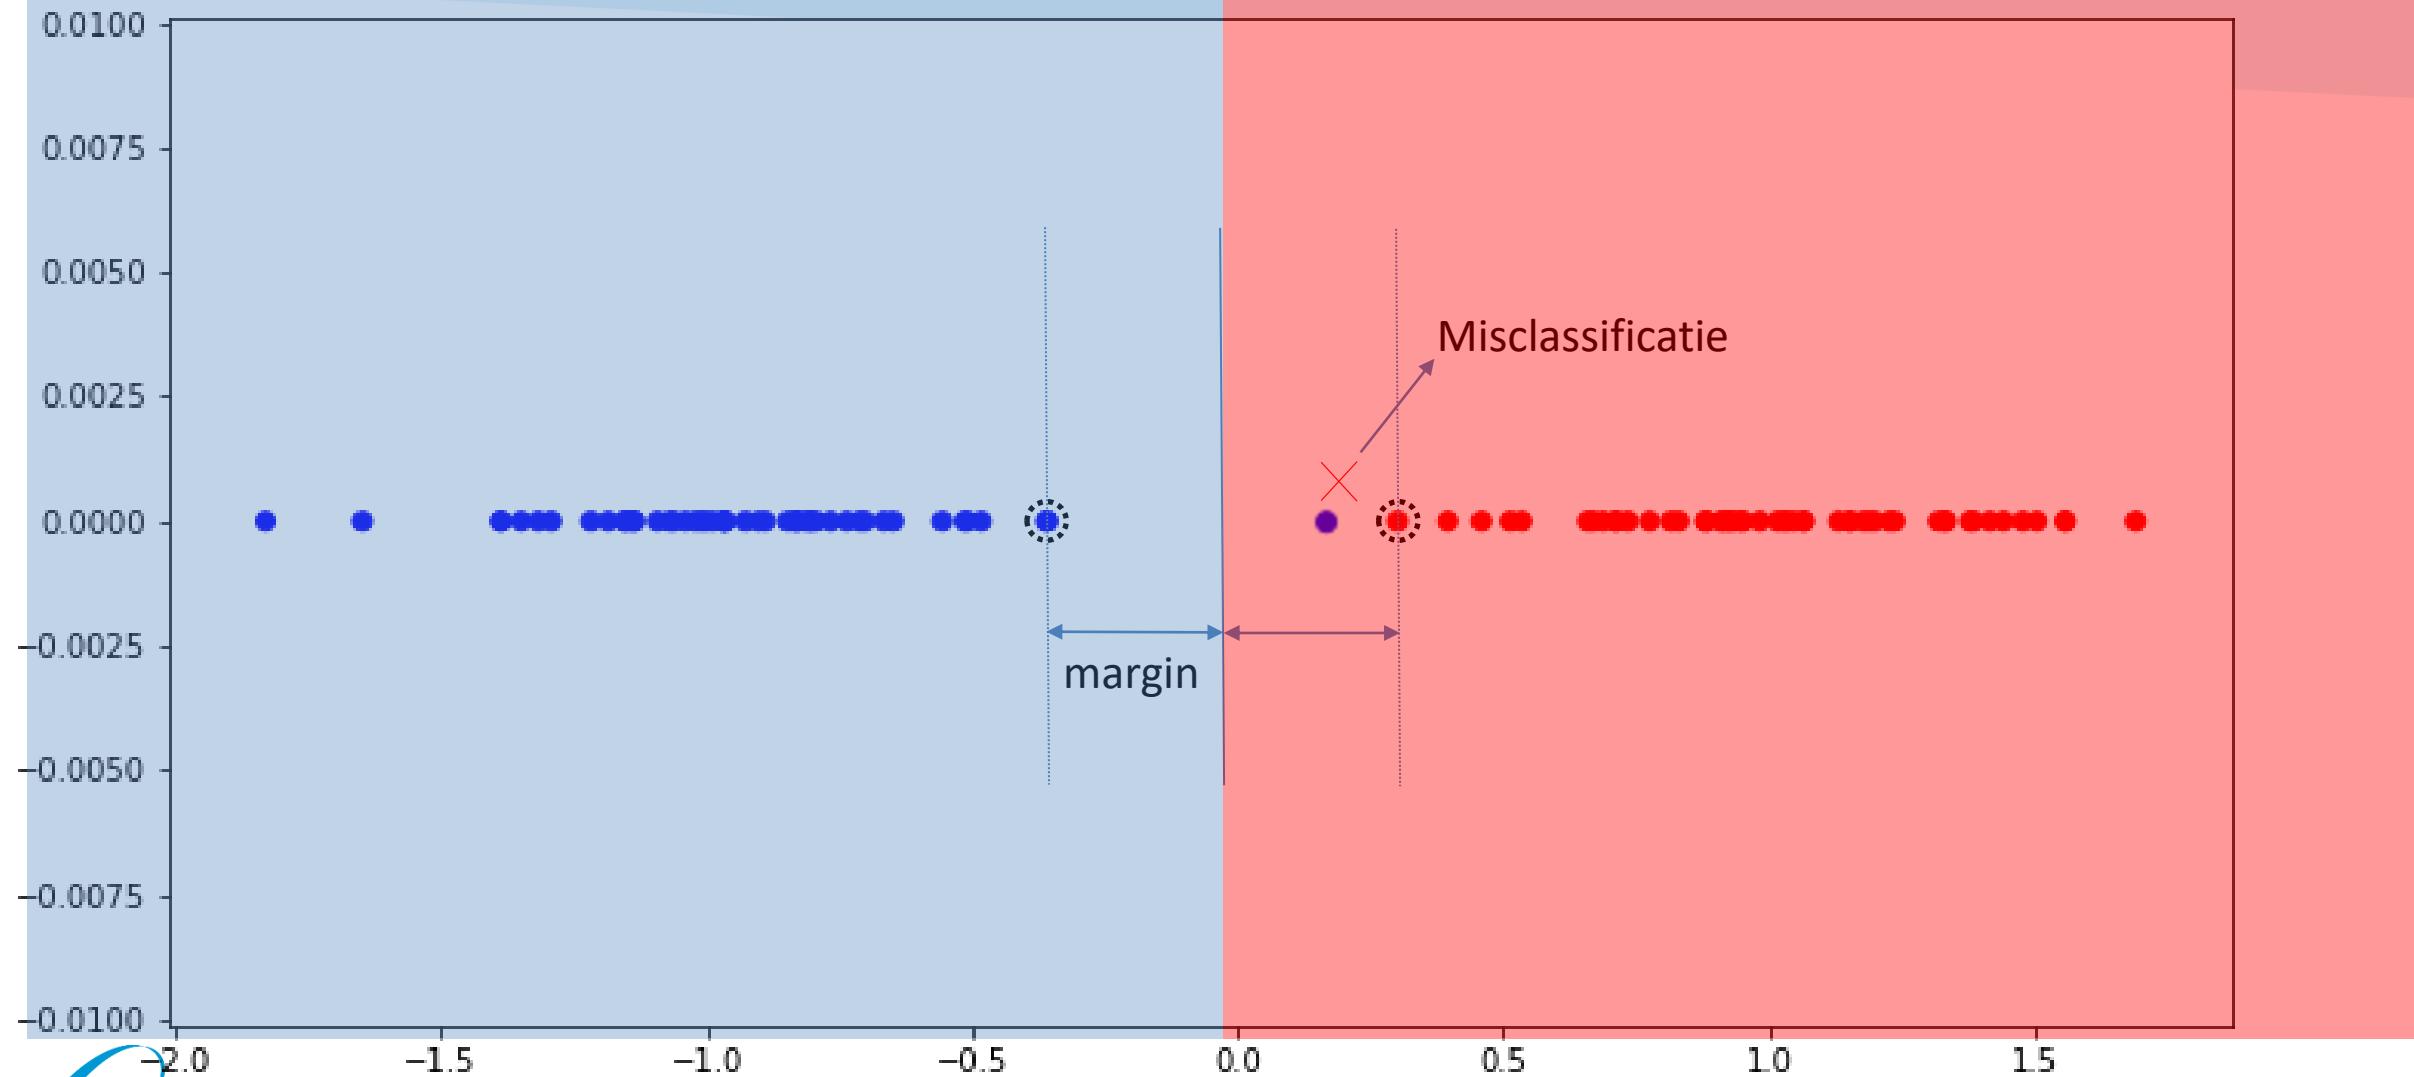
SUPPORT VECTOR MACHINES

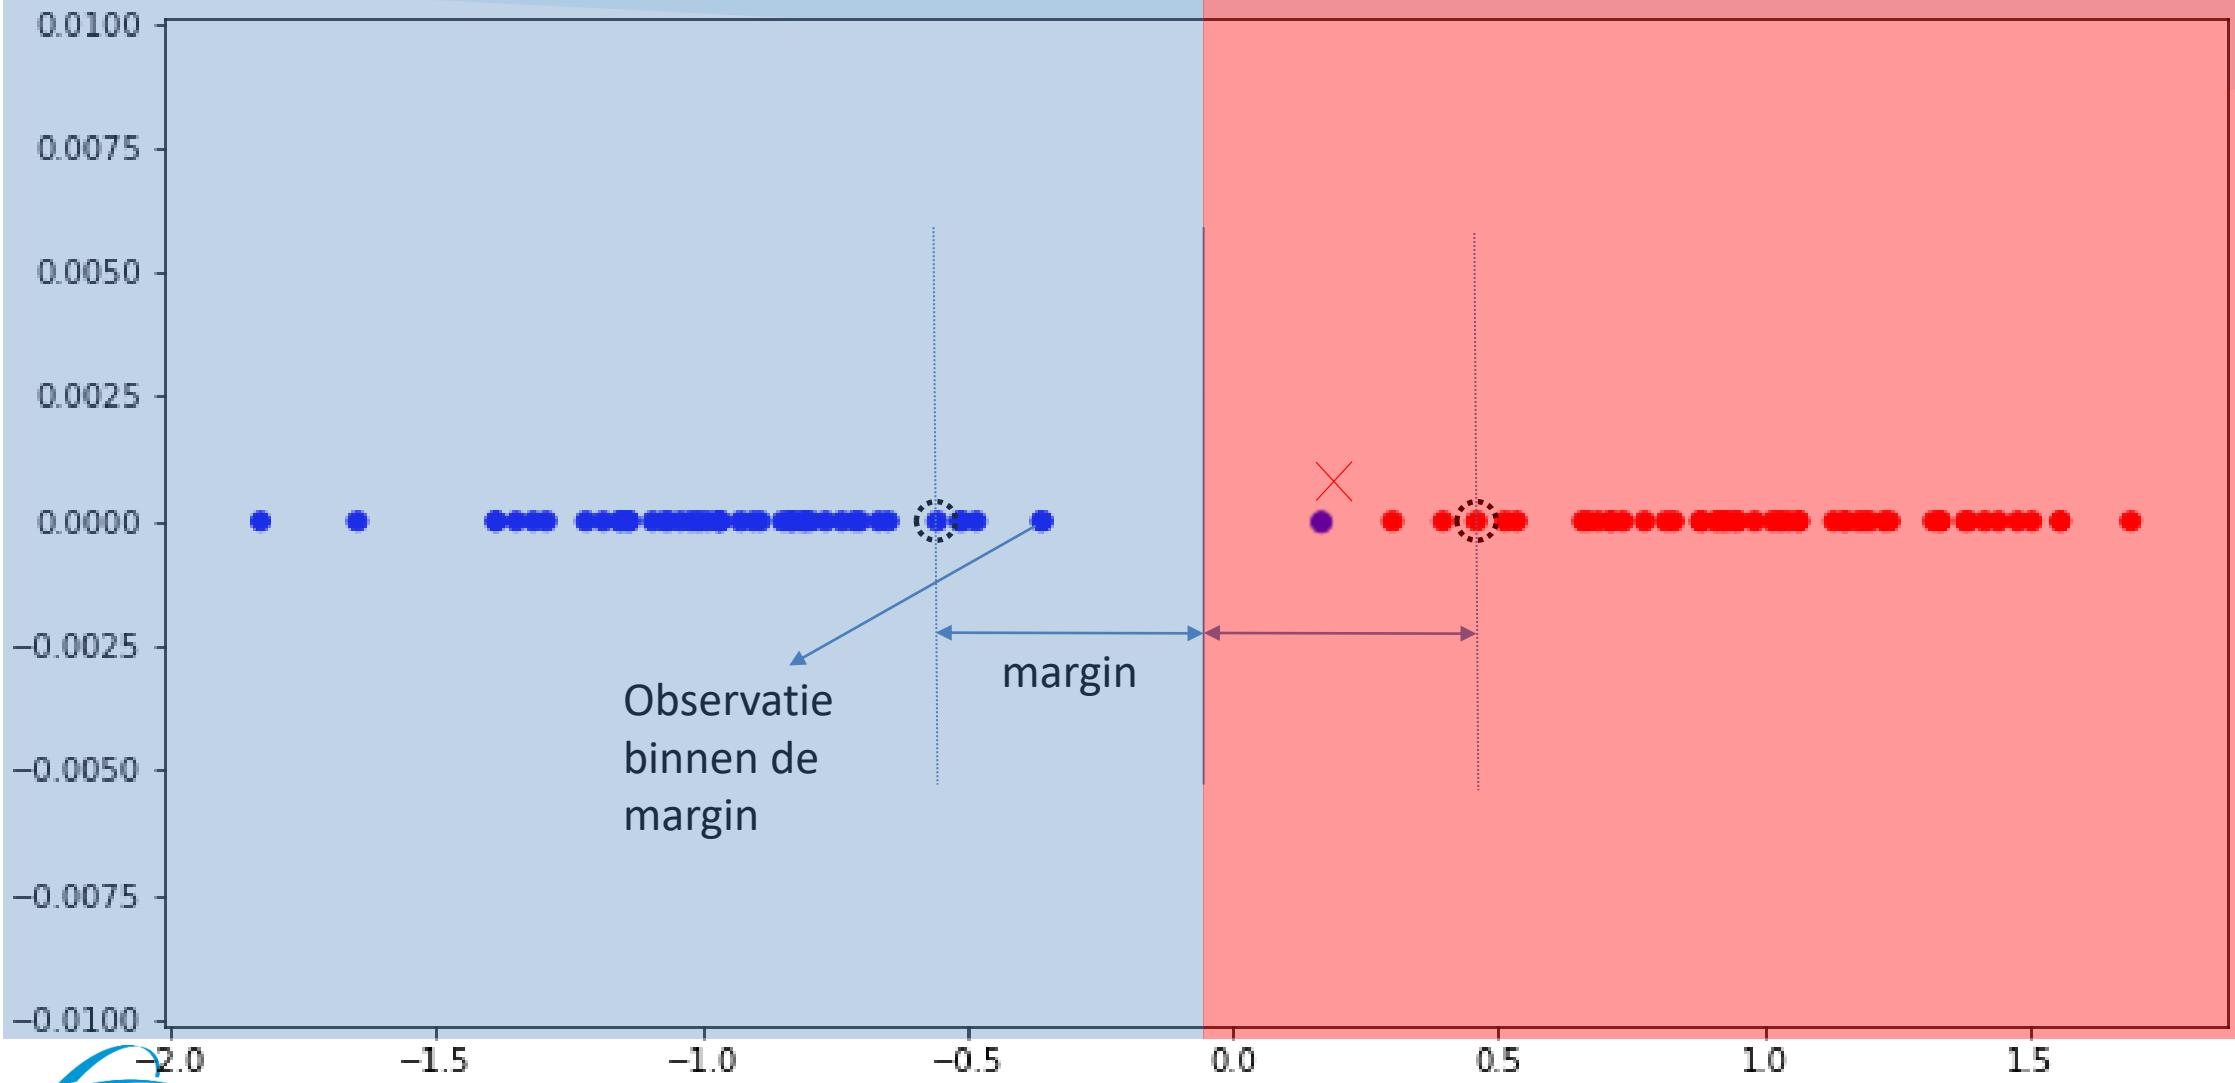


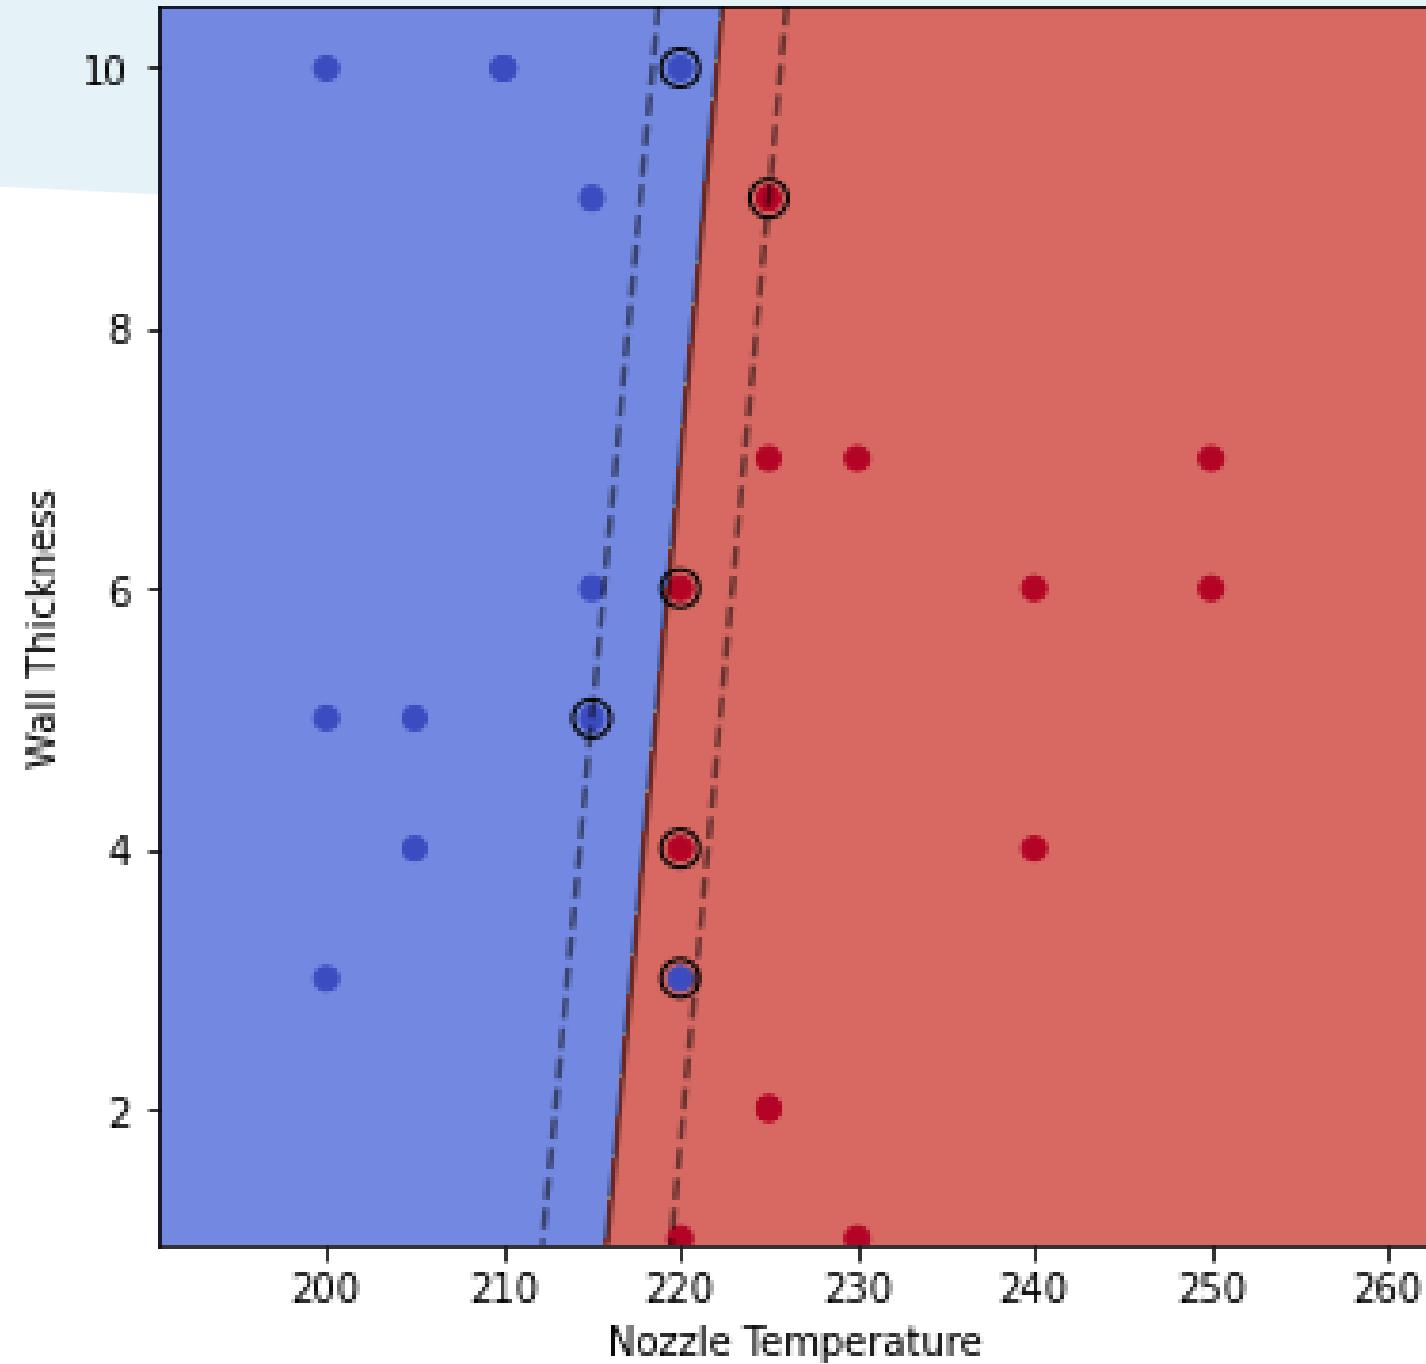
SVM



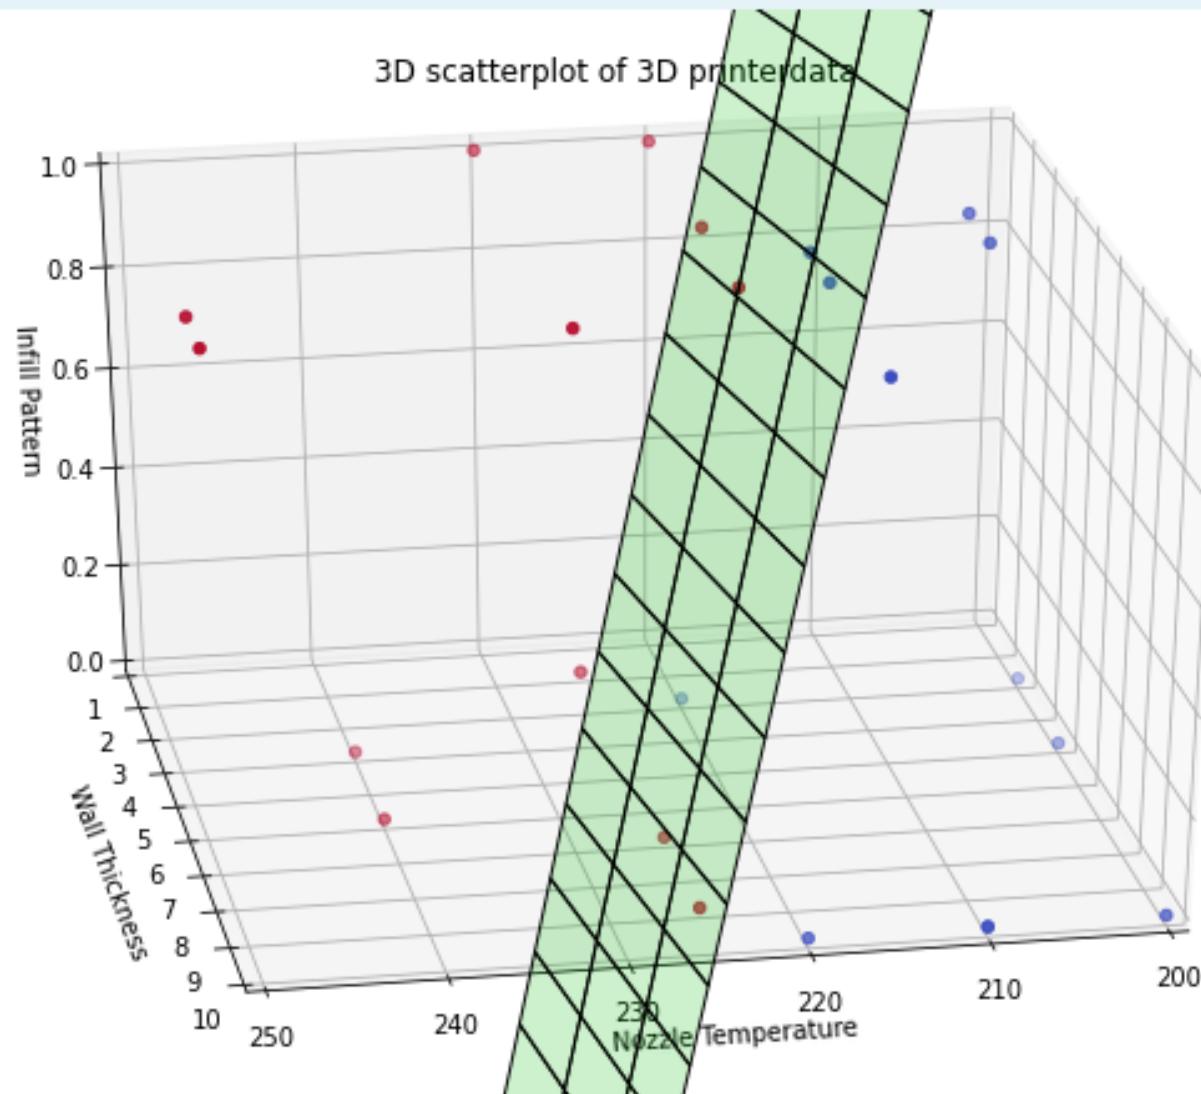




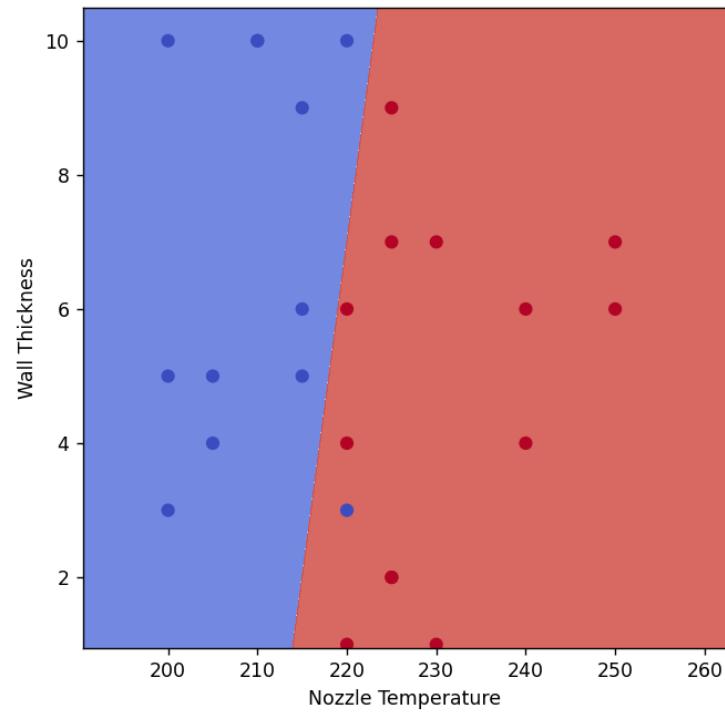




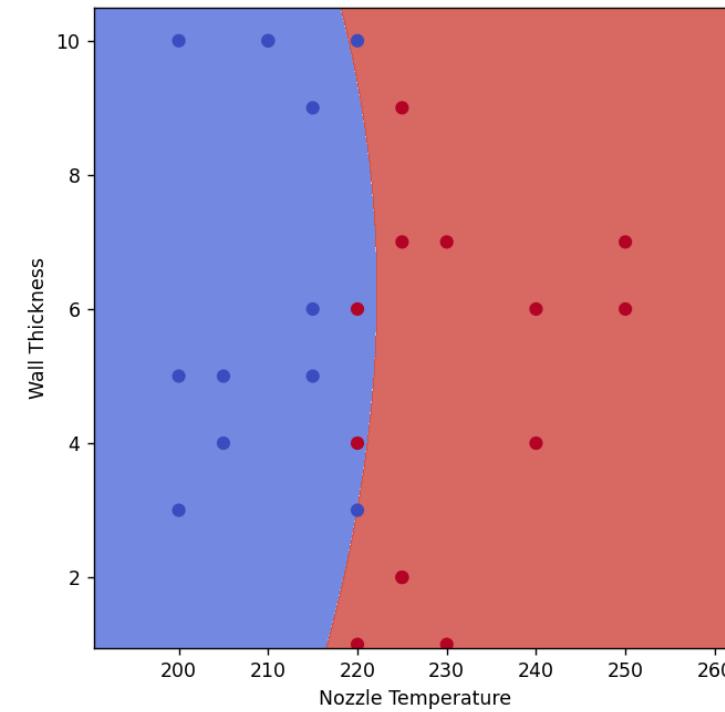
Scatterplot



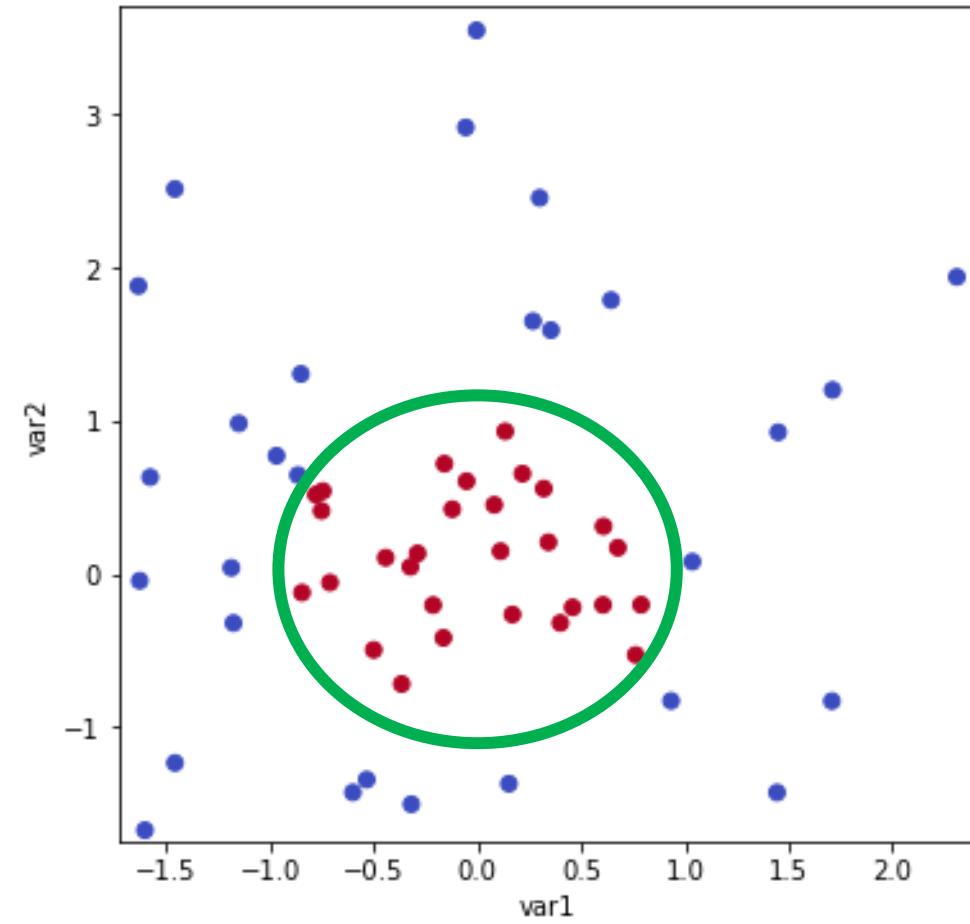
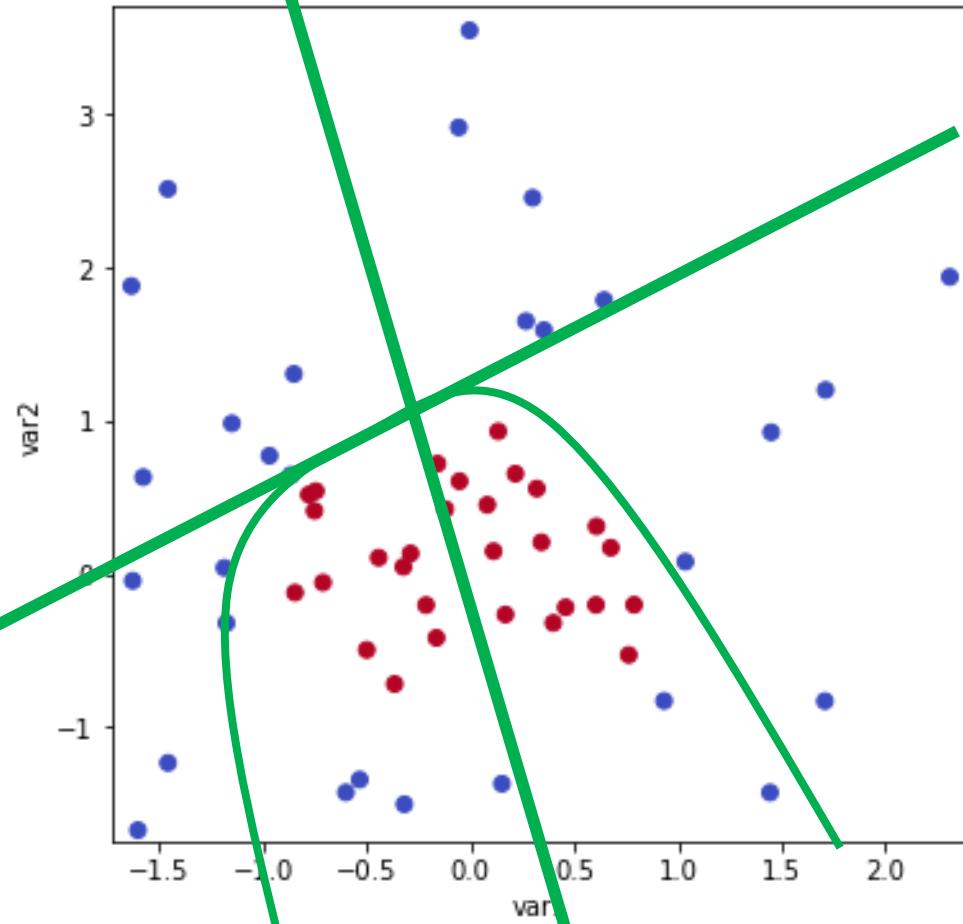
Lineair



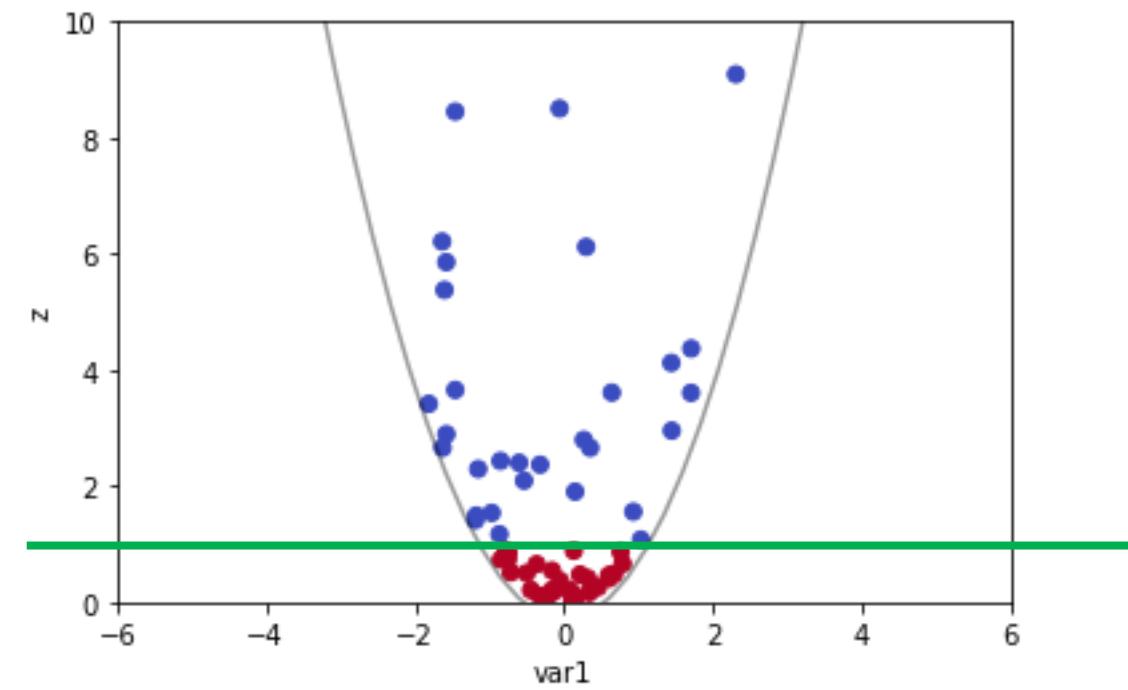
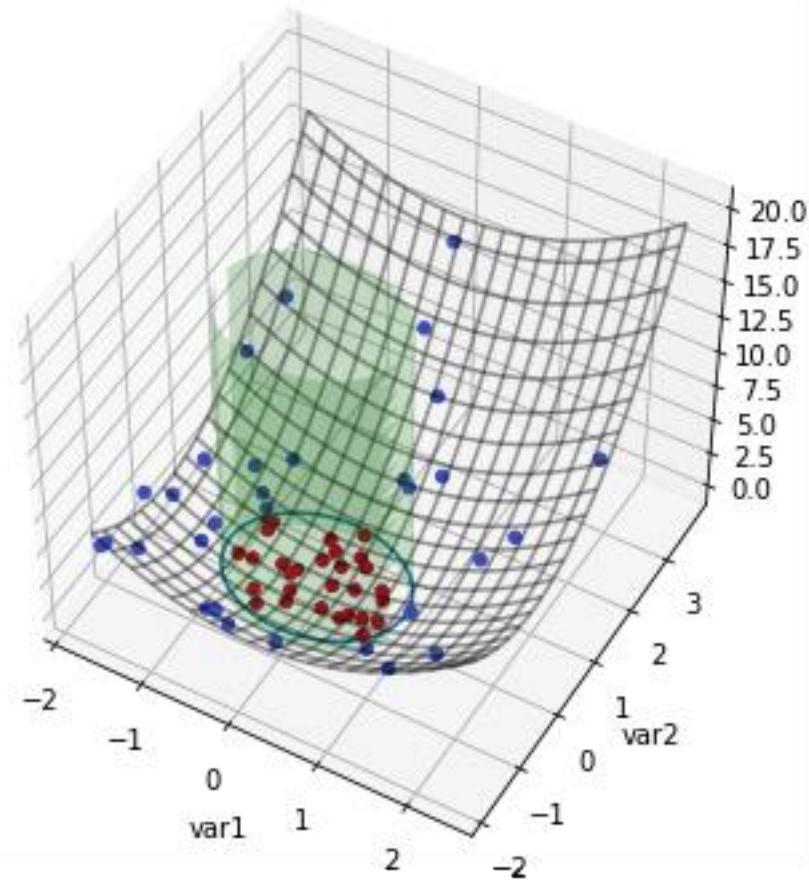
Polynoom

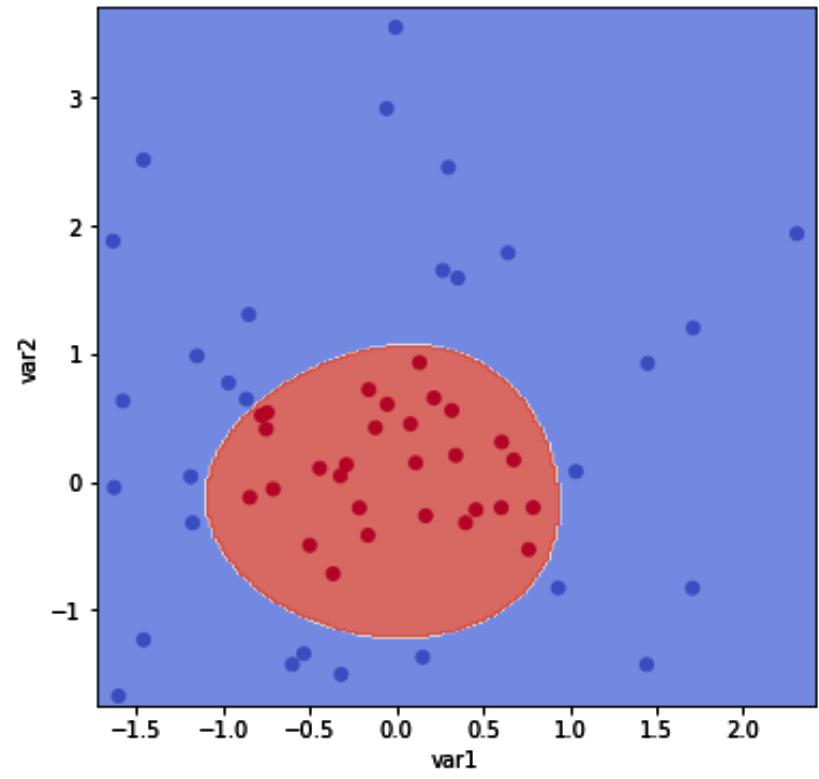
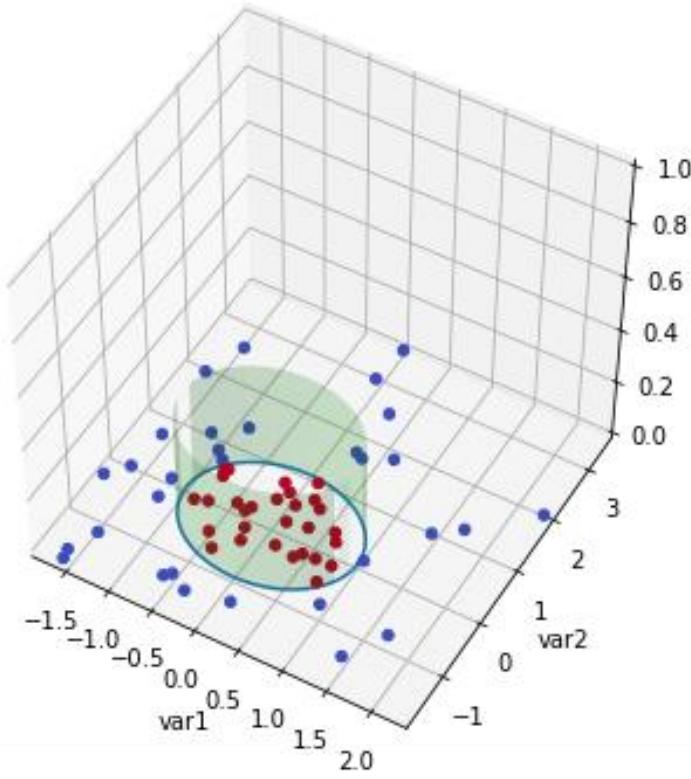
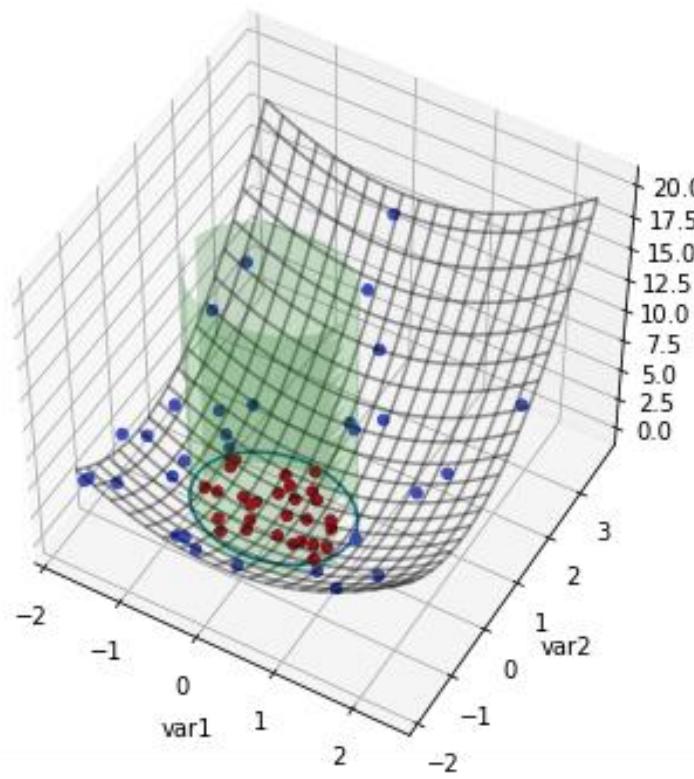


Radial Basis Function



Kernel trick





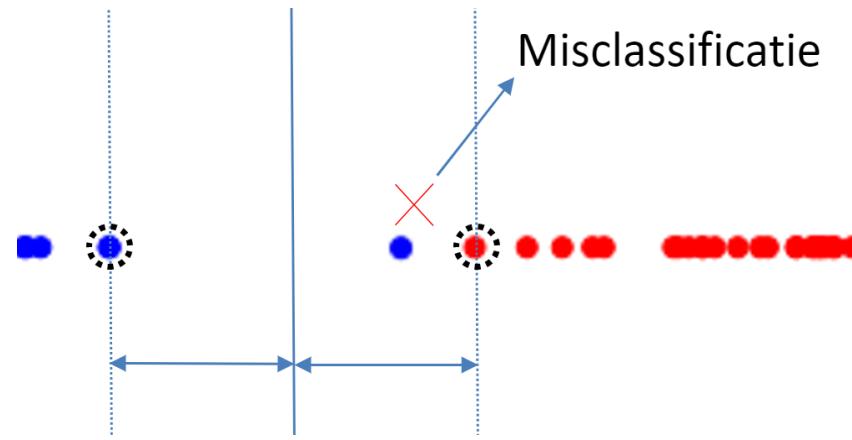
SVM optimization

- Kernel trick ook wiskundig handig
 - Makkelijk te optimaliseren
 - Geen informatie over de data nodig



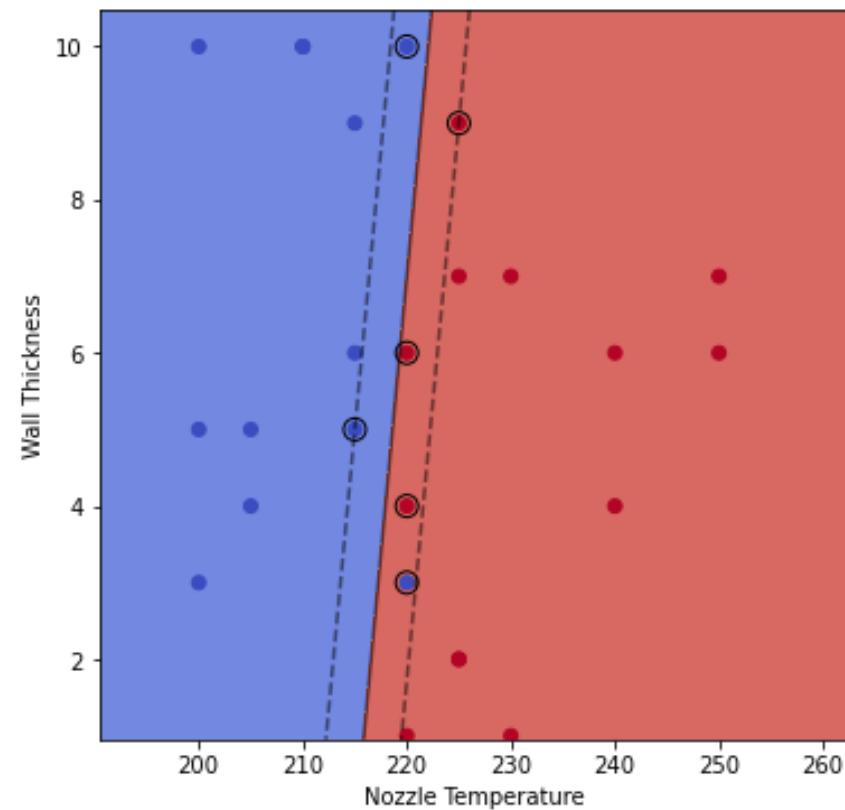
Hyperparameters

- Function/Kernel
['linear', 'poly', 'rbf']
- C → penalty/*regularization* voor misclassificatie
- Gamma γ (alleen voor 'poly' & 'rbf') → De invloed van 1 datapunt.

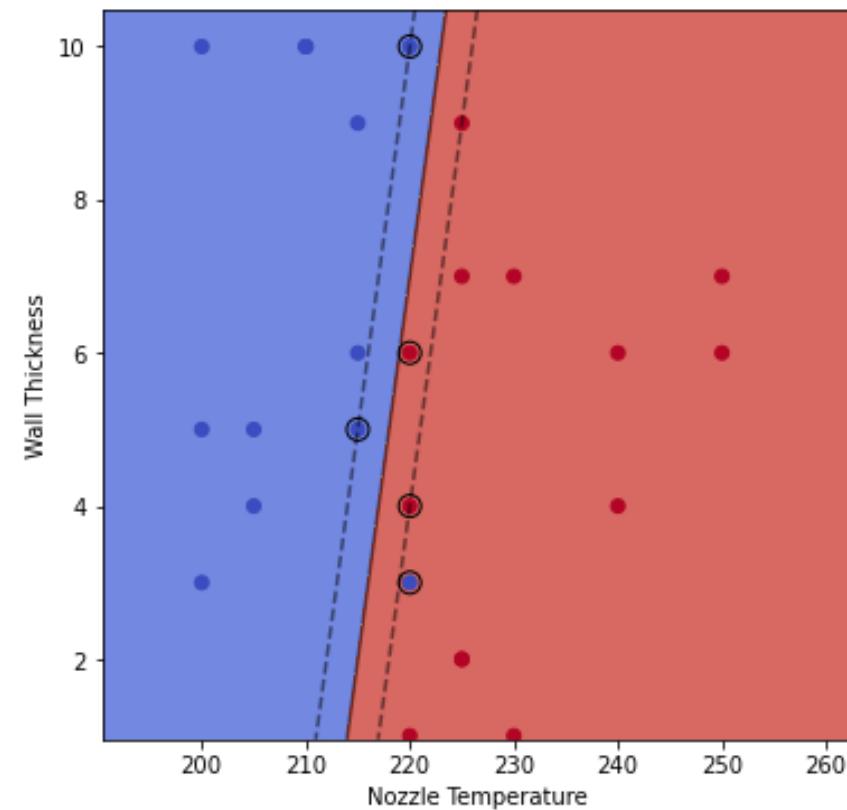


C (penalty/regularization)

C = 0.1

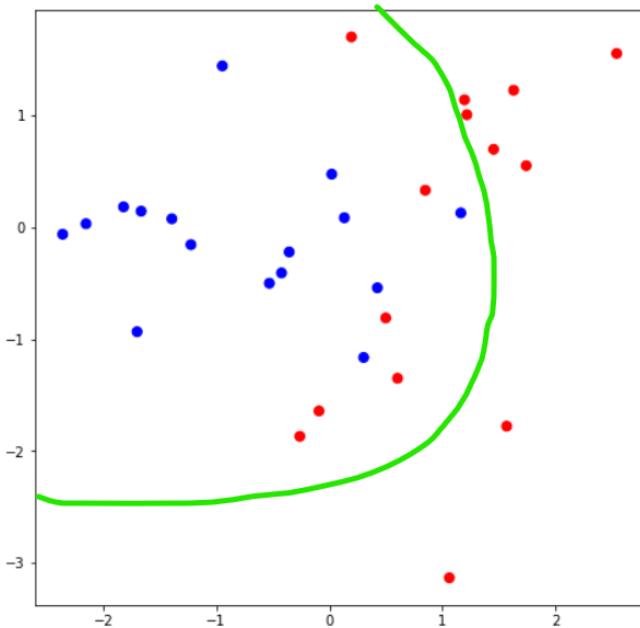


C = 10

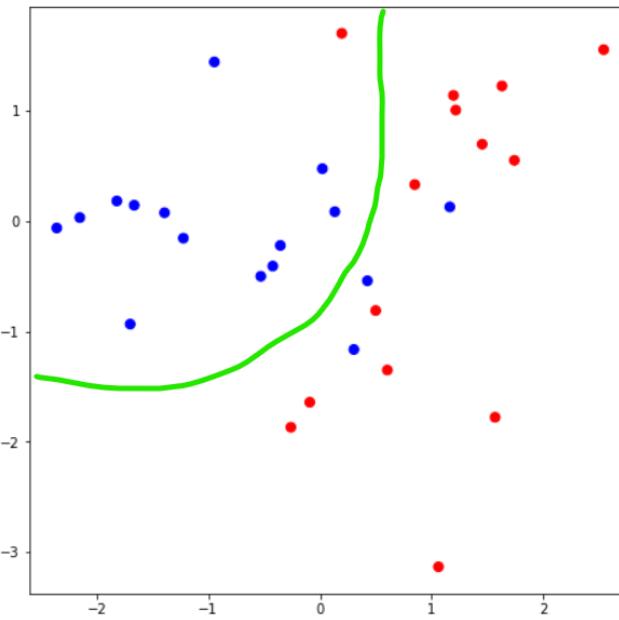


C (penalty/regularization)

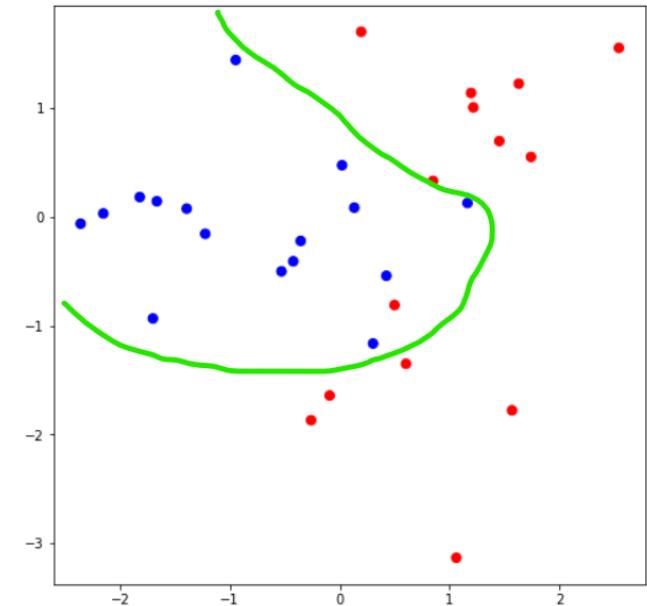
Lage C



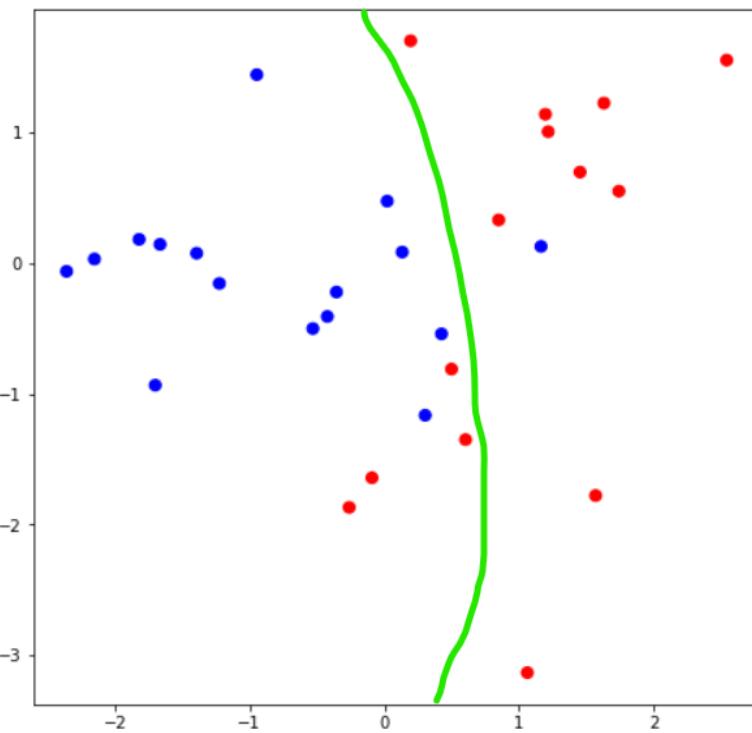
Middel C



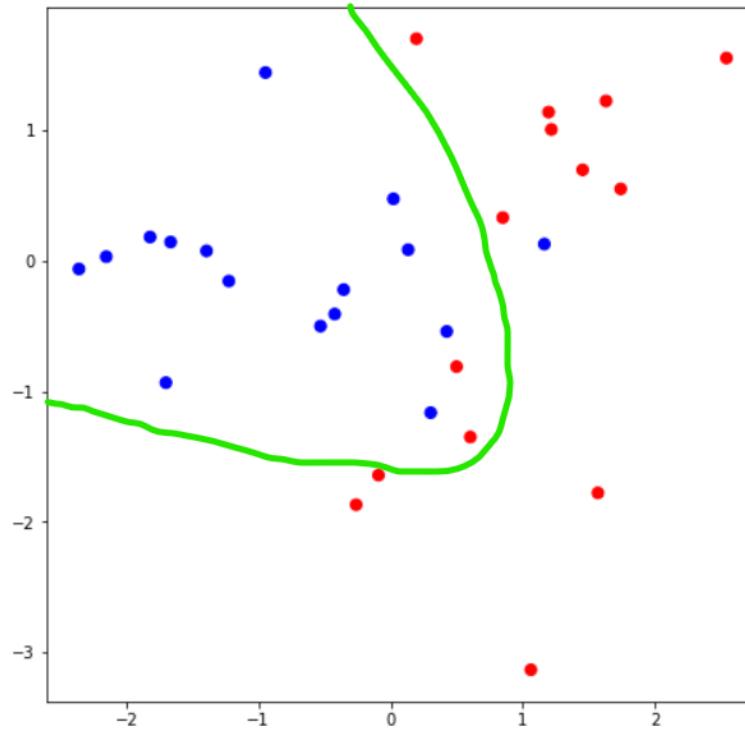
Hoge C



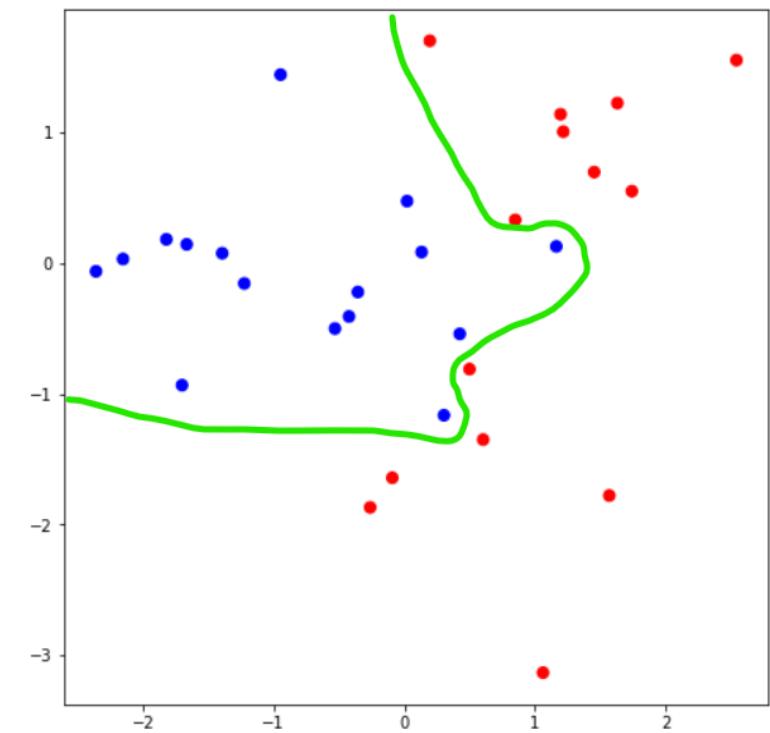
Lage Gamma



Middel Gamma



Hoge Gamma



$$\gamma = \text{"scale"} = \frac{1}{N_{\text{features}} \sigma}$$

$$\gamma = \text{"auto"} = \frac{1}{N_{\text{features}}}$$

Maak een *Support Vector Machine* classifier die voorspelt met welk materiaal is geprint

```
from sklearn.svm import SVC
```

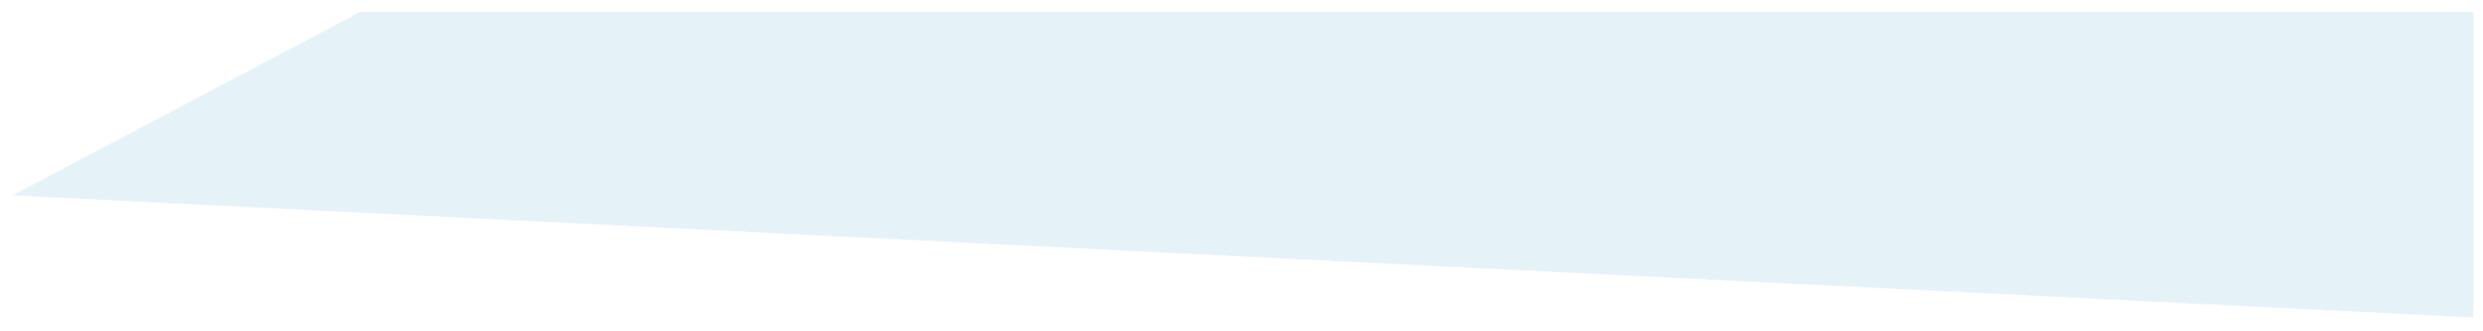
~ 5 minuten

Voordelen

- Goed wanneer er een duidelijke scheidingslijn is tussen de classes
- Werkt goed met weinig data en veel features
- Veel mogelijkheden met kernels

Nadelen

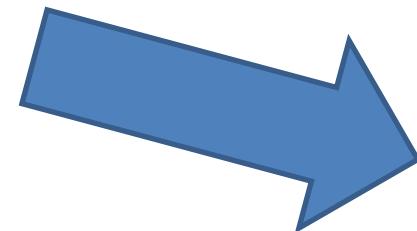
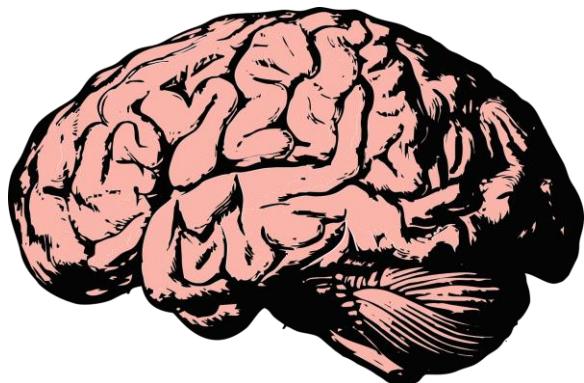
- Gevoelig voor ruis en overlappende classes
- Niet goed met grote datasets
- Veel mogelijkheden met kernels



Methode 3

NEURALE NETWERKEN

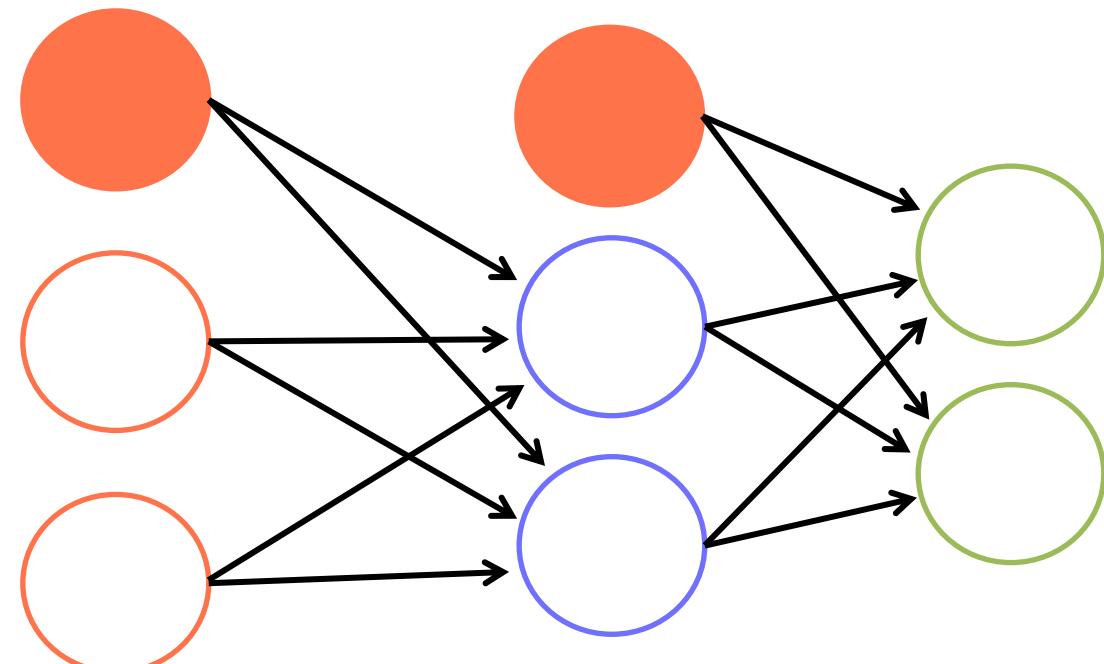
Neurale netwerken



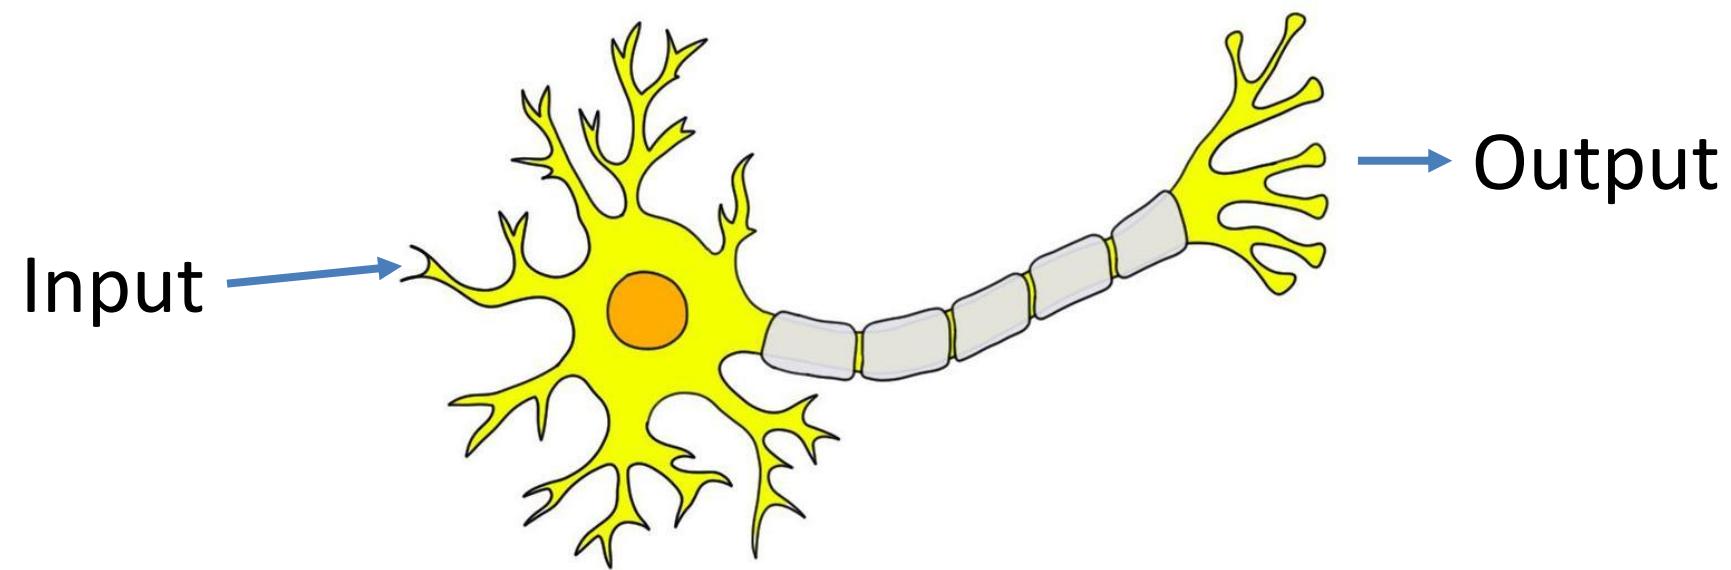
$$\begin{aligned}\frac{\partial E}{\partial \mathbf{W}^{(L)}} &= \frac{\partial E}{\partial \mathbf{y}} \frac{\partial f_L(\mathbf{W}^{(L)} \mathbf{x}^{(L-1)})}{\partial \mathbf{W}^{(L)}} \\ &= \frac{\partial E}{\partial \mathbf{y}} \odot \frac{\partial f_L(\mathbf{W}^{(L)} \mathbf{x}^{(L-1)})}{\partial \mathbf{W}^{(L)} \mathbf{x}^{(L-1)}} \frac{\partial \mathbf{W}^{(L)} \mathbf{x}^{(L-1)}}{\partial \mathbf{W}^{(L)}} \\ &= \frac{\partial E}{\partial \mathbf{y}} \odot f'_L(\mathbf{W}^{(L)} \mathbf{x}^{(L-1)}) (\mathbf{x}^{(L-1)})^T\end{aligned}$$

Substitute $\delta_L = \frac{\partial E}{\partial \mathbf{y}} \odot f'_L(\mathbf{W}^{(L)} \mathbf{x}^{(L-1)})$

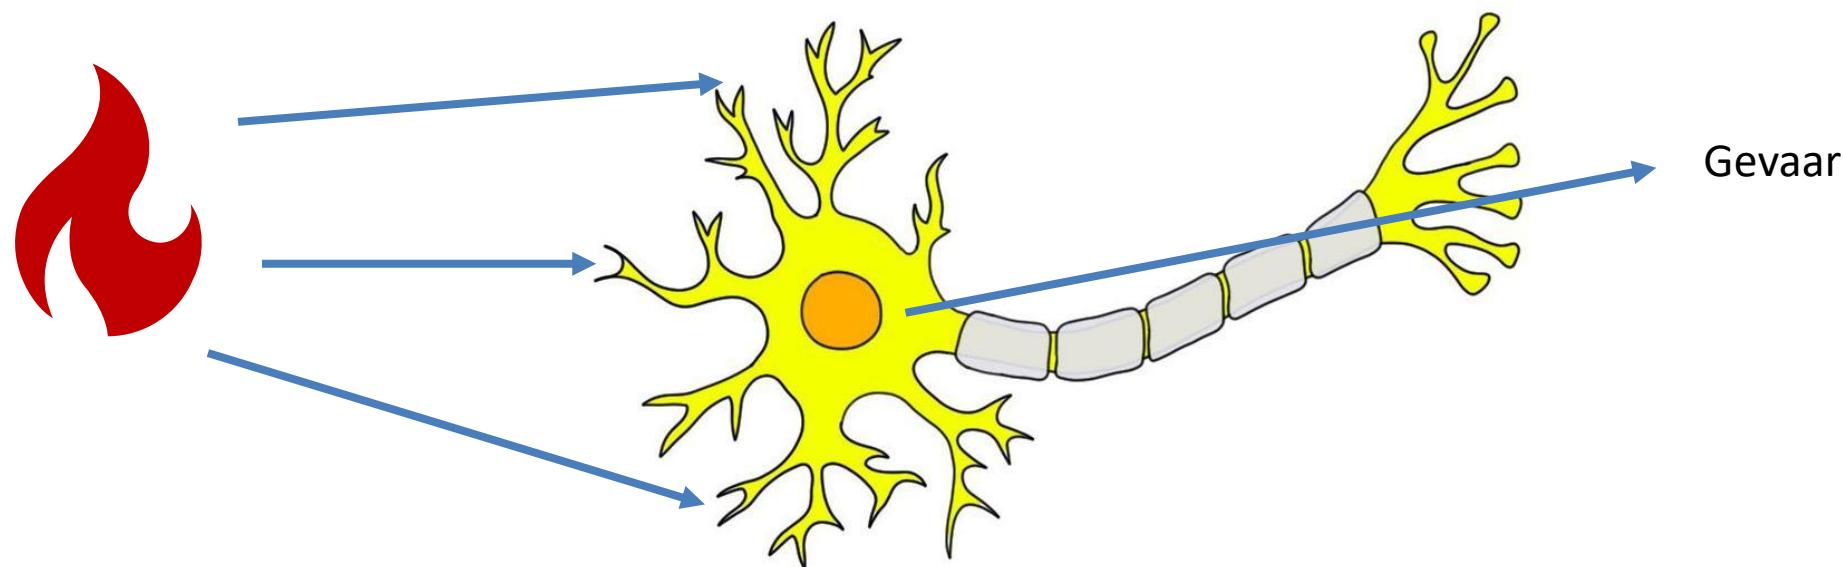
$$\frac{\partial E}{\partial \mathbf{W}^{(L)}} = \delta_L (\mathbf{x}^{(L-1)})^T$$



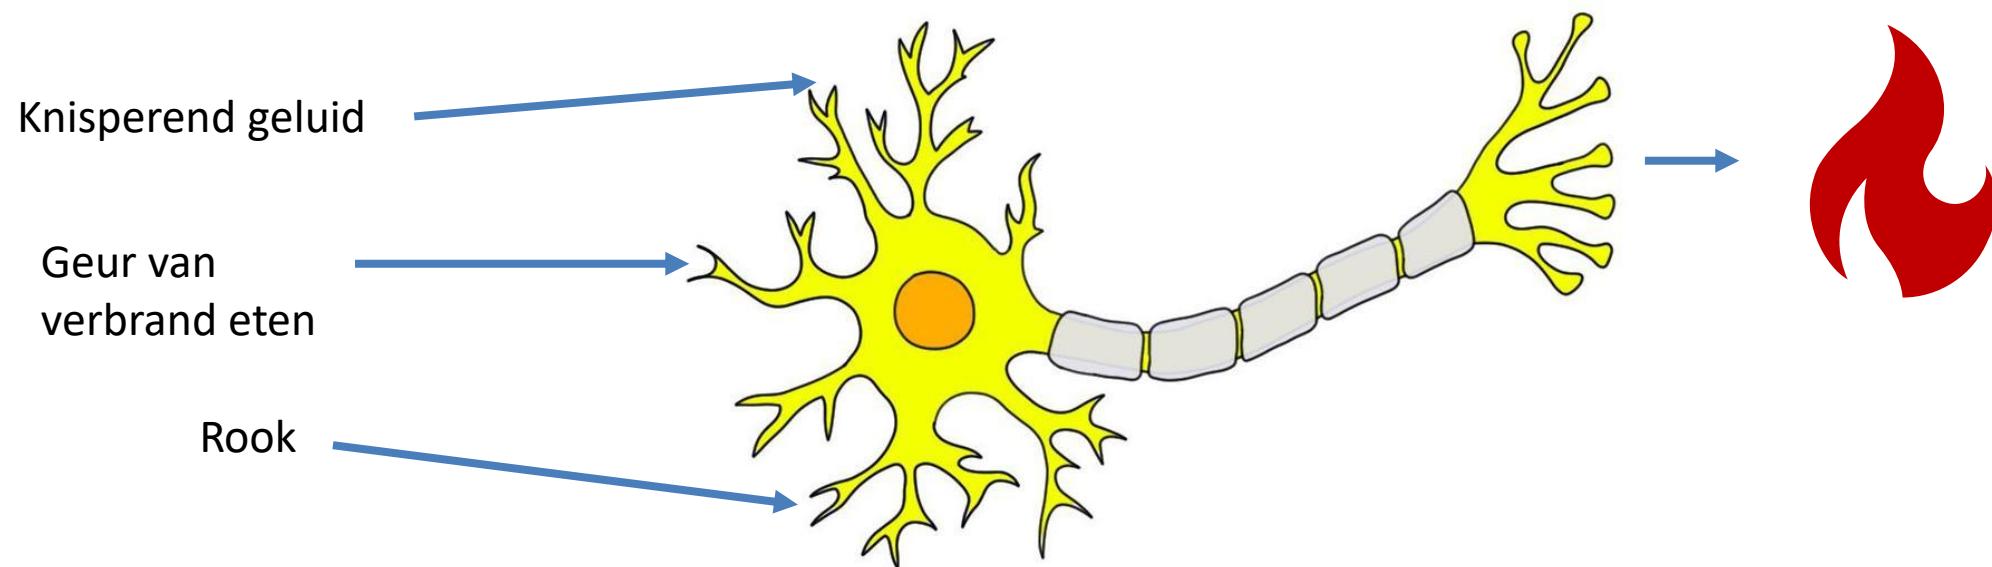
McCulloch, W. S., & Pitts, W. (1943). A logical calculus of the ideas immanent in nervous activity. *The bulletin of mathematical biophysics*, 5(4), 115–133.

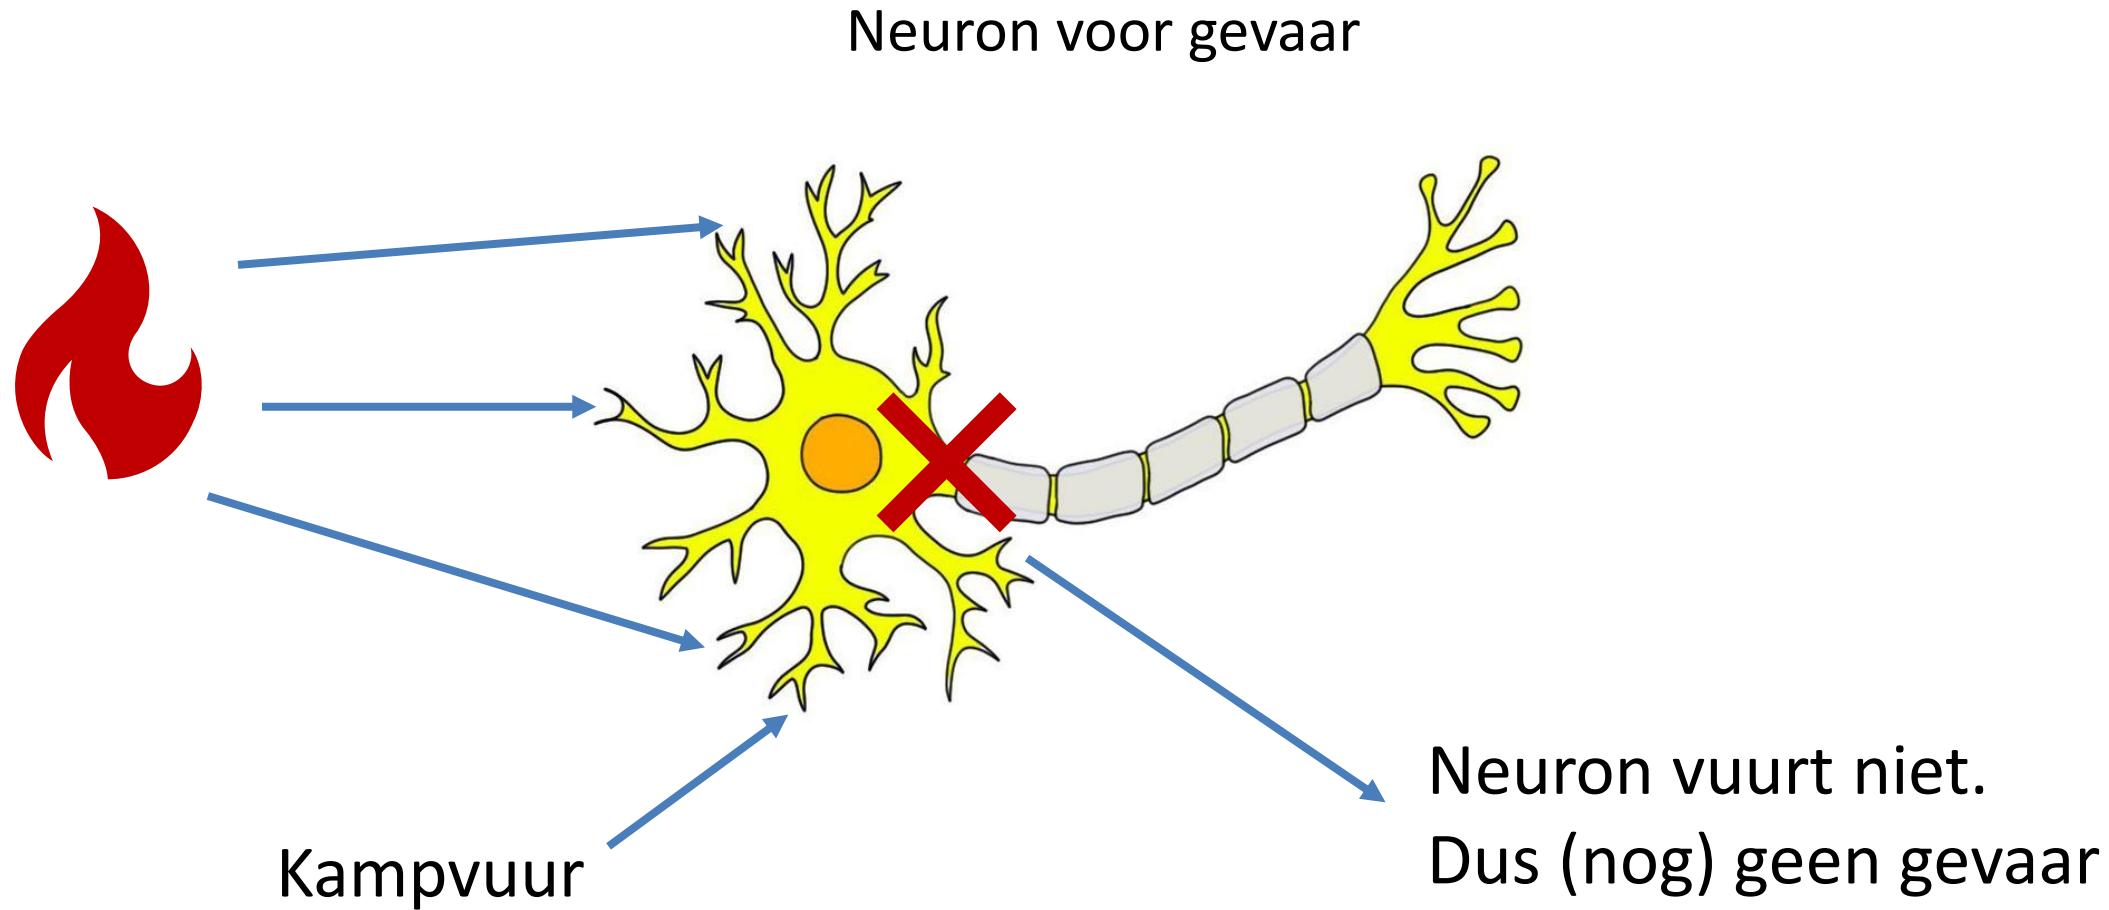


Neuron voor gevaar

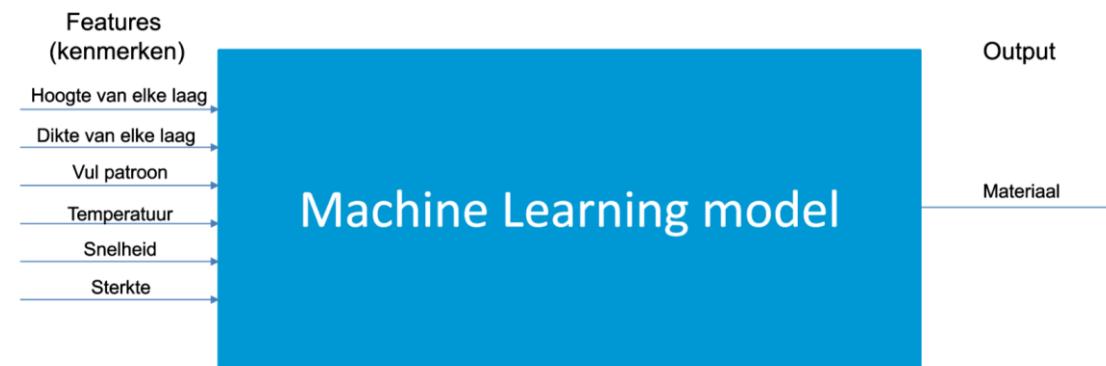
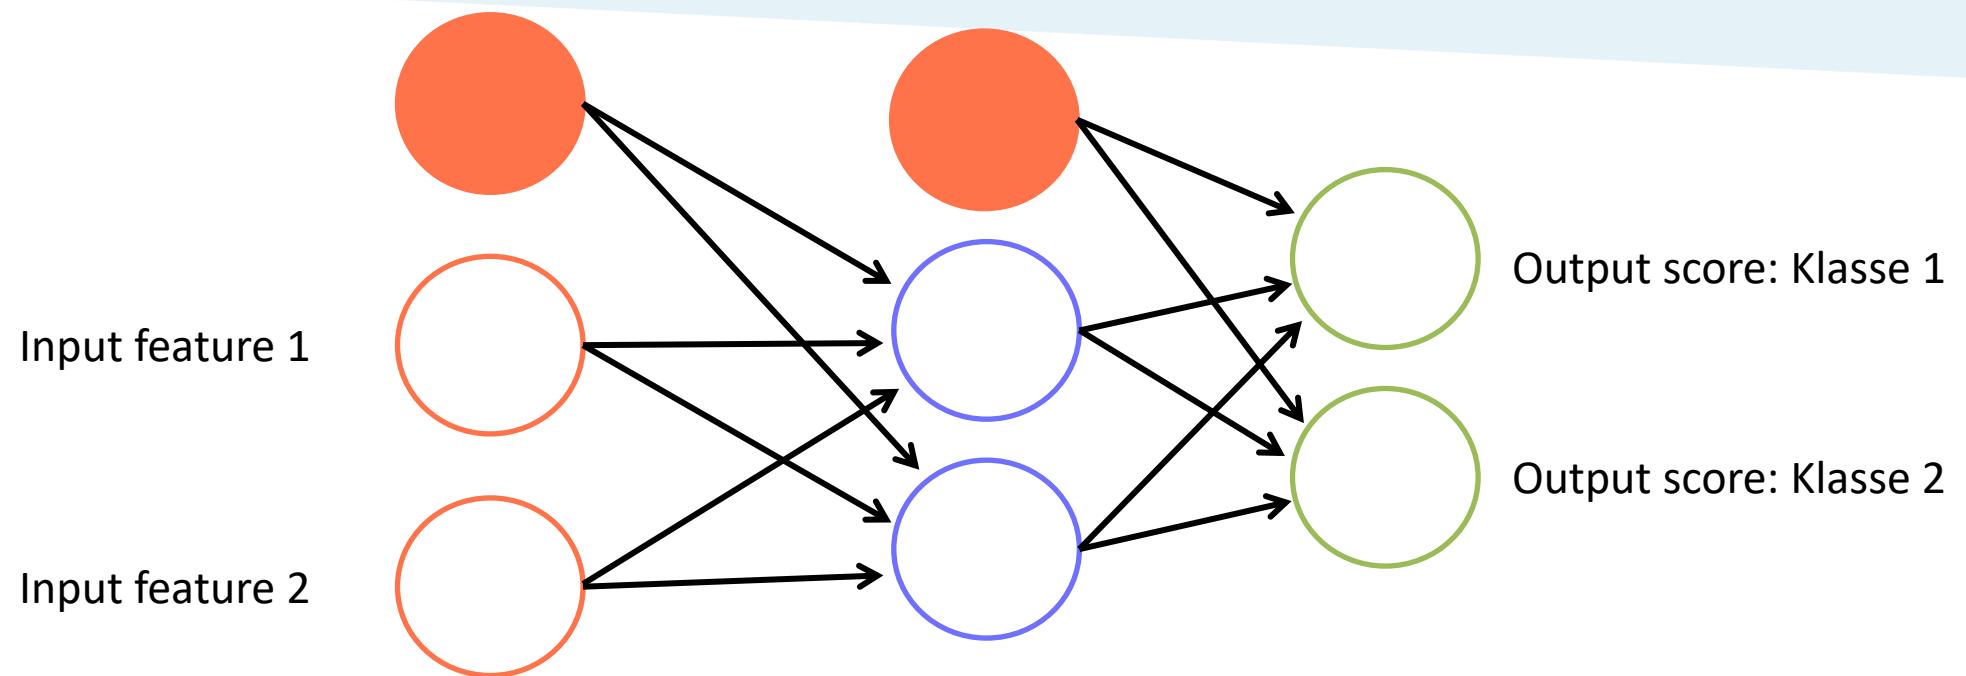


Neuron voor vuur

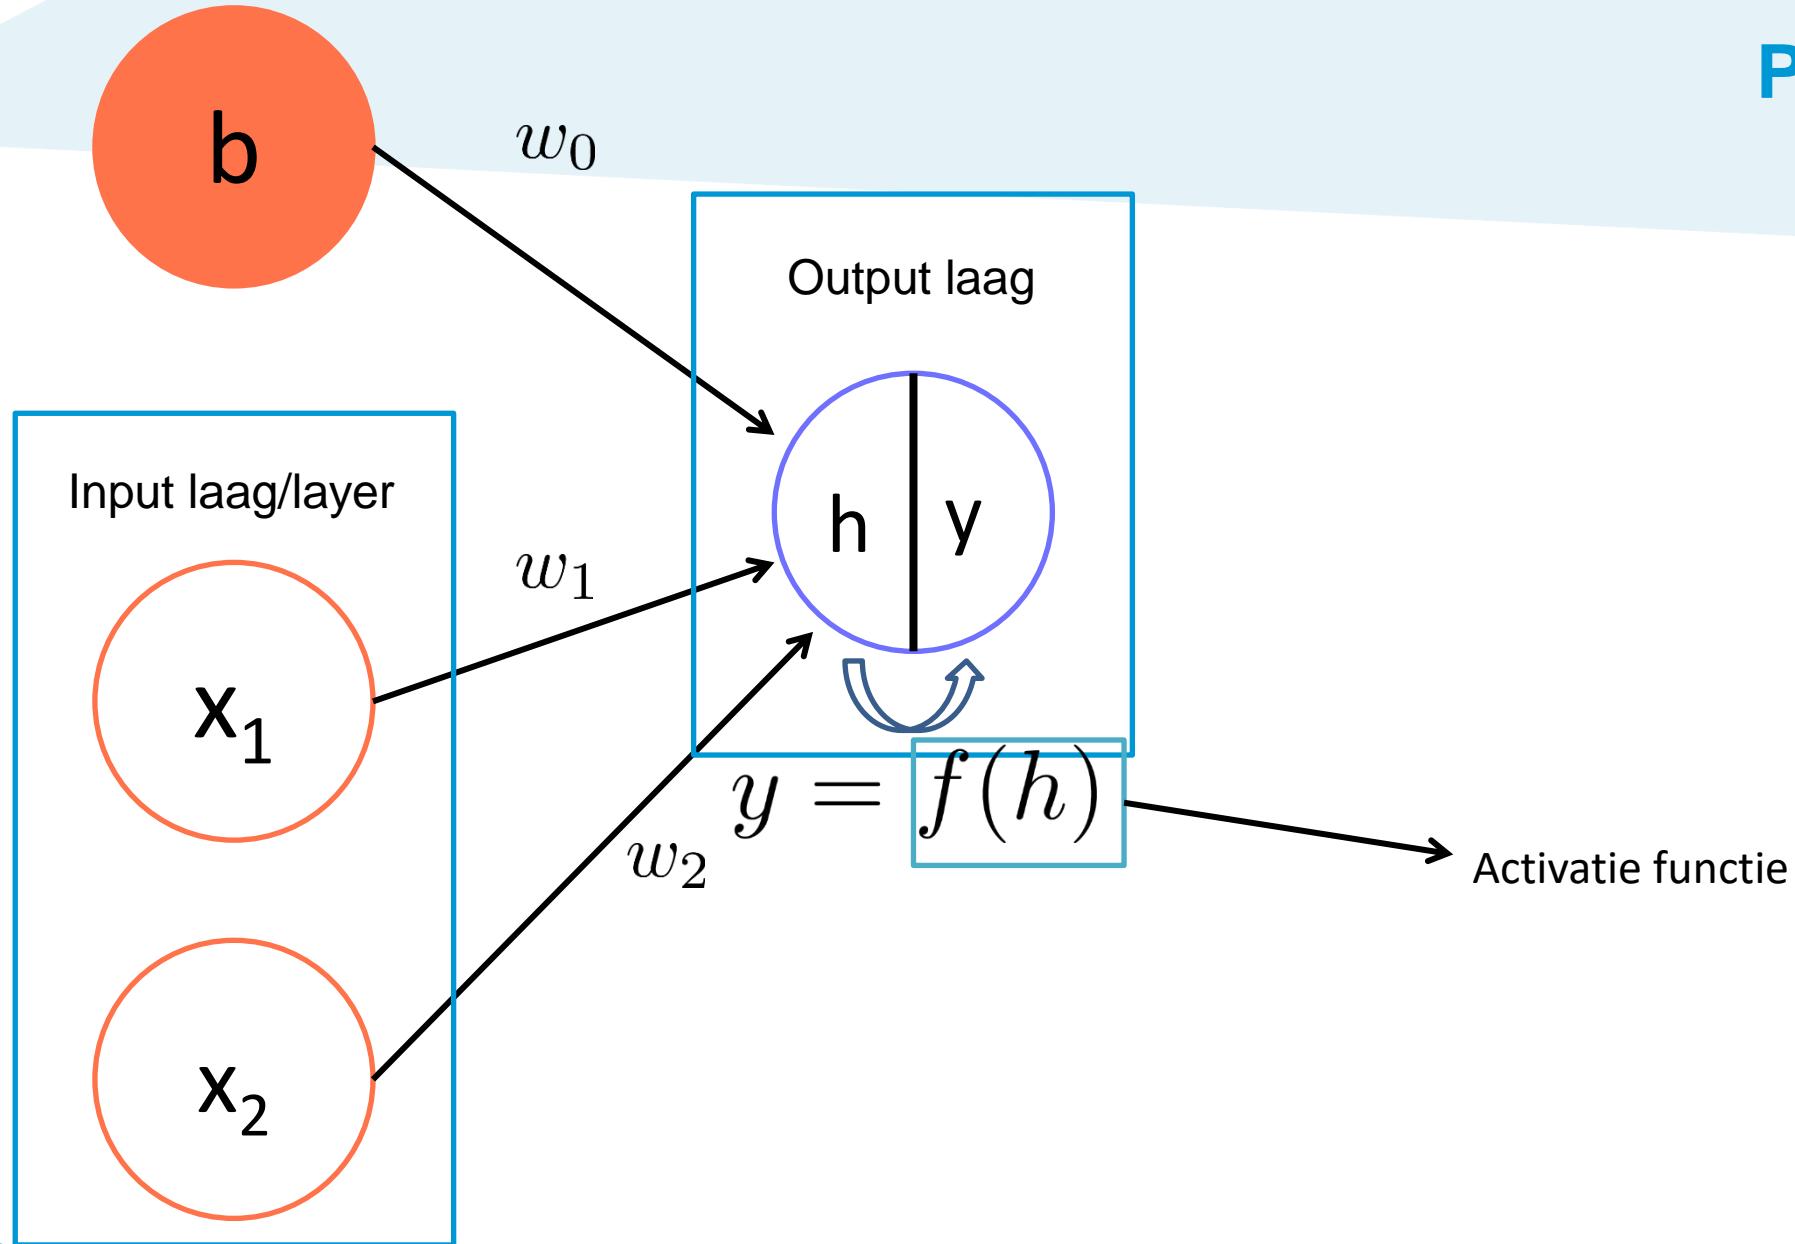




Neurale netwerk (basics)

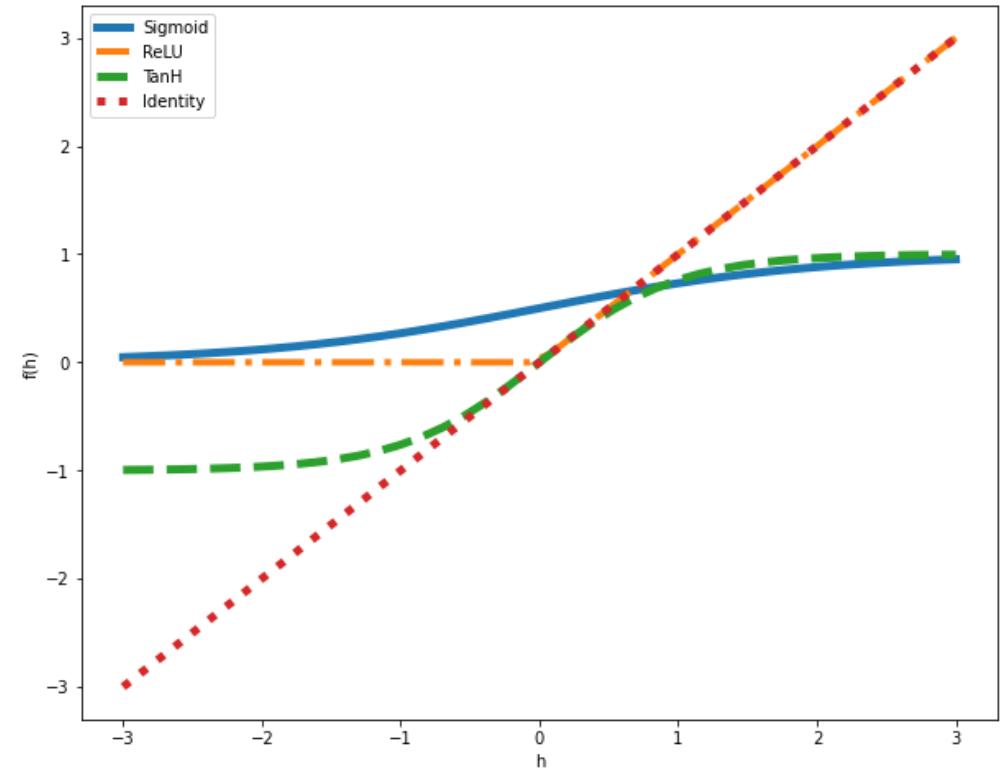


Perceptron

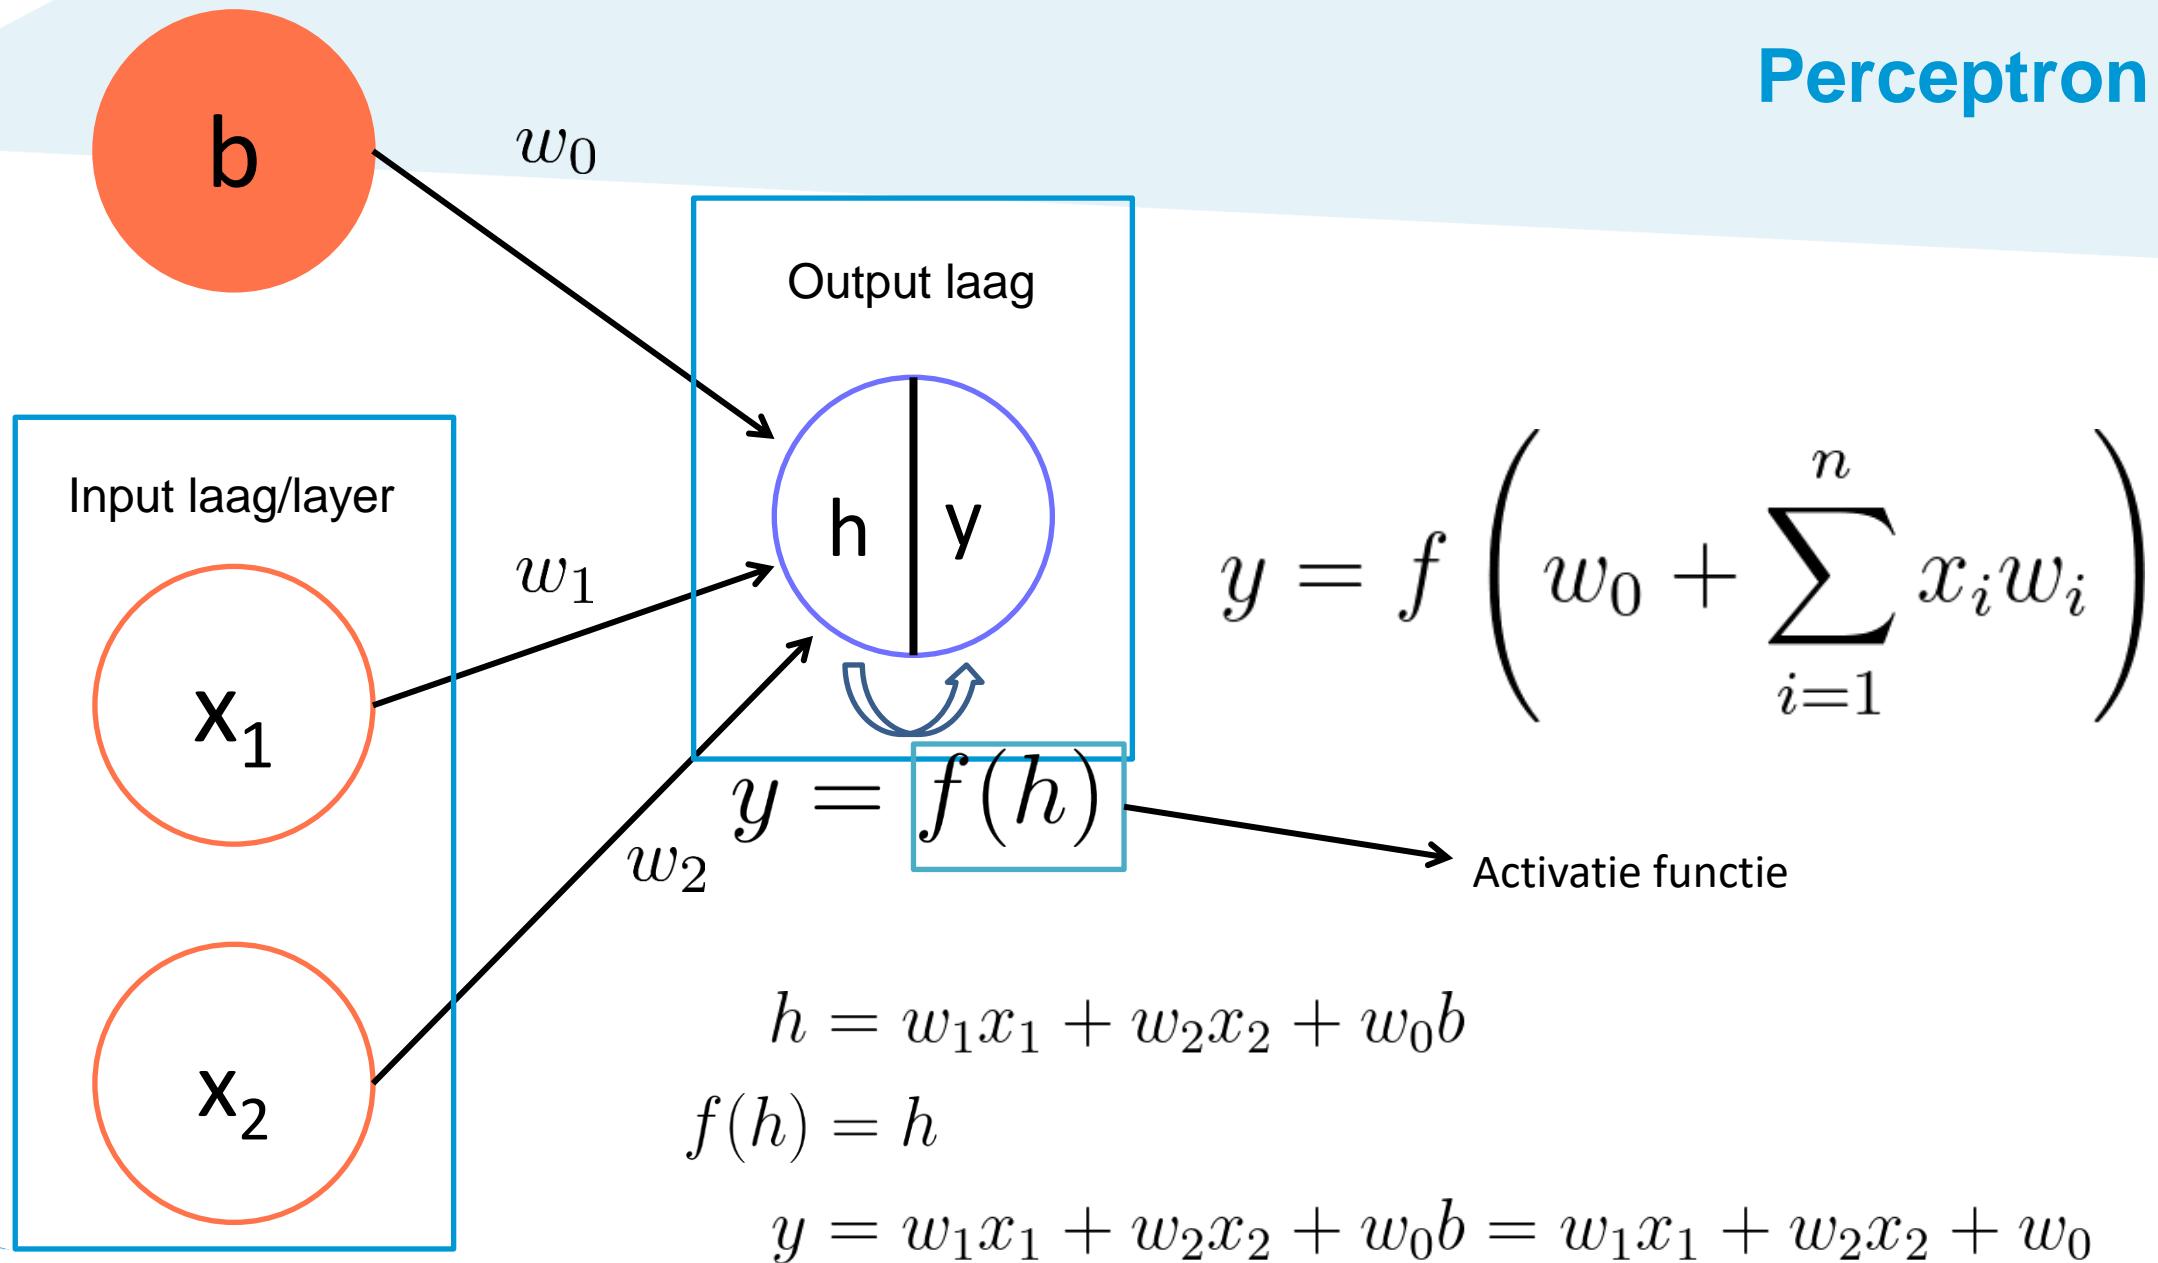


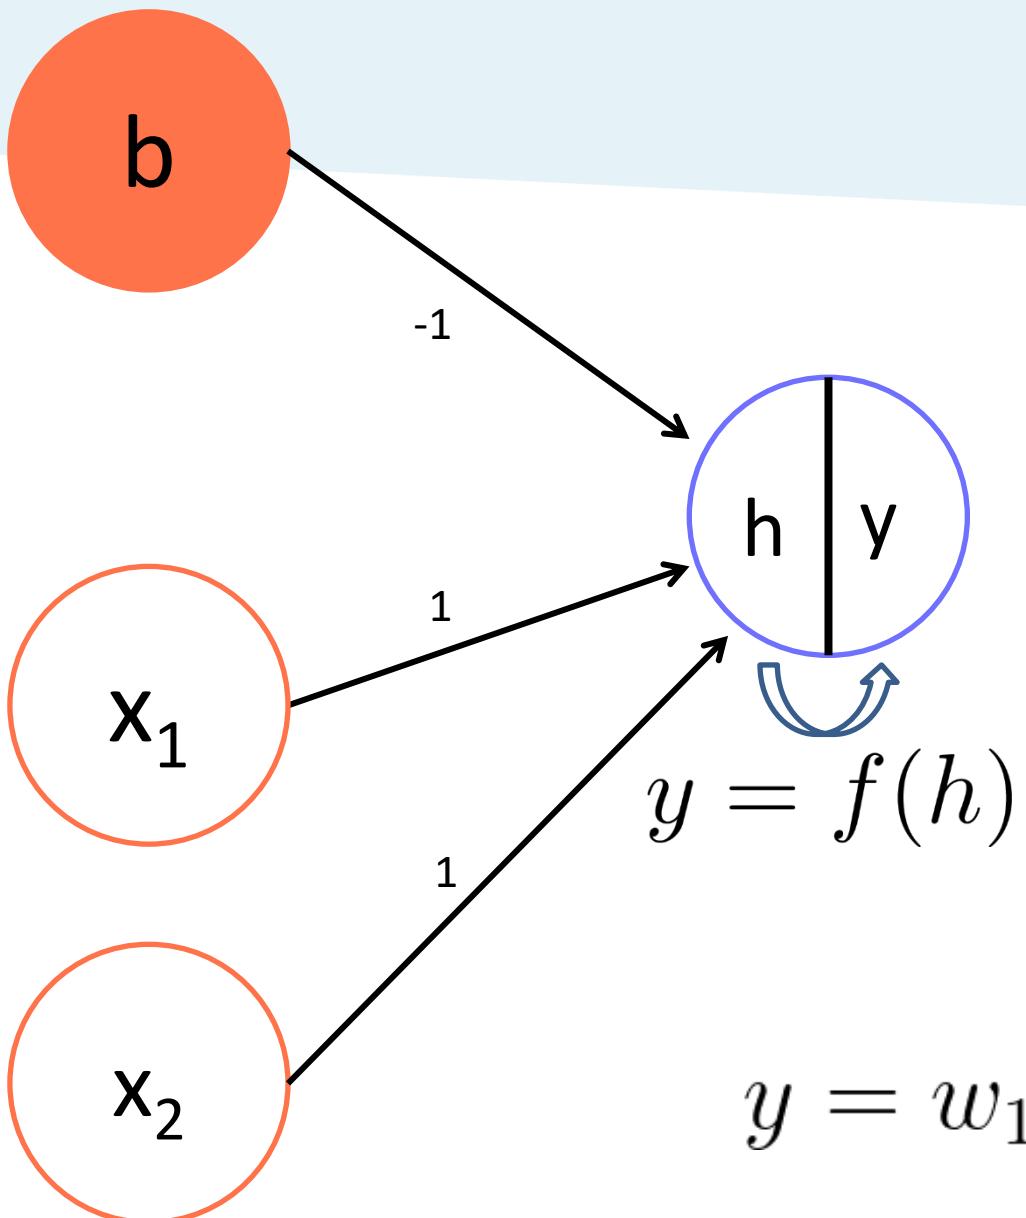
Activatie functies

Function name	Function
Linear	$f(x) = x$
Sigmoid/Logistic	$f(x) = \frac{1}{1+e^{-x}}$
ReLU	$f(x) = \max(0, x)$
Hyperbolic tangent	$f(x) = \tanh x$
Softmax	$\sigma(\mathbf{z})_i = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$



Perceptron





$$y = w_1x_1 + w_2x_2 + w_0b$$

AND-operator

$$\begin{aligned}(x_1=1) \cap (x_2=1) &= 1 \\ (x_1=1) \cap (x_2=0) &= 0 \\ (x_1=0) \cap (x_2=1) &= 0 \\ (x_1=0) \cap (x_2=0) &= 0\end{aligned}$$

$$y = w_1x_1 + w_2x_2 + w_0b$$

$$y = 1 \cdot x_1 + 1 \cdot x_2 - 1$$

$$1 = 1 \cdot 1 + 1 \cdot 1 - 1$$

$$0 = 1 \cdot 1 + 1 \cdot 0 - 1$$

$$0 = 1 \cdot 0 + 1 \cdot 1 - 1$$

$$-1 = 1 \cdot 0 + 1 \cdot 0 - 1 \text{ ~}$$

AND-operator

$$(x_1=1) \cap (x_2=1) = 1$$

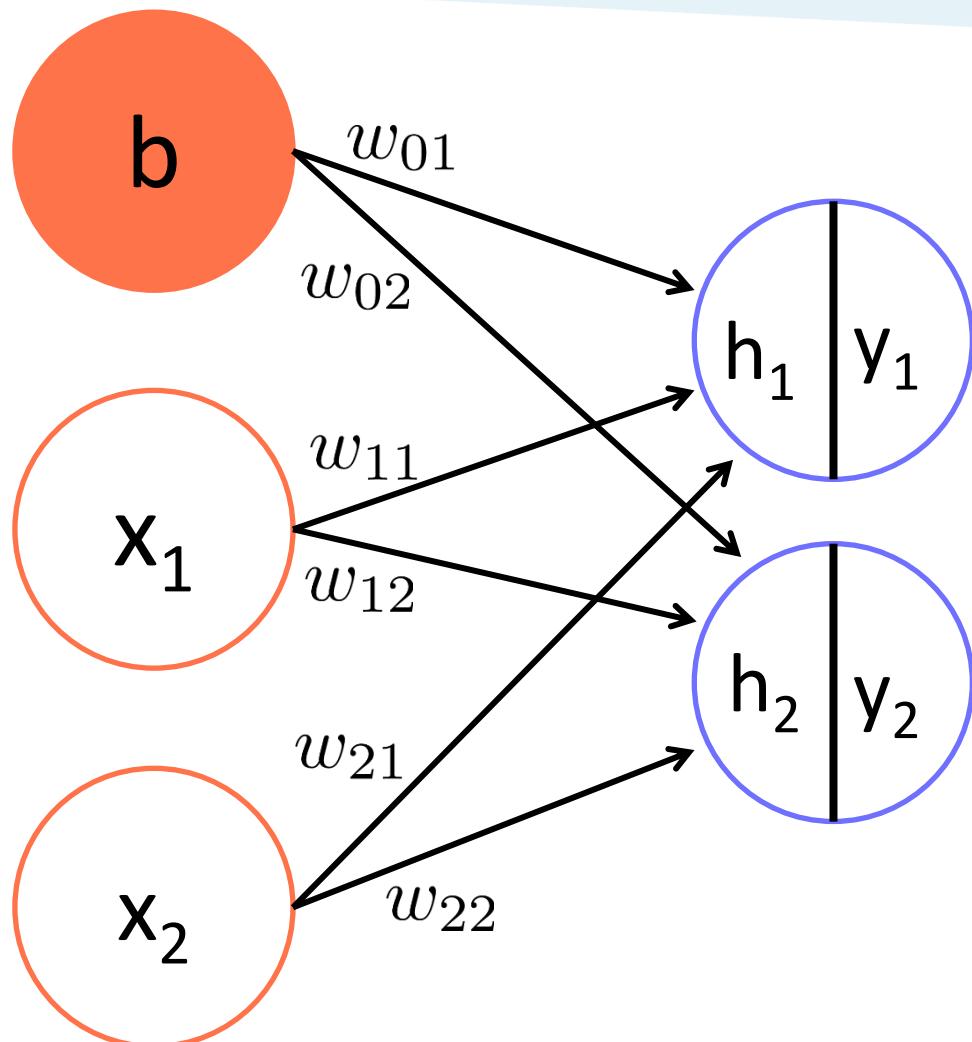
$$(x_1=1) \cap (x_2=0) = 0$$

$$(x_1=0) \cap (x_2=1) = 0$$

$$(x_1=0) \cap (x_2=0) = 0$$

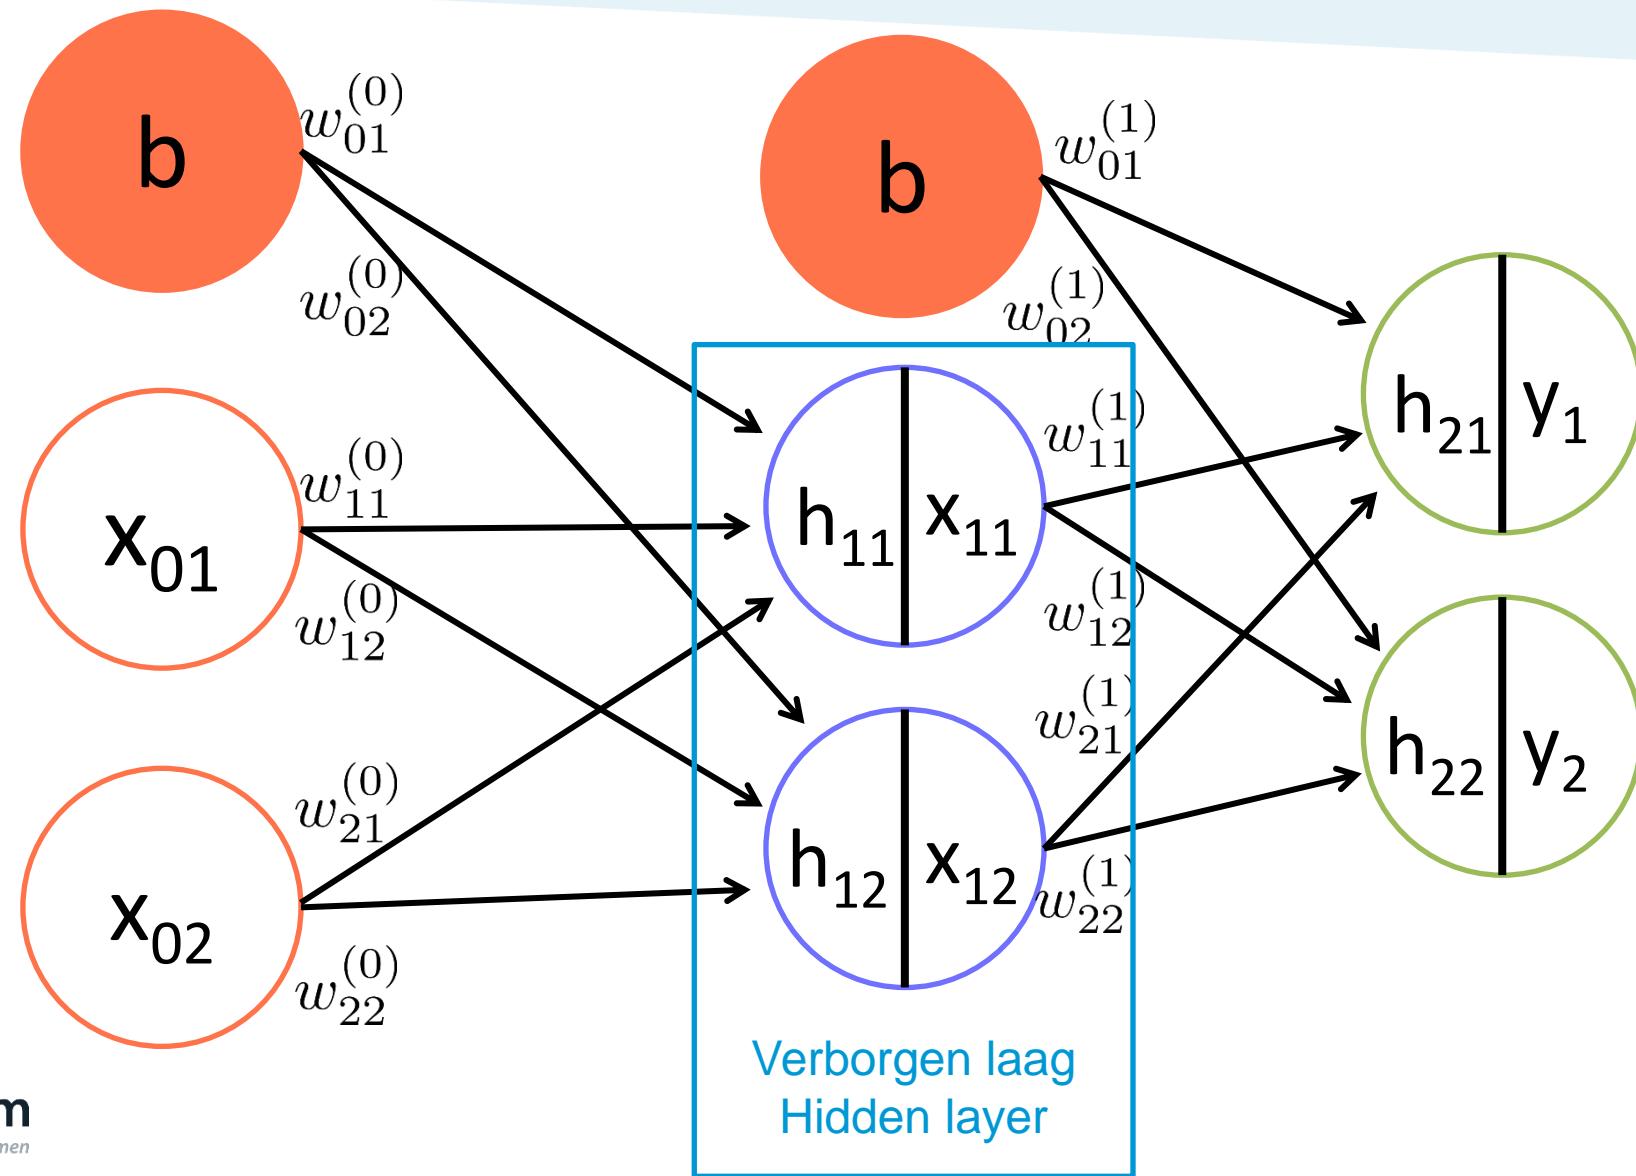
Function name	Function
Linear	$f(x) = x$
Sigmoid/Logistic	$f(x) = \frac{1}{1+e^{-x}}$
ReLU	$f(x) = \max(0, x)$
Hyperbolic tangent	$f(x) = \tanh x$
Softmax	$\sigma(\mathbf{z})_i = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$

Perceptron

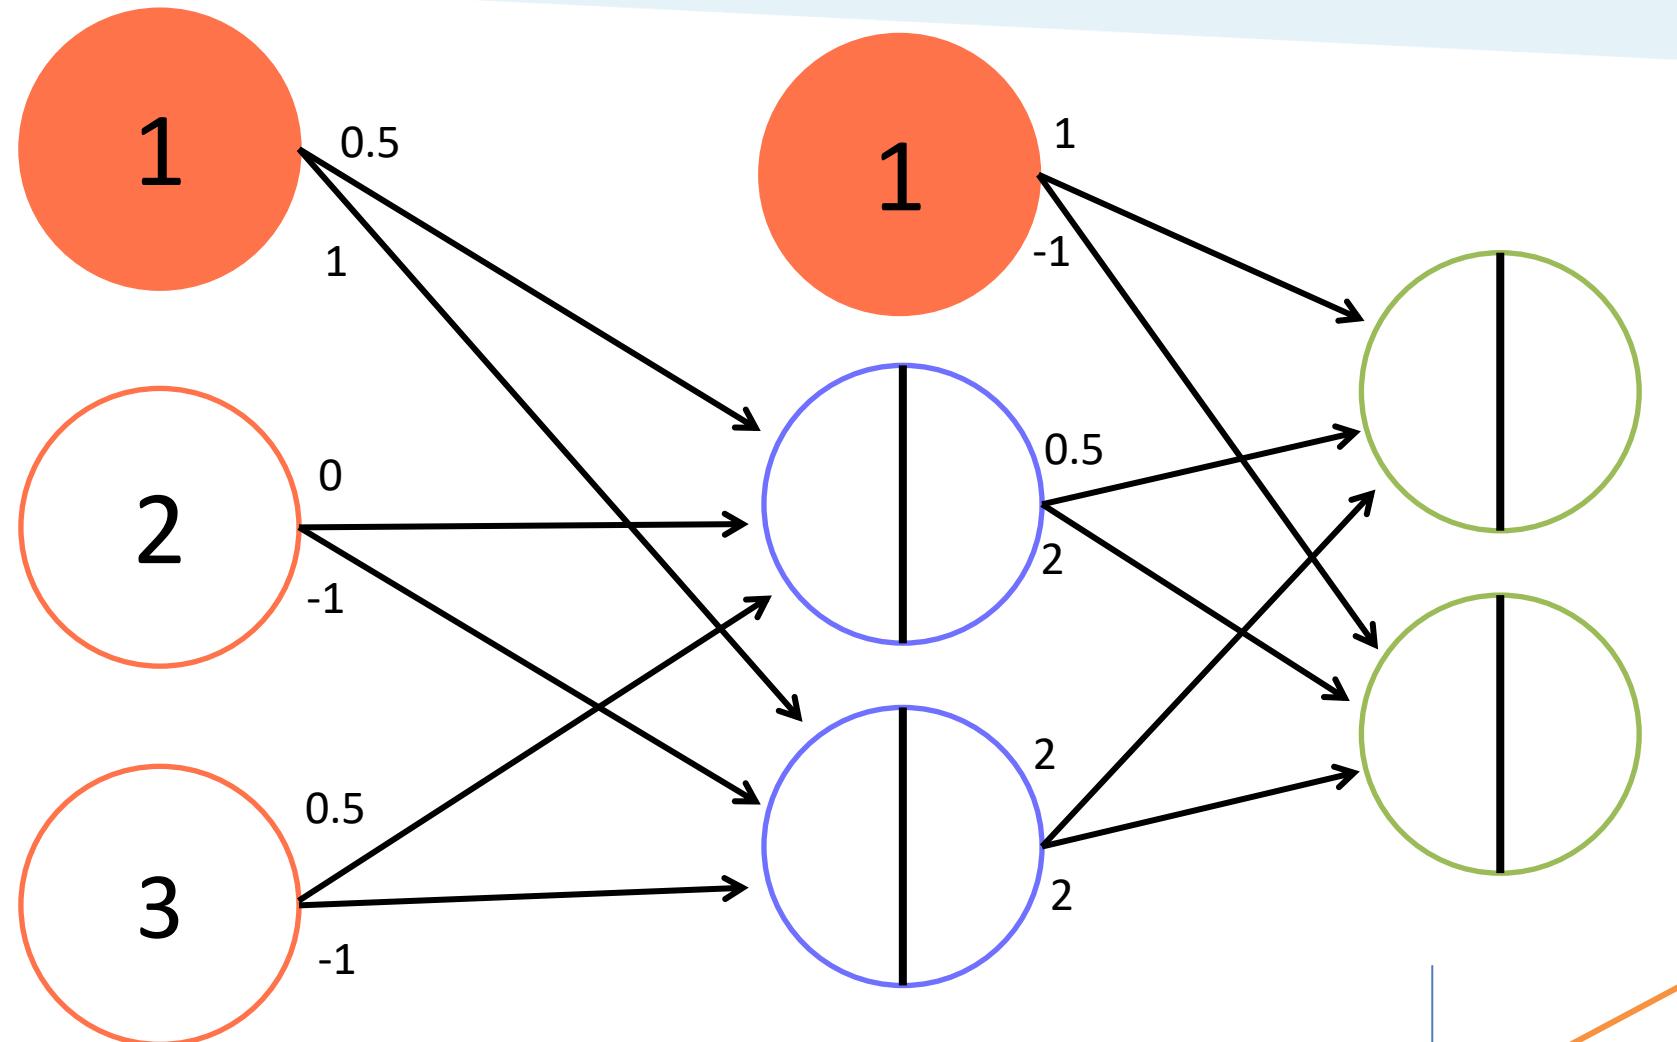


$$y_j = f \left(w_0 + \sum_{i=1}^n x_i w_{ij} \right)$$
$$\mathbf{Y} = \begin{bmatrix} y_1 \\ y_2 \end{bmatrix}$$

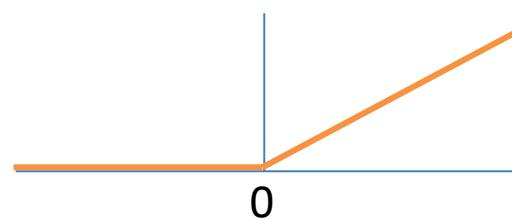
Multi-Layer Perceptron/Neural Network



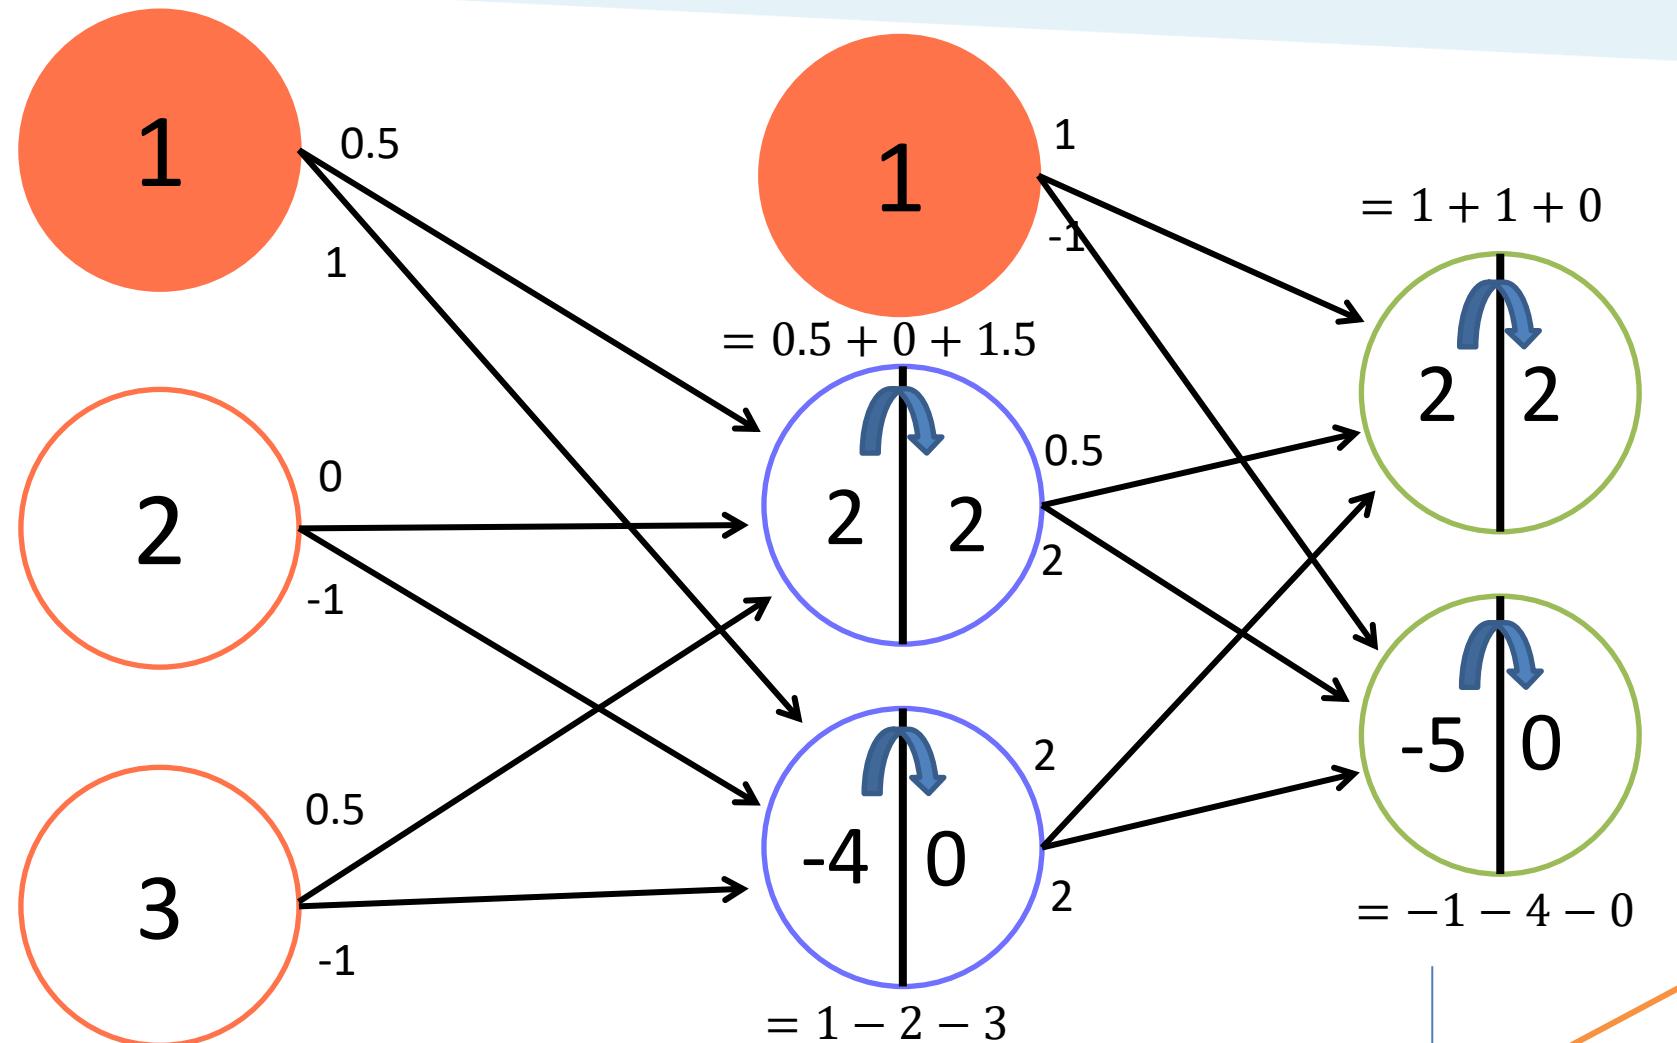
Bereken het neurale netwerk



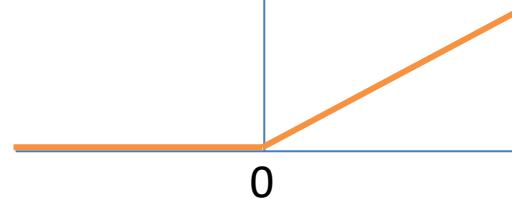
$$\text{ReLU} \quad f(x) = \max(0, x)$$



Multi-Layer Perceptron/Neural Network



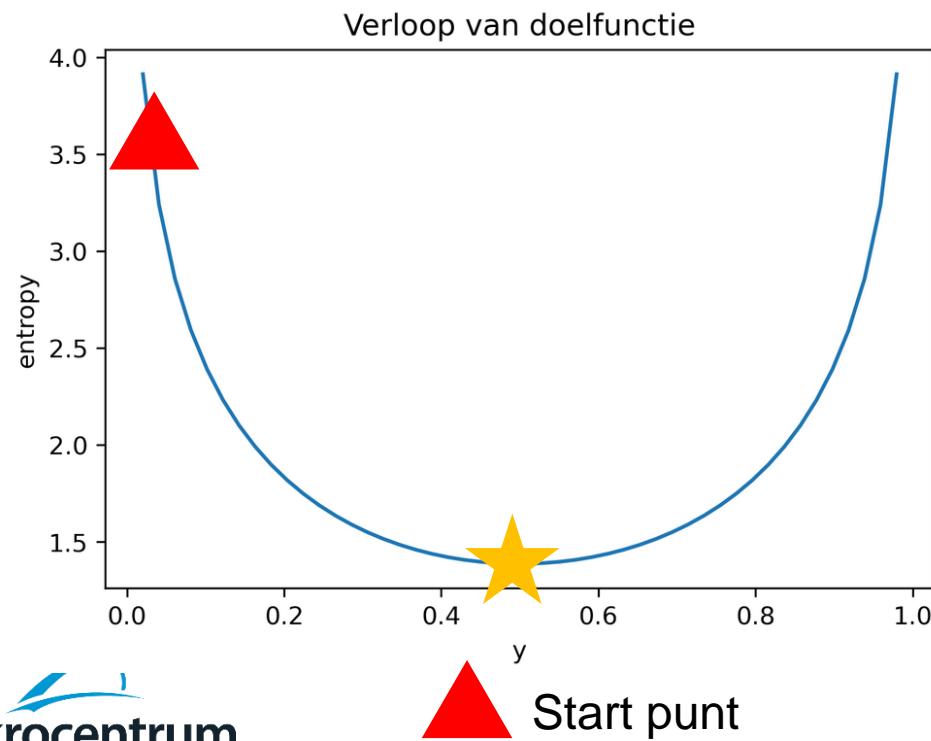
ReLU $f(x) = \max(0, x)$



Hoe leert een Neural Network?

$$\text{MSE} = \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

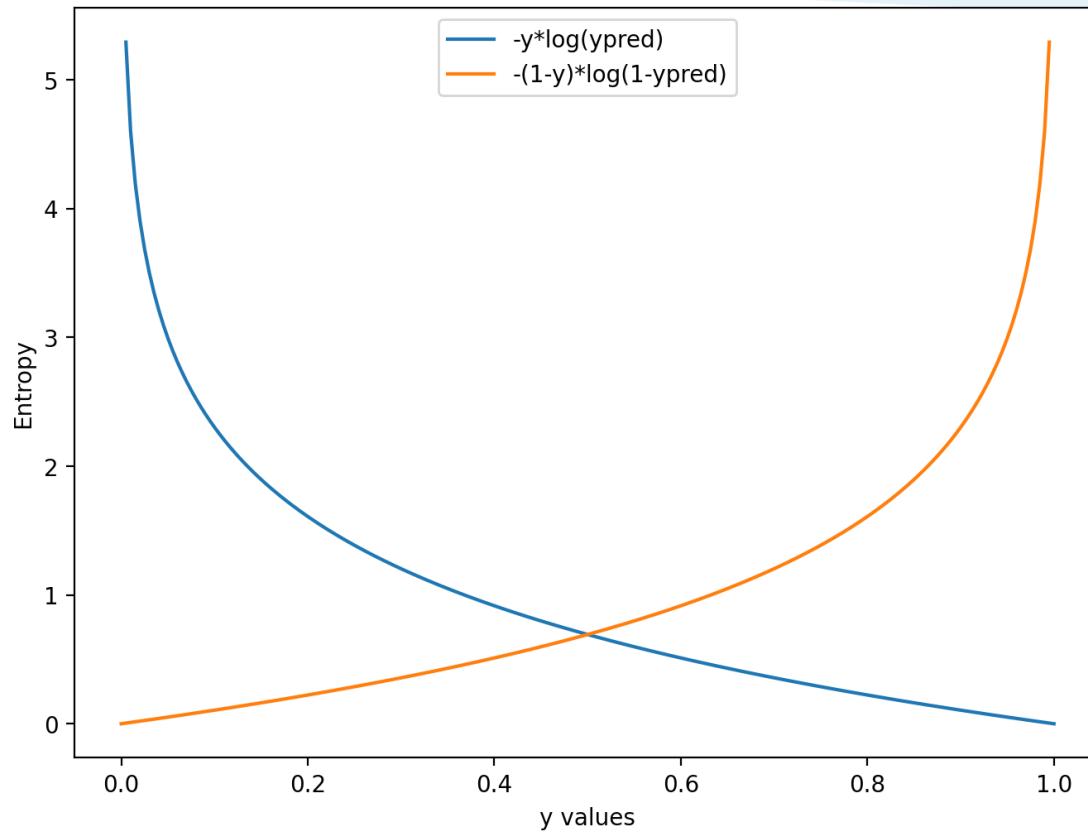
$$\text{Entropy} = - \sum_{i=1}^n y_i \log \hat{y}_i$$



$$\hat{y}_j = f \left(w_0 + \sum_{i=1}^n x_i w_i \right)$$

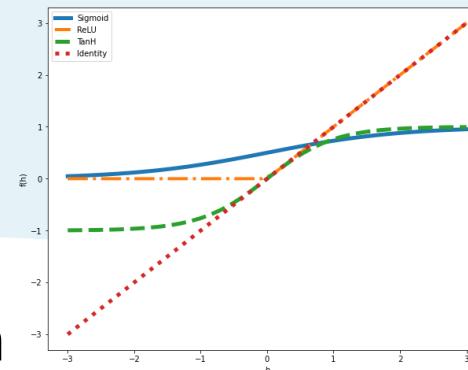
Voor welke \mathbf{W} is de entropy het kleinst?

Entropy function



$$\text{Entropy (binair)} = -y \log \hat{y} - (1 - y) \log(1 - \hat{y})$$

Hyperparameters



- activation = Activation function
- hidden_layer_sizes = (128, 256, 512)
Dit maakt een neuraal netwerk met 3 hidden layers met respectievelijk 128, 256 en 512 neuronen in de laag
- max_iter = Maximale hoeveelheid iteraties in de training
- tol = Hoeveel het model moet verbeteren om door te gaan met training (tolerantie)
- learning_rate= "adaptive" of learning_rate= "constant"
Dit bepaald of de learning rate aangepast wordt als het model minder verbeterd na een paar iteraties
- initial_learning_rate = 0.001
Learning rate die aan het begin wordt gebruikt

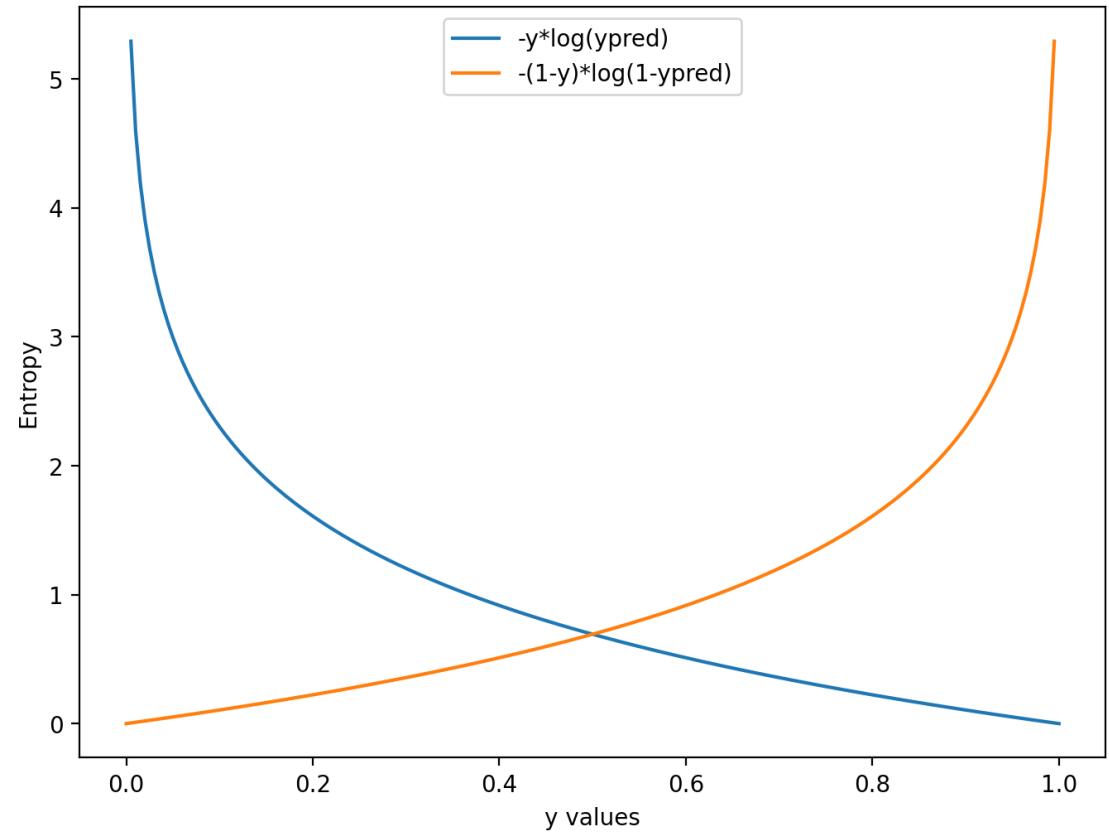
Maak een *Multi Layer Perceptron/Neural Network* classifier die voorspelt met welk materiaal is geprint

```
from sklearn.neural_network import MLPClassifier
```

~10 minuten

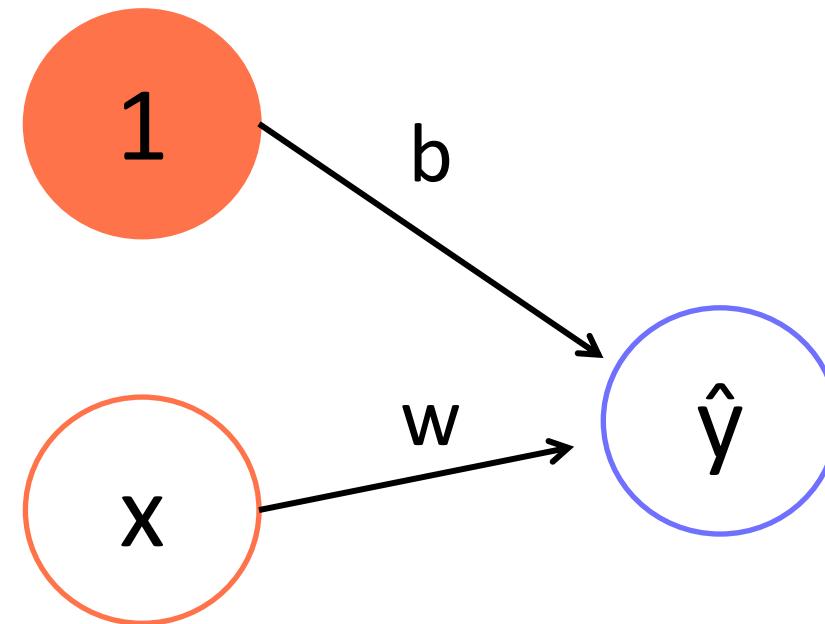
Wat zit achter een Neural Network?

- (Stochastic) Gradient descent
- Wees gewaarschuwd: wiskunde



$$\text{Entropy (binair)} = -y \log \hat{y} - (1 - y) \log(1 - \hat{y})$$

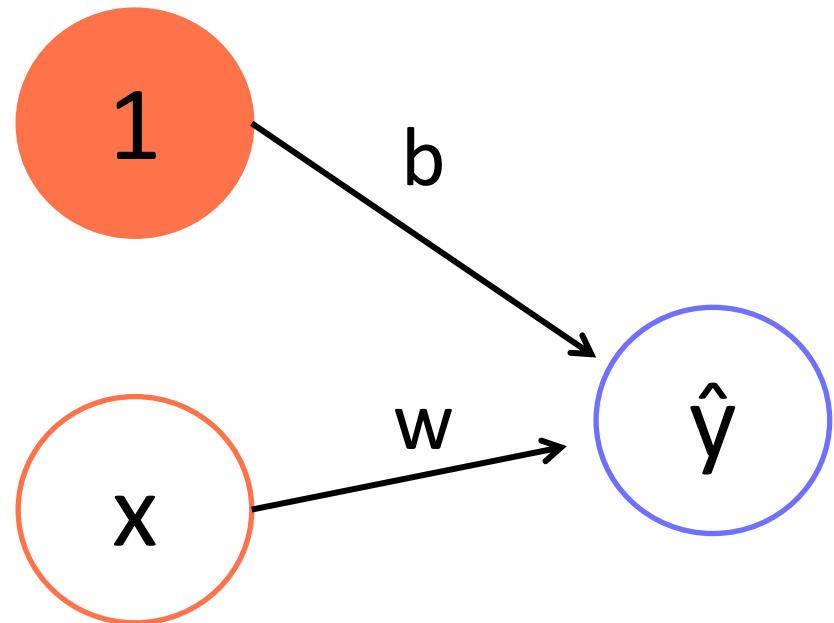
Gradient descend (basis)



$$E(w, b) = -y \log(\hat{y}) - (1 - y) \log(1 - \hat{y})$$

$$\hat{y} = wx + b$$

Gradient descend (basis)



$$\begin{aligned}\min_w E(w, b) \\ \frac{\partial E(w, b)}{\partial w} &= \frac{\partial E(w, b)}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial w} \\ \frac{\partial E(w, b)}{\partial b} &= \frac{\partial E(w, b)}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial b}\end{aligned}$$

$$E(w, b) = -y \log(\hat{y}) - (1 - y) \log(1 - \hat{y})$$

$$\hat{y} = wx + b$$

Gradient descent (basis)

$$E(w, b) = -y \log(\hat{y}) - (1 - y) \log(1 - \hat{y})$$

$$\hat{y} = wx + b$$

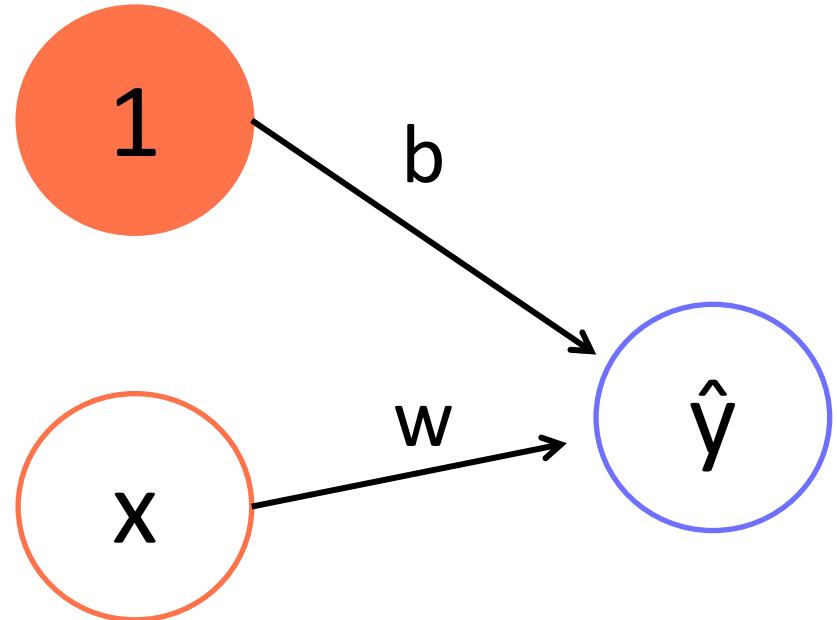
$$\frac{d \log x}{dx} = \frac{1}{x}$$

$$\frac{\partial E(w, b)}{\partial \hat{y}} = -\frac{y}{\hat{y}} + \frac{1 - y}{1 - \hat{y}}$$

$$\frac{\partial \hat{y}}{\partial w} = ?$$

$$\frac{\partial \hat{y}}{\partial b} = ?$$

Gradient descend (basis)



$$E(w, b) = -y \log(\hat{y}) - (1 - y) \log(1 - \hat{y})$$

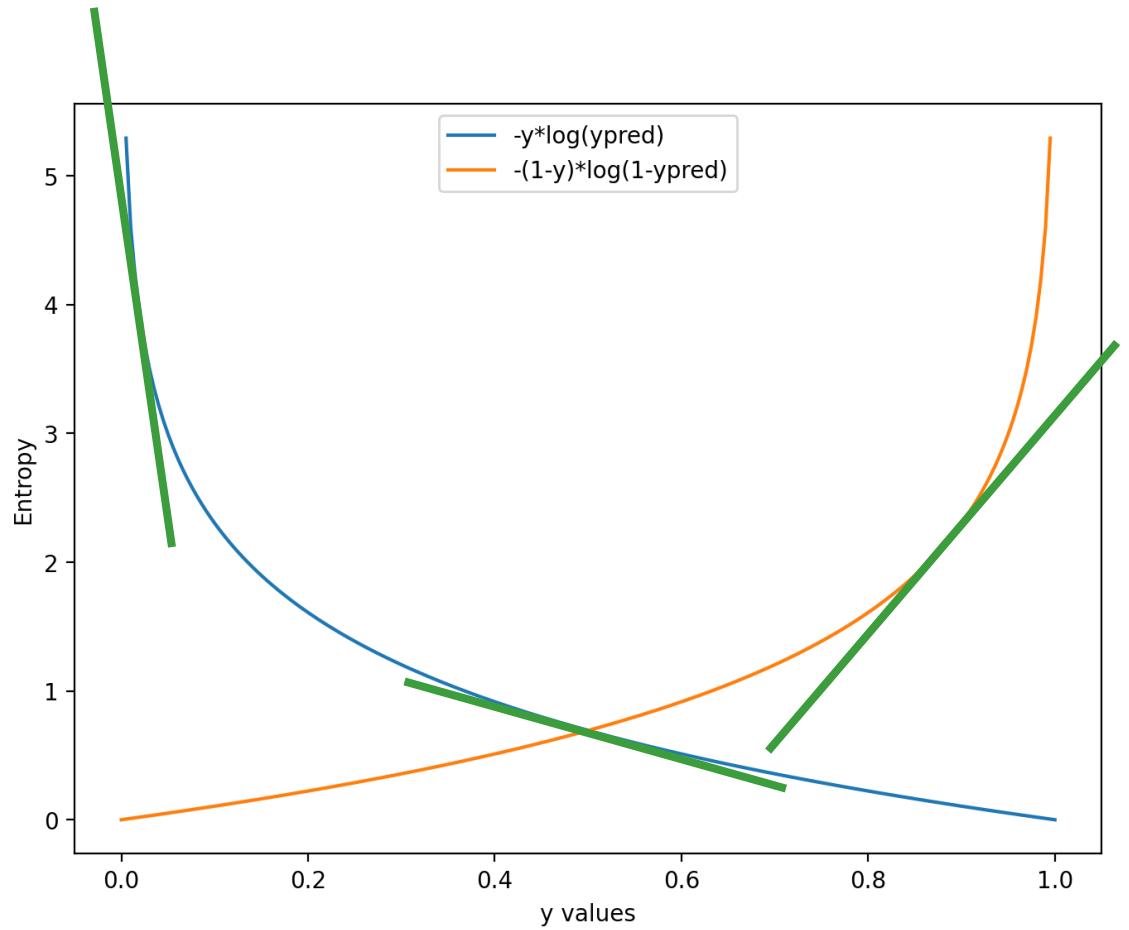
$$\hat{y} = wx + b$$

$$\min_w E(w, b)$$

$$\begin{aligned}\frac{\partial E(w, b)}{\partial w} &= \frac{\partial E(w, b)}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial w} \\ &= \left(\frac{-y}{\hat{y}} + \frac{1 - y}{1 - \hat{y}} \right) x\end{aligned}$$

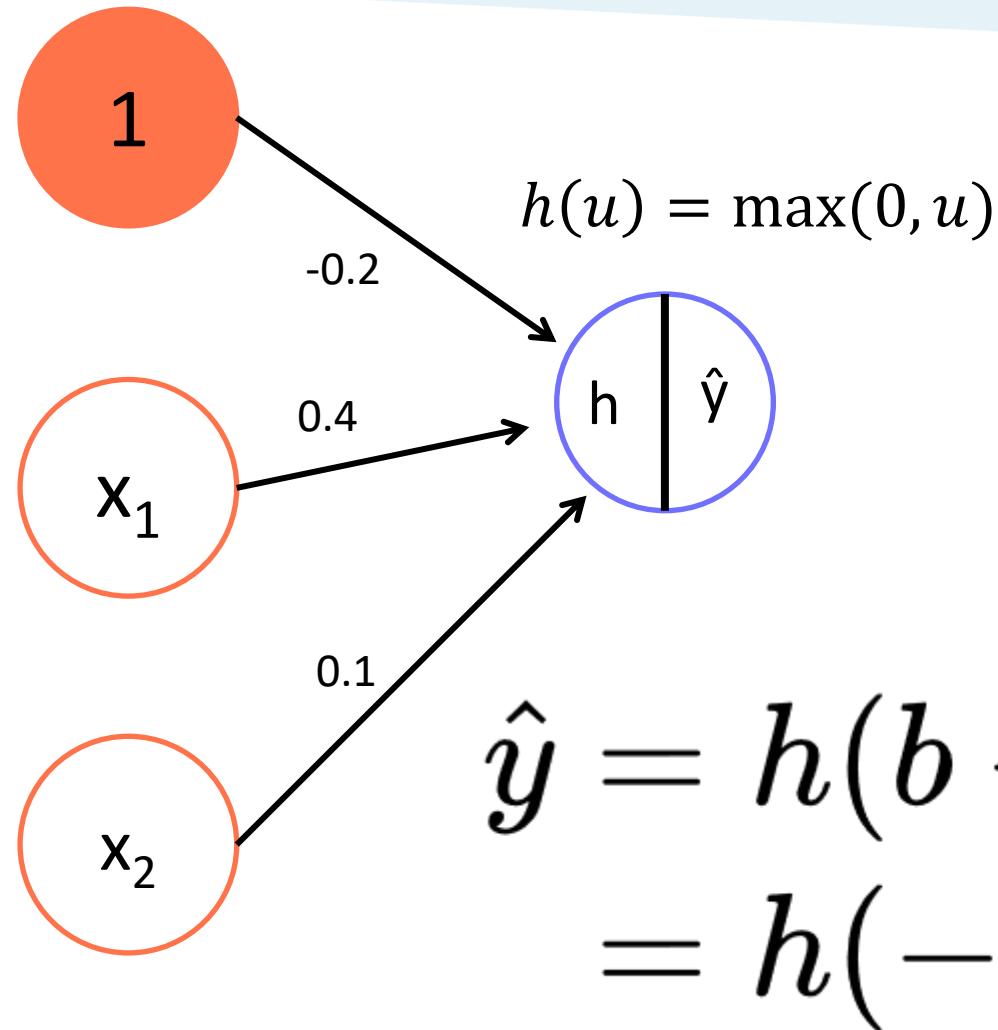
$$\begin{aligned}\frac{\partial E(w, b)}{\partial b} &= \frac{\partial E(w, b)}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial b} \\ &= \left(\frac{-y}{\hat{y}} + \frac{1 - y}{1 - \hat{y}} \right)\end{aligned}$$

Update gradient descend



$$w^+ = w - \alpha \frac{\partial E(w, b)}{\partial w}$$
$$b^+ = b - \alpha \frac{\partial E(w, b)}{\partial b}$$

AND operator trainen



AND-operator

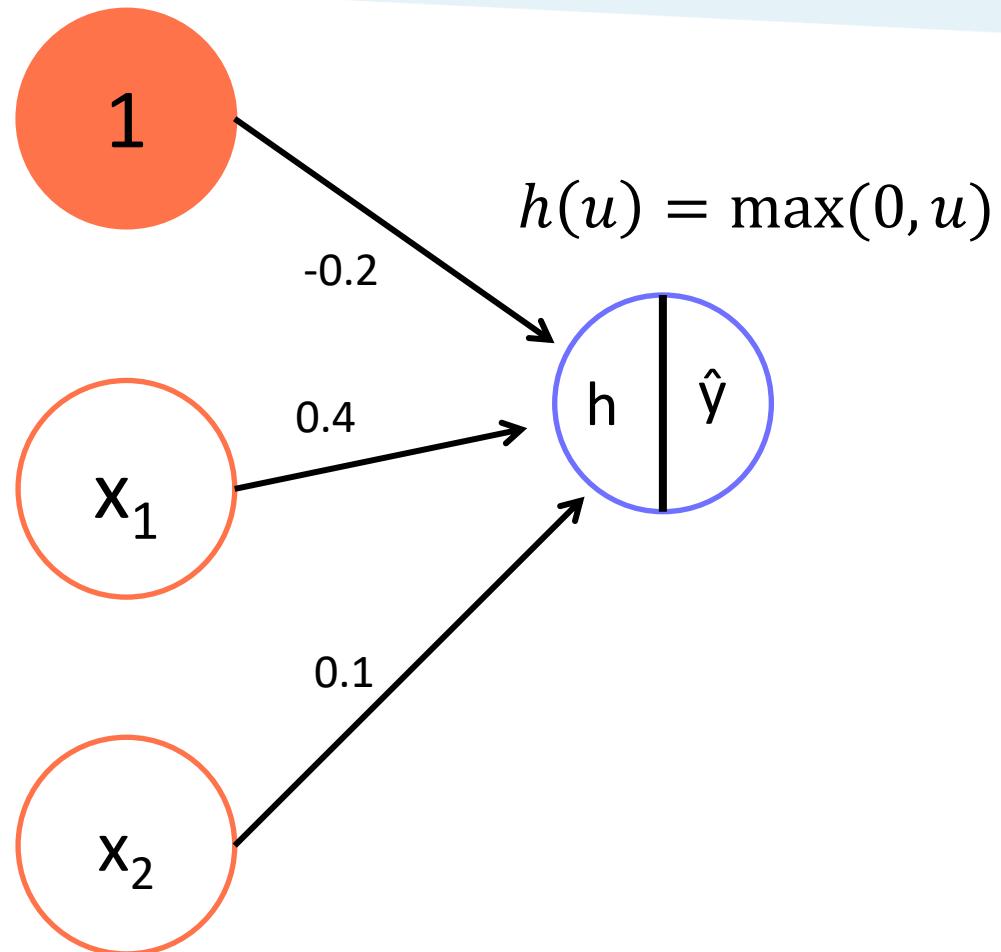
$$(x_1=1) \cap (x_2=1) = 1$$

$$(x_1=1) \cap (x_2=0) = 0$$

$$(x_1=0) \cap (x_2=1) = 0$$

$$(x_1=0) \cap (x_2=0) = 0$$

AND operator trainen



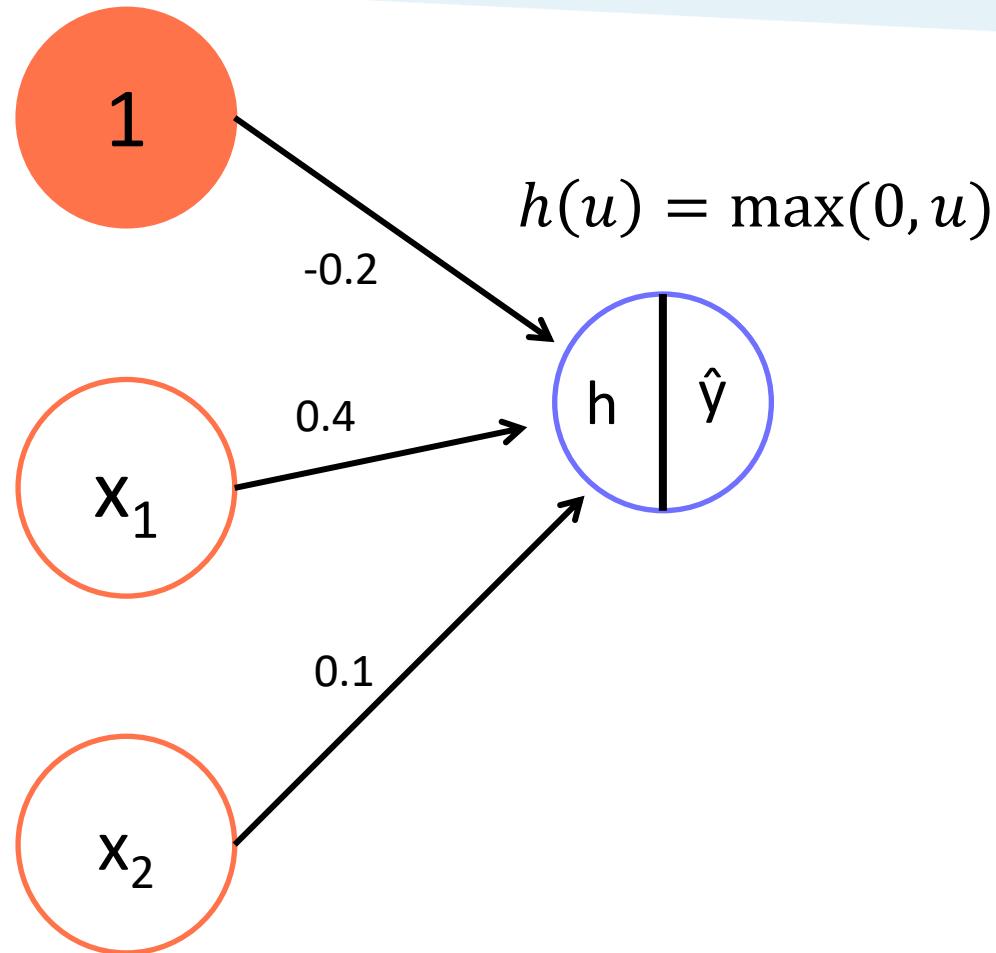
$$\begin{aligned}\hat{y} &= h(b + w_1x_1 + w_2x_2) \\ &= h(-0.2 + 0.4x_1 + 0.1x_2)\end{aligned}$$

$$\min_w E(w, b)$$

$$\frac{\partial E(w, b)}{\partial w} = \frac{\partial E(w, b)}{\partial \hat{y}} \boxed{\frac{\partial \hat{y}}{\partial h}} \frac{\partial h}{\partial w}$$

$$\frac{\partial E(w, b)}{\partial b} = \frac{\partial E(w, b)}{\partial \hat{y}} \boxed{\frac{\partial \hat{y}}{\partial h}} \frac{\partial h}{\partial b}$$

AND operator trainen



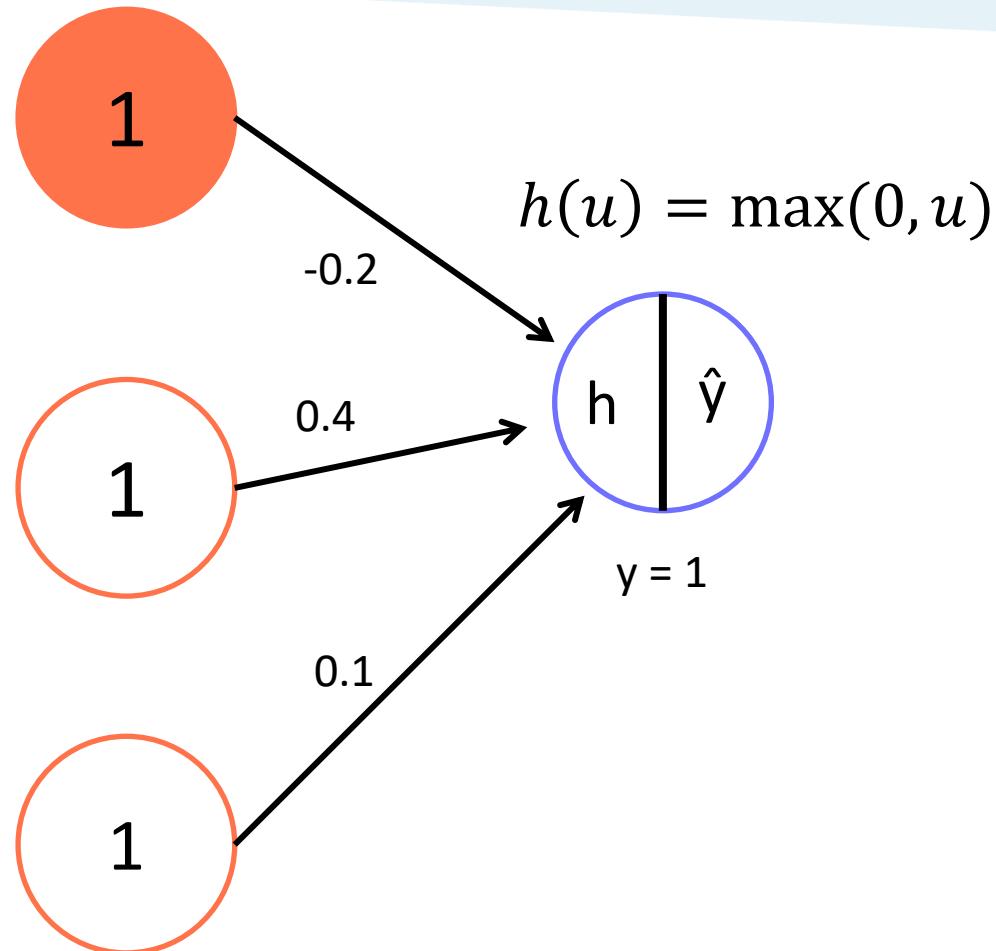
$$\begin{aligned}\hat{y} &= h(b + w_1x_1 + w_2x_2) \\ &= h(-0.2 + 0.4x_1 + 0.1x_2)\end{aligned}$$

$$\frac{\partial E(w, b)}{\partial w} = \left(\frac{-y}{\hat{y}} + \frac{1 - y}{1 - \hat{y}} \right) h'(x, w, b)x$$

$$\frac{\partial E(w, b)}{\partial b} = \left(\frac{-y}{\hat{y}} + \frac{1 - y}{1 - \hat{y}} \right) h'(x, w, b)$$

$$h'(x, w, b) = \begin{cases} 1 & \text{als } b + w_1x_1 + w_2x_2 > 0 \\ 0 & \text{anders} \end{cases}$$

AND operator trainen



$$\begin{aligned}\hat{y} &= h(b + w_1x_1 + w_2x_2) \\ &= h(-0.2 + 0.4x_1 + 0.1x_2)\end{aligned}$$

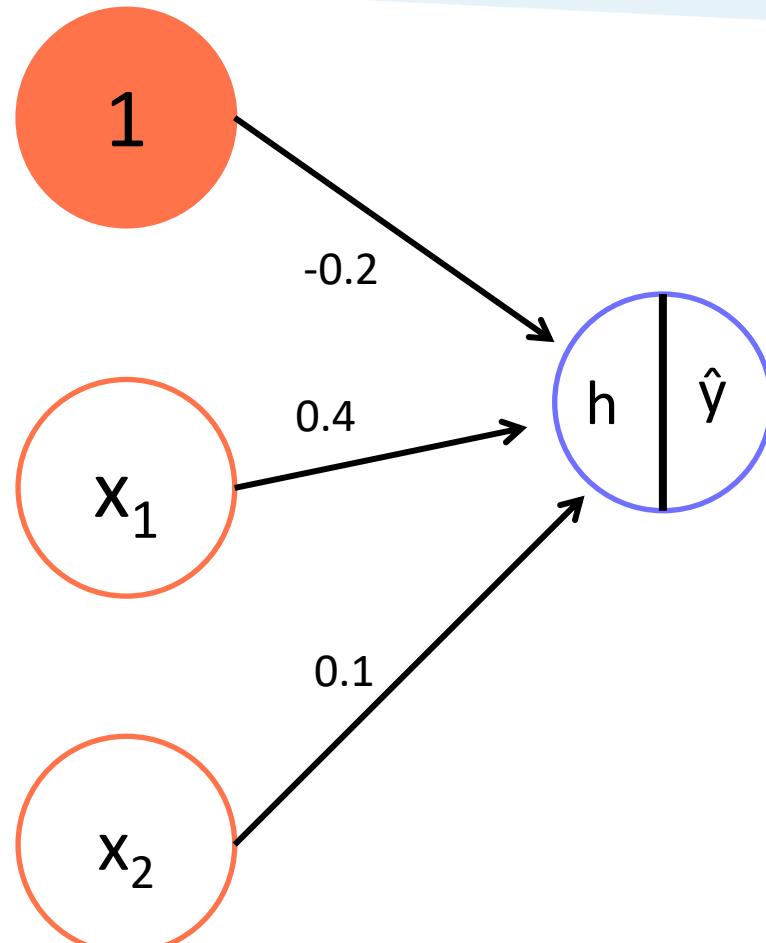
$$\frac{\partial E(w, b)}{\partial w} = \left(\frac{-y}{\hat{y}} + \frac{1-y}{1-\hat{y}} \right) h'(x, w, b)x$$

$$\frac{\partial E(w, b)}{\partial b} = \left(\frac{-y}{\hat{y}} + \frac{1-y}{1-\hat{y}} \right) h'(x, w, b)$$

$$h'(x, w, b) = \begin{cases} 1 & \text{als } b + w_1x_1 + w_2x_2 > 0 \\ 0 & \text{anders} \end{cases}$$

Opdracht: bereken de gradienten

AND operator trainen



$$h(u) = \max(0, u)$$

AND-operator

$$(x_1=1) \cap (x_2=1) = 1$$

$$(x_1=1) \cap (x_2=0) = 0$$

$$(x_1=0) \cap (x_2=1) = 0$$

$$(x_1=0) \cap (x_2=0) = 0$$

$$\hat{y} = h(-0.2 + 0.4 \cdot 1 + 0.1 \cdot 1) = 0.3$$

$$\hat{y} = h(-0.2 + 0.4 \cdot 1 + 0.1 \cdot 0) = 0.2$$

$$\hat{y} = h(-0.2 + 0.4 \cdot 0 + 0.1 \cdot 1) = 0$$

$$\hat{y} = h(-0.2 + 0.4 \cdot 0 + 0.1 \cdot 0) = 0$$

AND operator trainen

$$\hat{y} = h(-0.2 + 0.4 \cdot 1 + 0.1 \cdot 1) = 0.3$$

$$\hat{y} = h(-0.2 + 0.4 \cdot 1 + 0.1 \cdot 0) = 0.2$$

$$\hat{y} = h(-0.2 + 0.4 \cdot 0 + 0.1 \cdot 1) = 0$$

$$\hat{y} = h(-0.2 + 0.4 \cdot 0 + 0.1 \cdot 0) = 0$$

AND-operator

$$(x_1=1) \cap (x_2=1) = 1$$

$$(x_1=1) \cap (x_2=0) = 0$$

$$(x_1=0) \cap (x_2=1) = 0$$

$$(x_1=0) \cap (x_2=0) = 0$$

$$\frac{\partial E(w, b)}{\partial w_1} = \left(\frac{-y}{\hat{y}} + \frac{1 - y}{1 - \hat{y}} \right) h'(x, w, b) x_1 = \frac{-1}{0.3} \cdot 1 \cdot 1$$

$$\frac{\partial E(w, b)}{\partial w_2} = \left(\frac{-y}{\hat{y}} + \frac{1 - y}{1 - \hat{y}} \right) h'(x, w, b) x_2 = \frac{-1}{0.3} \cdot 1 \cdot 1$$

$$\frac{\partial E(w, b)}{\partial b} = \left(\frac{-y}{\hat{y}} + \frac{1 - y}{1 - \hat{y}} \right) h'(x, w, b) = \frac{-1}{0.3} \cdot 1$$

$$h'(x, w, b) = \begin{cases} 1 & \text{als } b + w_1 x_1 + w_2 x_2 > 0 \\ 0 & \text{anders} \end{cases}$$

AND operator trainen

$$\hat{y} = h(-0.2 + 0.4 \cdot 1 + 0.1 \cdot 1) = 0.3$$

$$\hat{y} = h(-0.2 + 0.4 \cdot 1 + 0.1 \cdot 0) = 0.2$$

$$\hat{y} = h(-0.2 + 0.4 \cdot 0 + 0.1 \cdot 1) = 0$$

$$\hat{y} = h(-0.2 + 0.4 \cdot 0 + 0.1 \cdot 0) = 0$$

AND-operator

$$(x_1=1) \cap (x_2=1) = 1$$

$$(x_1=1) \cap (x_2=0) = 0$$

$$(x_1=0) \cap (x_2=1) = 0$$

$$(x_1=0) \cap (x_2=0) = 0$$

$$\frac{\partial E(w, b)}{\partial w_1} = \frac{1}{0.8} \cdot 1 \cdot 1$$

$$\frac{\partial E(w, b)}{\partial w_2} = \frac{1}{0.8} \cdot 1 \cdot 0$$

$$\frac{\partial E(w, b)}{\partial b} = \frac{1}{0.8} \cdot 1$$

AND operator trainen

$$\hat{y} = h(-0.2 + 0.4 \cdot 1 + 0.1 \cdot 1) = 0.3$$

$$\hat{y} = h(-0.2 + 0.4 \cdot 1 + 0.1 \cdot 0) = 0.2$$

$$\hat{y} = h(-0.2 + 0.4 \cdot 0 + 0.1 \cdot 1) = 0$$

$$\hat{y} = h(-0.2 + 0.4 \cdot 0 + 0.1 \cdot 0) = 0$$

AND-operator

$$(x_1=1) \cap (x_2=1) = 1$$

$$(x_1=1) \cap (x_2=0) = 0$$

$$(x_1=0) \cap (x_2=1) = 0$$

$$(x_1=0) \cap (x_2=0) = 0$$

$$\frac{\partial E(w, b)}{\partial w_1} = \frac{1}{1} \cdot 0 \cdot 0$$

$$\frac{\partial E(w, b)}{\partial w_2} = \frac{1}{1} \cdot 0 \cdot 0$$

$$\frac{\partial E(w, b)}{\partial b} = \frac{1}{1} \cdot 0$$

$$\frac{\partial E(w, b)}{\partial w_1} = \left(\frac{-y}{\hat{y}} + \frac{1-y}{1-\hat{y}} \right) h'(x, w, b) x_1 = \frac{-1}{0.3} \cdot 1 \cdot 1$$

$$\frac{\partial E(w, b)}{\partial w_2} = \left(\frac{-y}{\hat{y}} + \frac{1-y}{1-\hat{y}} \right) h'(x, w, b) x_2 = \frac{-1}{0.3} \cdot 1 \cdot 1$$

$$\frac{\partial E(w, b)}{\partial b} = \left(\frac{-y}{\hat{y}} + \frac{1-y}{1-\hat{y}} \right) h'(x, w, b) = \frac{-1}{0.3} \cdot 1$$

$$\frac{\partial E(w, b)}{\partial w_1} = \frac{1}{0.8} \cdot 1 \cdot 1$$

$$\frac{\partial E(w, b)}{\partial w_2} = \frac{1}{0.8} \cdot 1 \cdot 0$$

$$\frac{\partial E(w, b)}{\partial b} = \frac{1}{0.8} \cdot 1$$

$$\frac{\partial E(w, b)}{\partial w_1} = \frac{1}{1} \cdot 0 \cdot 0$$

$$\frac{\partial E(w, b)}{\partial w_2} = \frac{1}{1} \cdot 0 \cdot 0$$

$$\frac{\partial E(w, b)}{\partial b} = \frac{1}{1} \cdot 0$$

$$\alpha = 0.1$$

$$w_1^+ = w_1 - \alpha \left(-\frac{1}{0.3} + \frac{1}{0.8} \right) = 0.6083$$

$$w_2^+ = w_2 - \alpha \left(-\frac{1}{0.3} \right) = 0.4333$$

$$b^+ = b - \alpha \left(-\frac{1}{0.3} + \frac{1}{0.8} \right) = 0.0083$$

In python

```
epoch = 0
w1 = 0.4
w2 = 0.1
b = -0.2

X1 = np.array([1, 1, 0, 0])
X2 = np.array([1, 0, 1, 0])
Y = X1 & X2

dEdw1 = 0
dEdw2 = 0
dEdb = 0
E = 0

for x1, x2, y in zip(X1, X2, Y):
    h = b + w1*x1 + w2*x2
    yhat = max(0, h)
    E += -np.log(yhat) if y else - np.log(1-yhat)
    dEdy = -1/yhat if y else 1/(1-yhat)
    dEdw1 += dEdy*(h > 0) * x1
    dEdw2 += dEdy*(h > 0) * x2
    dEdb += dEdy*(h > 0)

alpha = 0.1
w1new = w1 - alpha * dEdw1
w2new = w2 - alpha * dEdw2
bnew = b - alpha * dEdb
print(f'Iteration {epoch}: loss = {E:.4f}')
print(w1, w1new)
print(w2, w2new)
print(b, bnew)
```

```
Iteration 0: loss = 1.4271
0.4 0.6083333333333334
0.1 0.4333333333333335
-0.2 0.0083333333333304
```

```

epoch = 0
w1 = 0.4
w2 = 0.1
b = -0.2

X1 = np.array([1, 1, 0, 0])
X2 = np.array([1, 0, 1, 0])
Y = X1 & X2

dEdw1 = 0
dEdw2 = 0
dEdb = 0
E = 0

for x1, x2, y in zip(X1, X2, Y):
    h = b + w1*x1 + w2*x2
    yhat = max(0, h)
    E += -np.log(yhat) if y else - np.log(1-yhat)
    dEdy = -1/yhat if y else 1/(1-yhat)
    dEdw1 += dEdy*(h > 0) * x1
    dEdw2 += dEdy*(h > 0) * x2
    dEdb += dEdy*(h > 0)

alpha = 0.1
w1new = w1 - alpha * dEdw1
w2new = w2 - alpha * dEdw2
bnew = b - alpha * dEdb
print(f'Iteration {epoch}: loss = {E:.4f}')
print(w1, w1new)
print(w2, w2new)
print(b, bnew)

```

Iteration 0: loss = 1.4271
0.4 0.608333333333334
0.1 0.433333333333335
-0.2 0.0083333333333304

```

epoch += 1
w1 = w1new
w2 = w2new
b = bnew

dEdw1 = 0
dEdw2 = 0
dEdb = 0
E = 0

for x1, x2, y in zip(X1, X2, Y):
    h = b + w1*x1 + w2*x2
    yhat = max(0, h)
    E += -np.log(yhat) if y else - np.log(1-yhat)
    dEdy = -1/yhat if y else 1/(1-yhat)
    dEdw1 += dEdy*(h > 0) * x1
    dEdw2 += dEdy*(h > 0) * x2
    dEdb += dEdy*(h > 0)

w1new = w1 - alpha * dEdw1
w2new = w2 - alpha * dEdw2
bnew = b - alpha * dEdb
print(f'Iteration {epoch}: loss = {E:.4f}')
print(w1, w1new)
print(w2, w2new)
print(b, bnew)

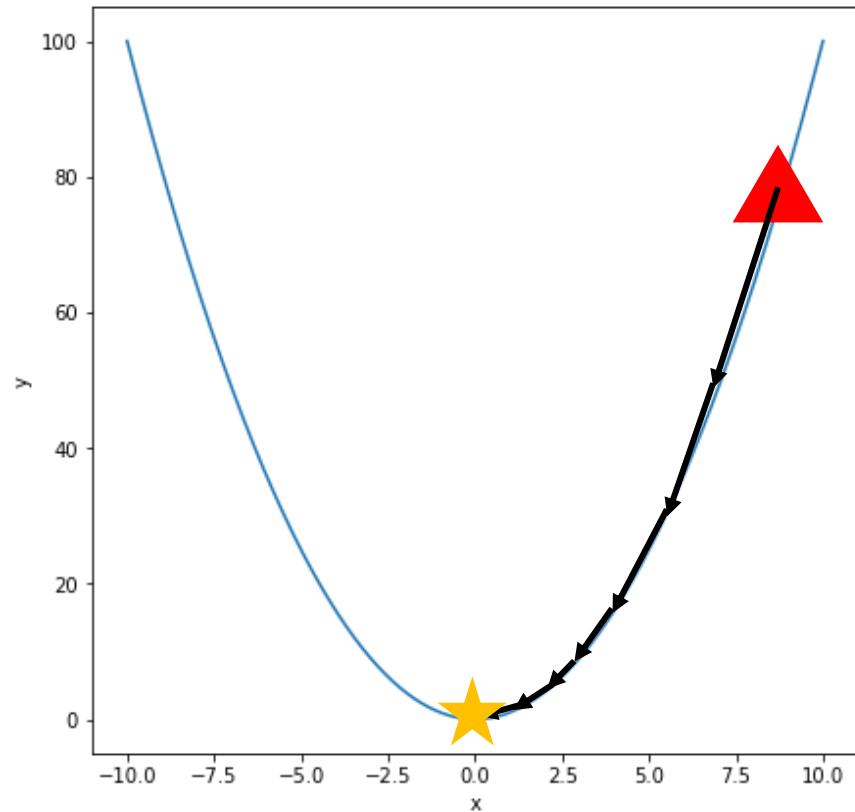
```

Iteration 1: loss = 1.5012
0.608333333333334 0.4427018633540373
0.4333333333333335 0.34946695095948827
0.0083333333333304 -0.43724295039235694

In python

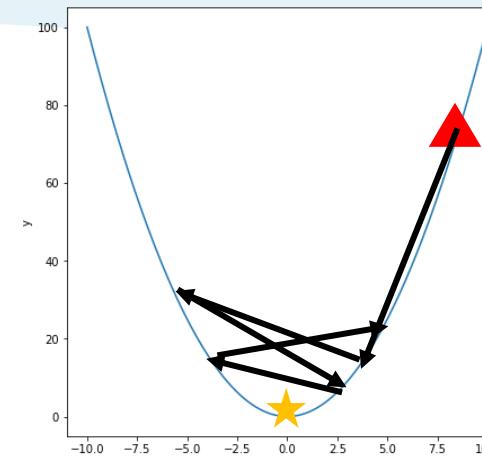
Learning rate

$LR = 0.001$

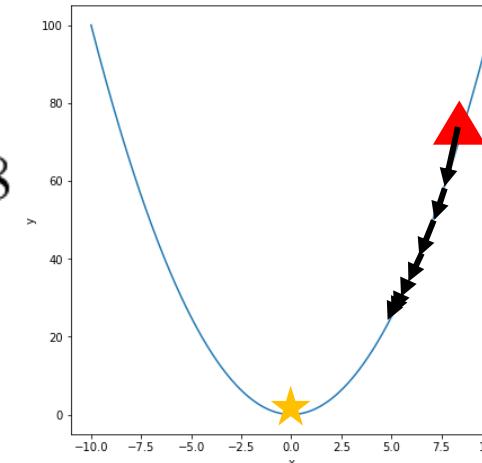


▲ Start punt

$LR = 1$



$LR = 10^{-8}$



★ Optimum

Voordelen

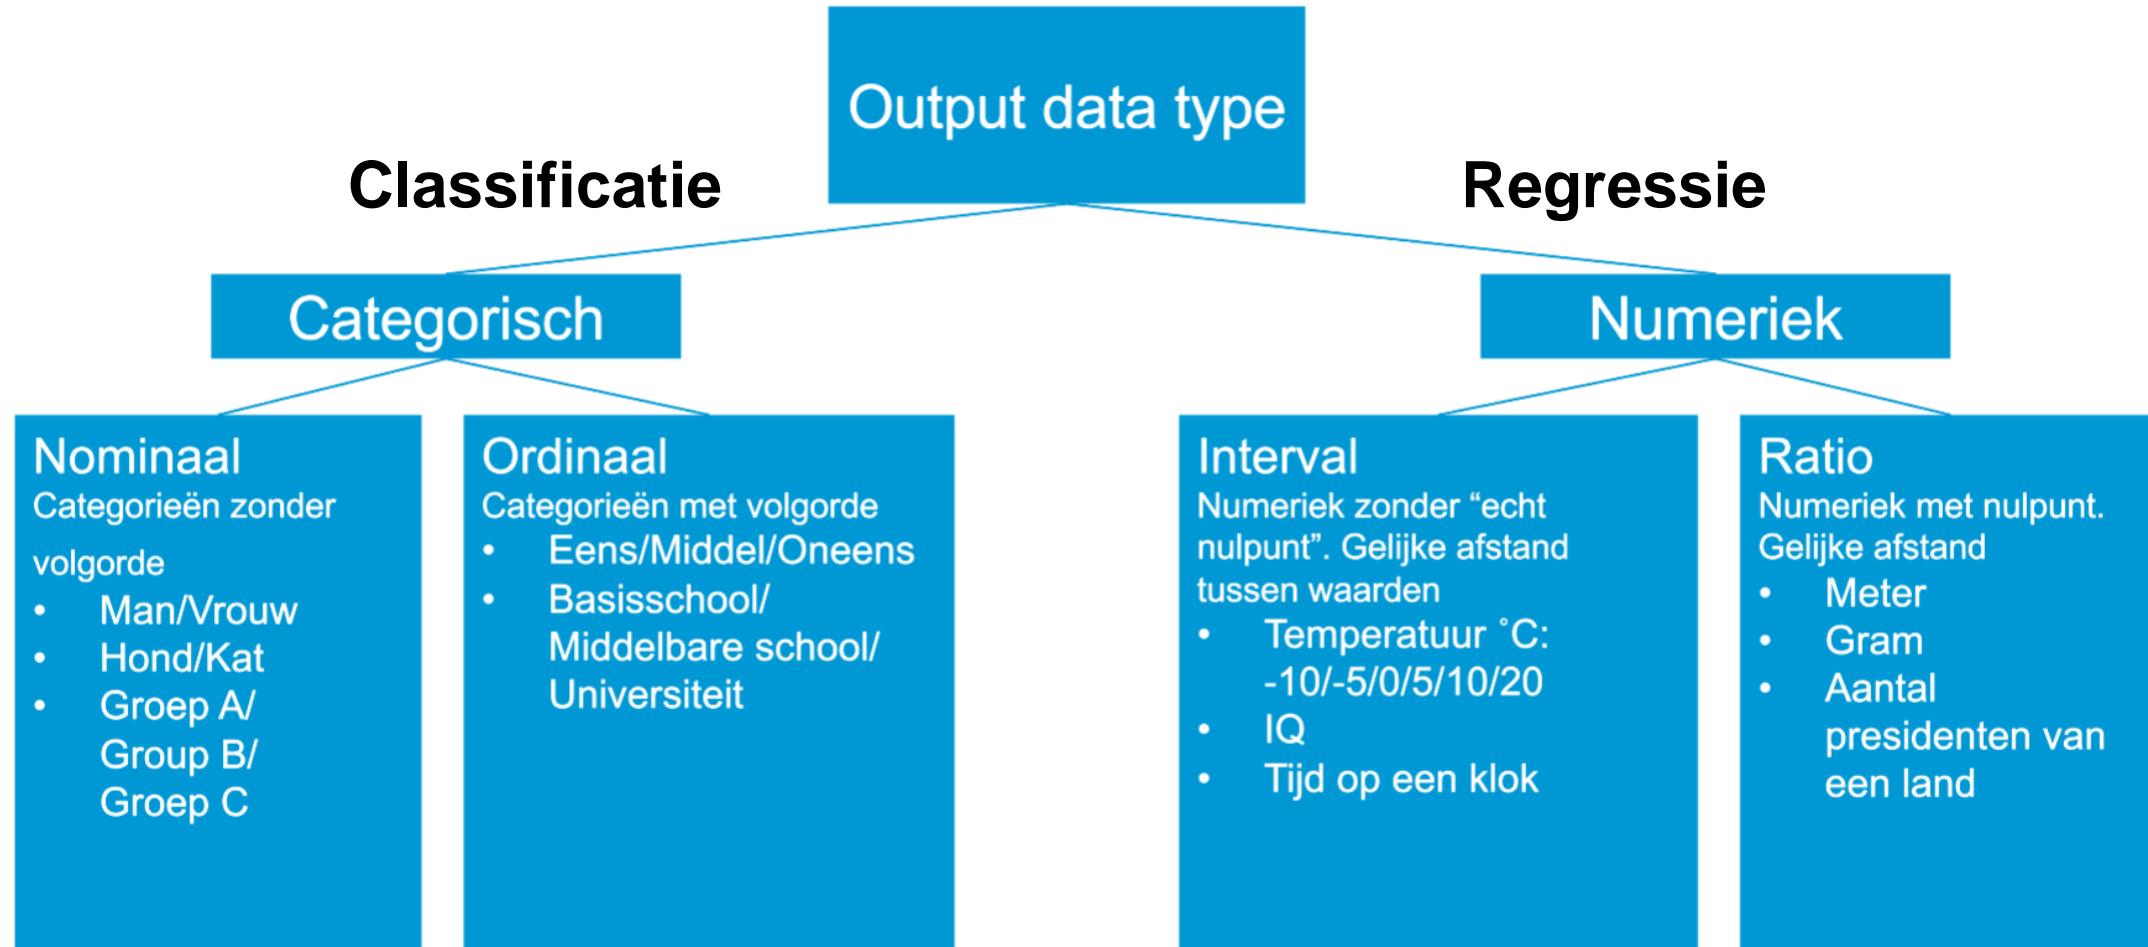
- Goed in het herkennen van complexe patronen
- Generiek voor nieuwe data
- Niet vatbaar voor de dataverdelingen

Nadelen

- Lastig om de hyperparameters te bepalen
- Trainen kan soms erg lang duren
- Gevoelig voor multicollineariteit (komt later in de cursus). Zorgt dat het netwerk minder snel tot een oplossing komt
- Mag door wetgeving niet gebruikt worden met persoonsgegevens

REGRESSIE

Verschil classificatie & regressie



Wat is het verschil tussen regressie en classificatie?

Classificatie

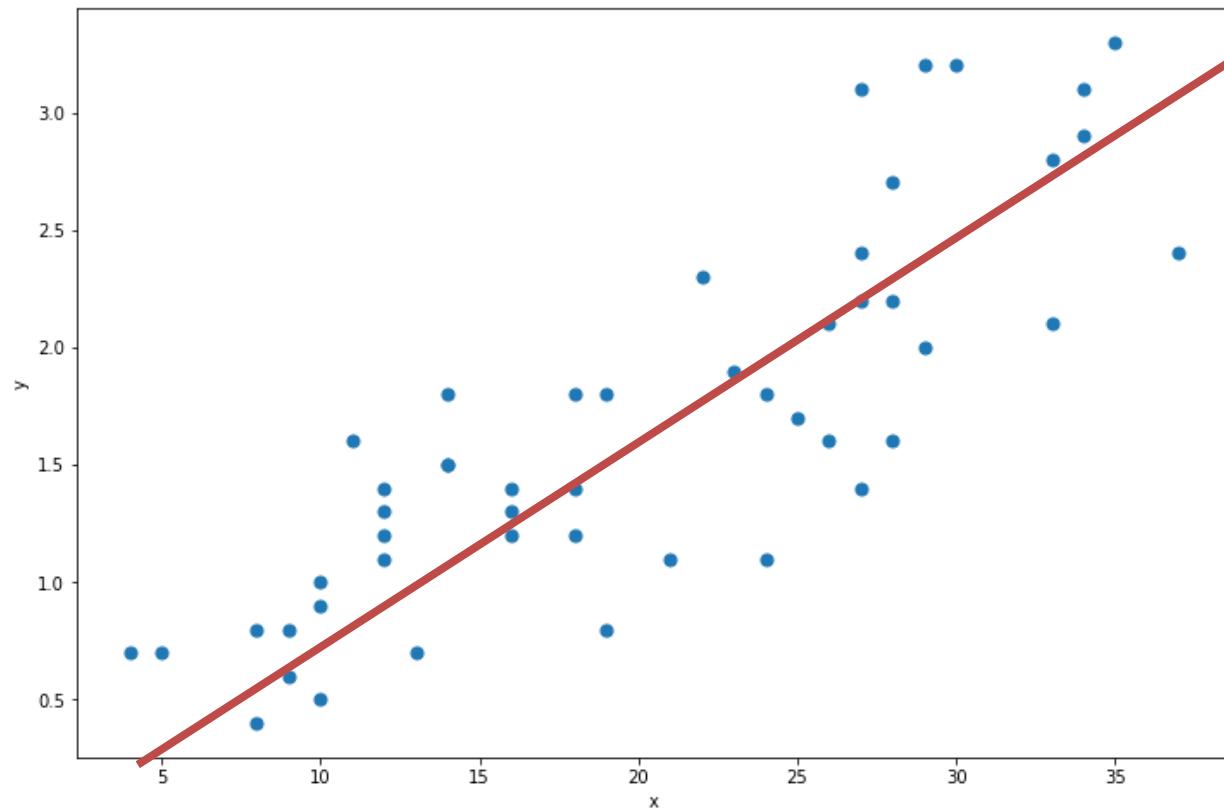
- Voorspellen van groepen/klassen
- Niet opzoek naar de relatie tussen de verschillende variabelen
- Welke variabelen zijn onderscheidend?
- Discreet

Regressie

- Voorspellen van waardes
- Opzoek naar een relatie tussen de afhankelijk variabele en de voorsteller.
- Welke variabelen zijn verklarend?
- Continue

Voorbeelden: Classificatie of regressie?

- Voor de kwaliteit van een product wordt een score aan gekoppeld van 0 tot 100. Voorspel de kwaliteit.
Regressie of classificatie?
- Voor de kwaliteitscontrole wordt een product gescand met als uitkomst: Fout, Goed, laten controleren door een expert.
Regressie of classificatie?
- Er wordt voorspeld of een robot snel vervangen moet worden.
Regressie of classificatie?



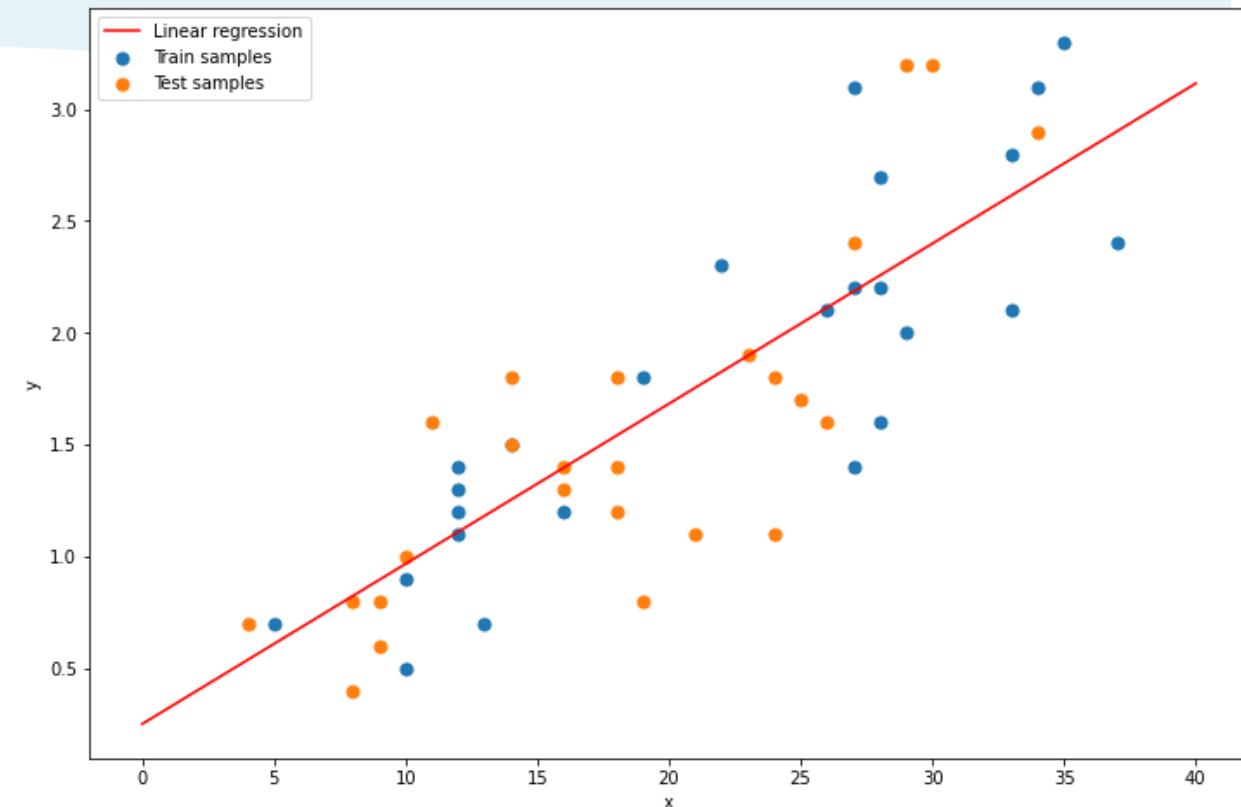
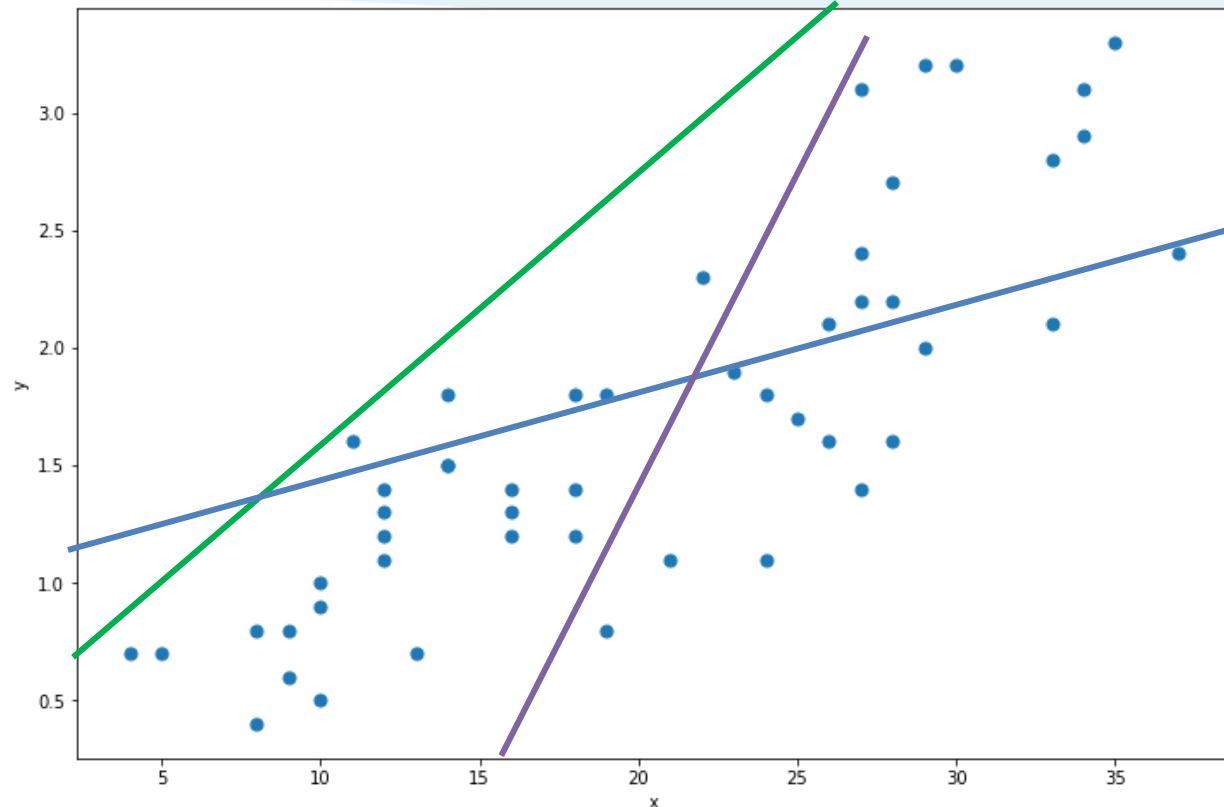
$$\hat{y} = \beta_1 x + \beta_0$$

Lineaire Regressie

$$\hat{y} = \beta_2 x_2 + \beta_1 x_1 + \beta_0$$

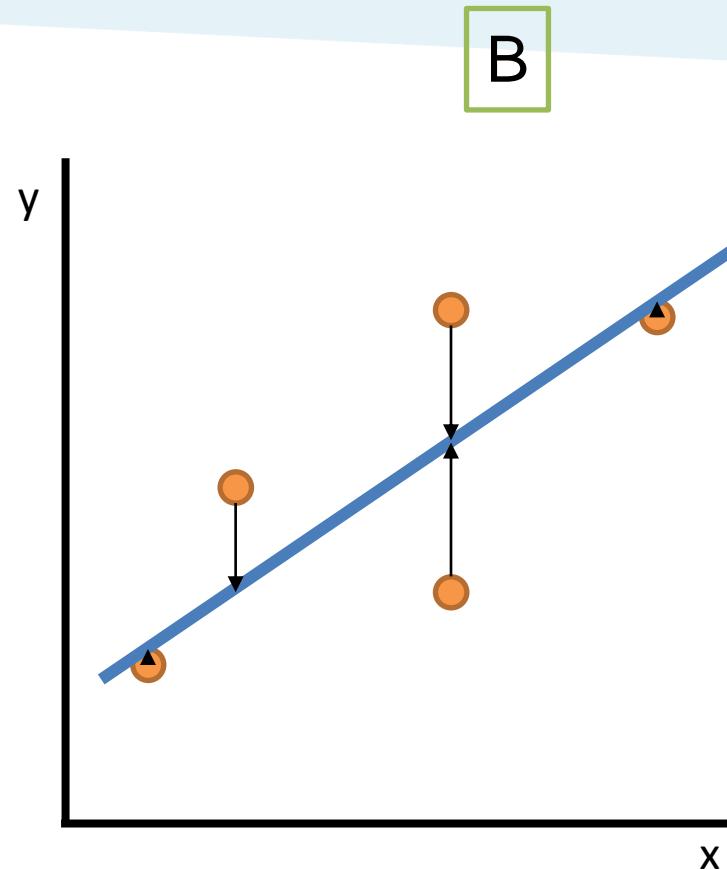
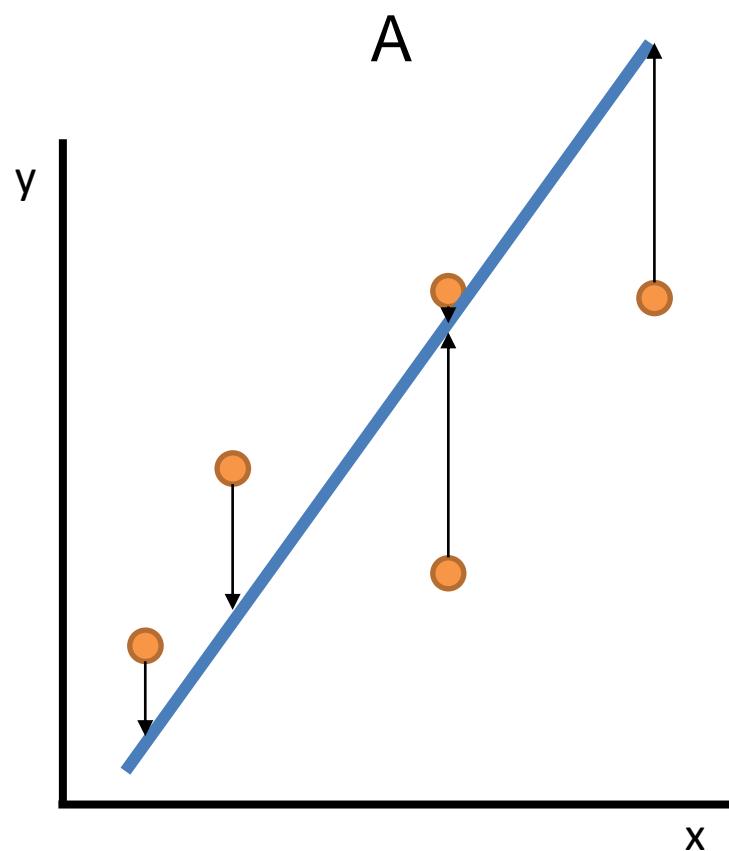
Multi Lineaire
Regressie

Lineaire Regressie



$$\hat{y} = \beta_1 x + \beta_0$$
$$\min_{\beta} (y - \hat{y})^2$$

Least squares method



Welke is beter?

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

Wat is de gemiddelde kwadratische afstand?
Range tussen 0 en oneindig

$$R^2 = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2}$$

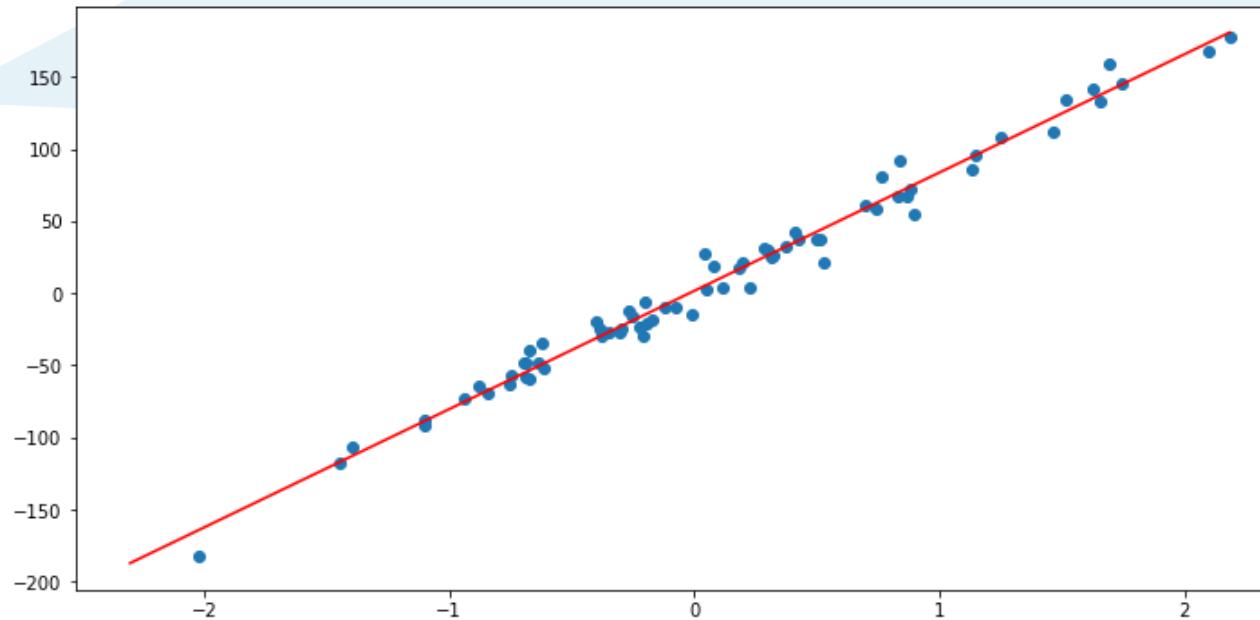
Hoeveel van de spreiding kan
verklaard worden door het model?

Range tussen 0 en 1

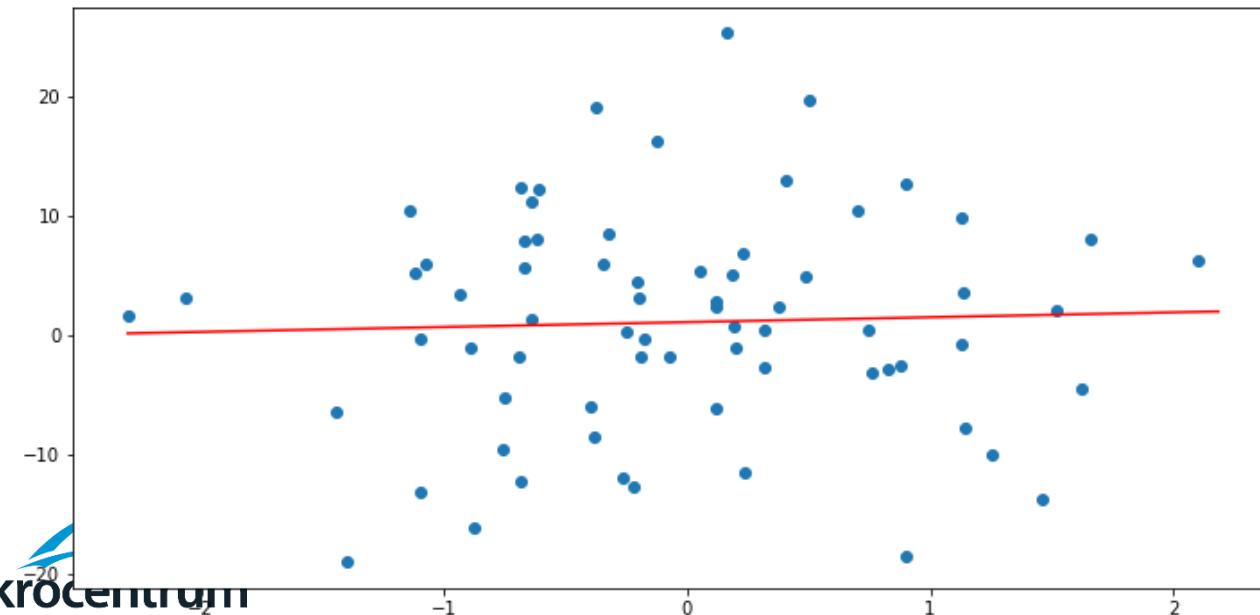
$$\sum_{i=1}^5 i = 1 + 2 + 3 + 4 + 5$$

\hat{y} = voorspelling van y
 \bar{y} = gemiddelde van y

Goede en slechte regressie



$MSE = 94.4$
 $R^2 = 0.98$



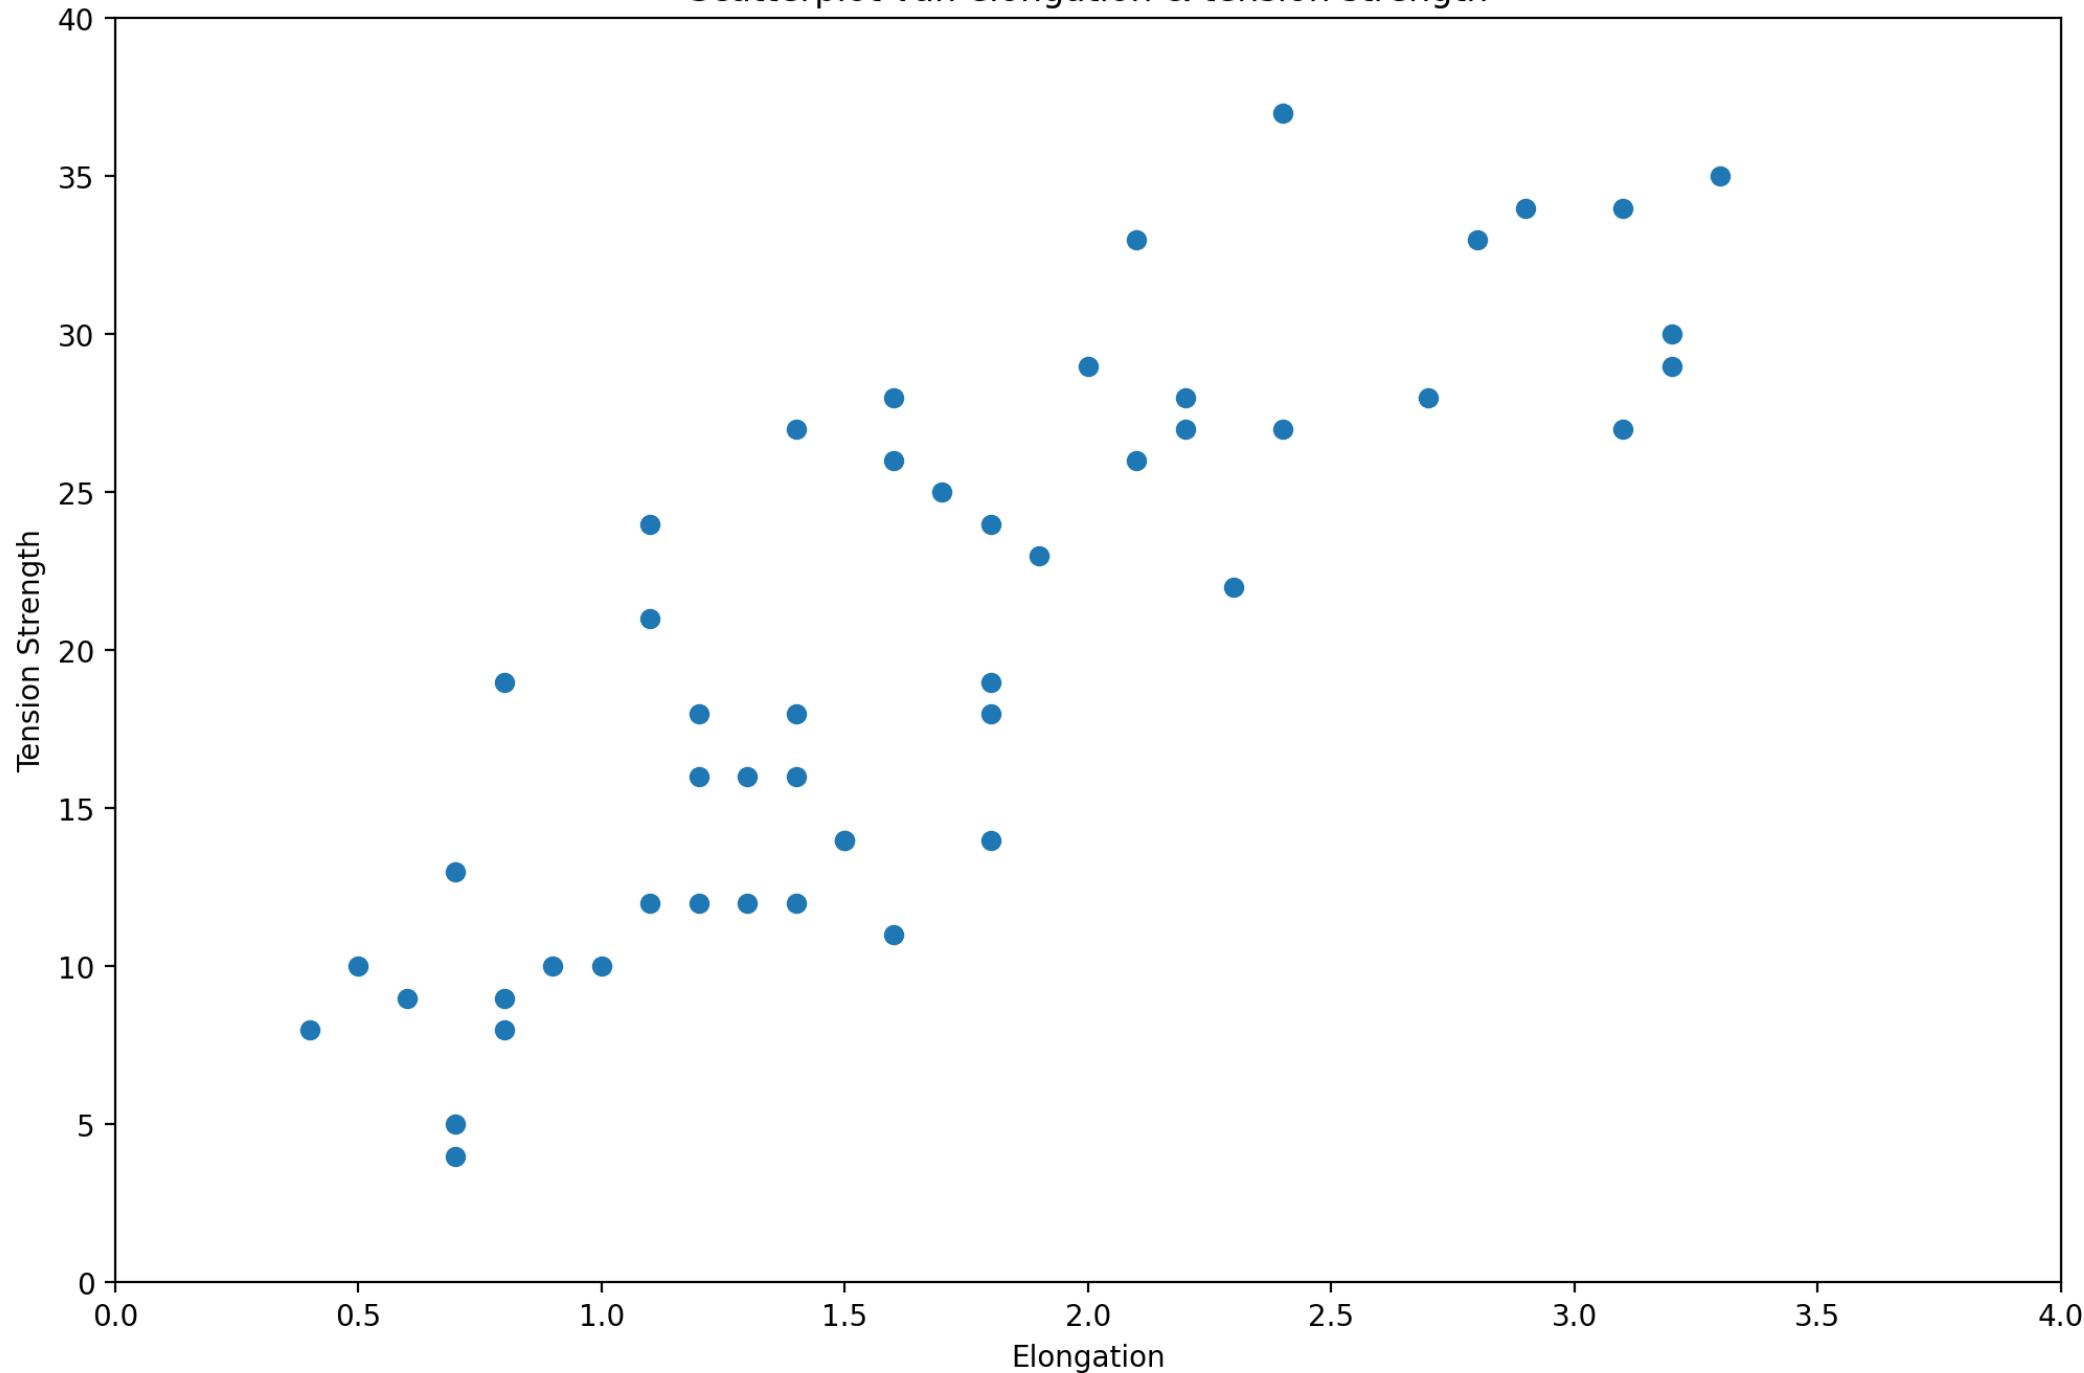
$MSE = 97.5$
 $R^2 = -0.009$

R² kan negatief zijn als de voorspellingen slechter zijn als het gemiddelde

- Met de 3d printer dataset, voorspel de waarde van “tension_strength” met de variabel "elongation“ (of gebruik de scatterplot op de volgende dia)
- Doe dit door waardes van β te schatten zonder sklearn en een zo laag mogelijke MSE te krijgen
- Maak een scatterplot van de voorspellingen en de daadwerkelijke waarden
- Maak een visualisatie

10 minuten de tijd

Scatterplot van elongation & tension strength



Correlatie is geen causatie

$$\hat{y} = \beta_1 x + \beta_0$$

β_1 = Hoeveel x
correleert met y

CORRELATIE IS GEEN CAUSATIE

Correlatie is geen causatie

Ik zie veel paraplu's op straat wanneer het regent. Dus regen wordt veroorzaakt door paraplu's.

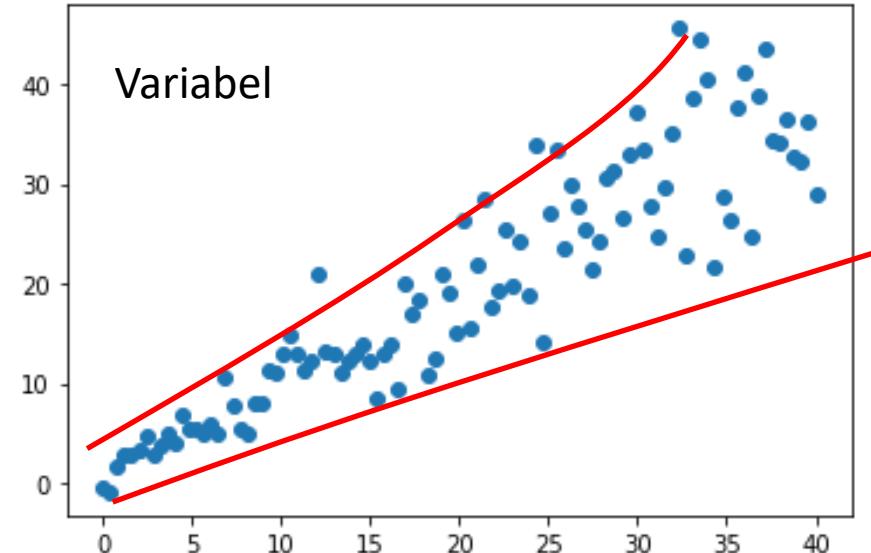
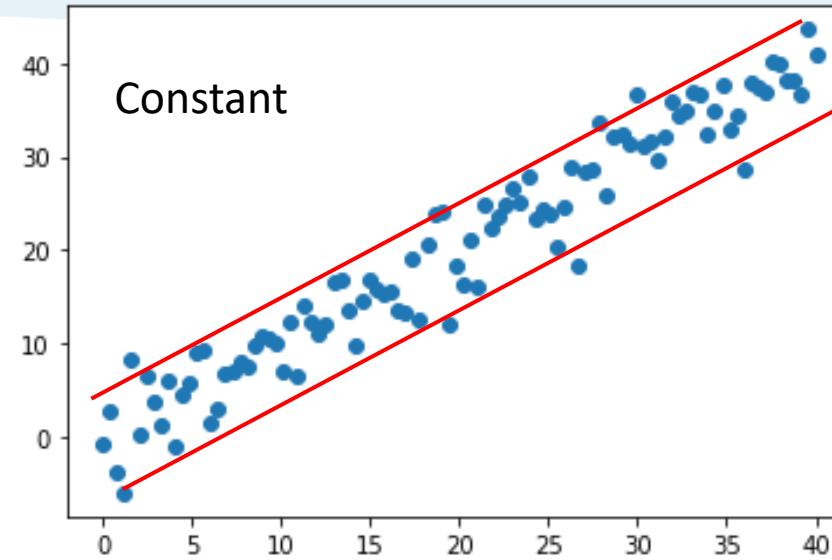
(tegenovergestelde causaliteit)

Slapen met je sokken aan zorgt voor wakker worden met hoofdpijn.
(causaliteit door een verbogen variabele)

De hoeveelheid verkochte staal voorspeld het aantal zetels van de Duitse socialistische partij (SPD)
(puur toeval)

Aannamens

- De relatie tussen x en y is lineair (duh)
- De variabel x bevat geen fouten (zoals meetfouten)
- Alle observaties zijn onafhankelijk van elkaar
- Constante variantie



Multi variabele lineaire regressie

$$\hat{y} = \beta_2 x_2 + \beta_1 x_1 + \beta_0$$

$$\min_{\beta} (y - \hat{y})^2$$

`sklearn.linear_model.LinearRegression`

Paar opmerkingen

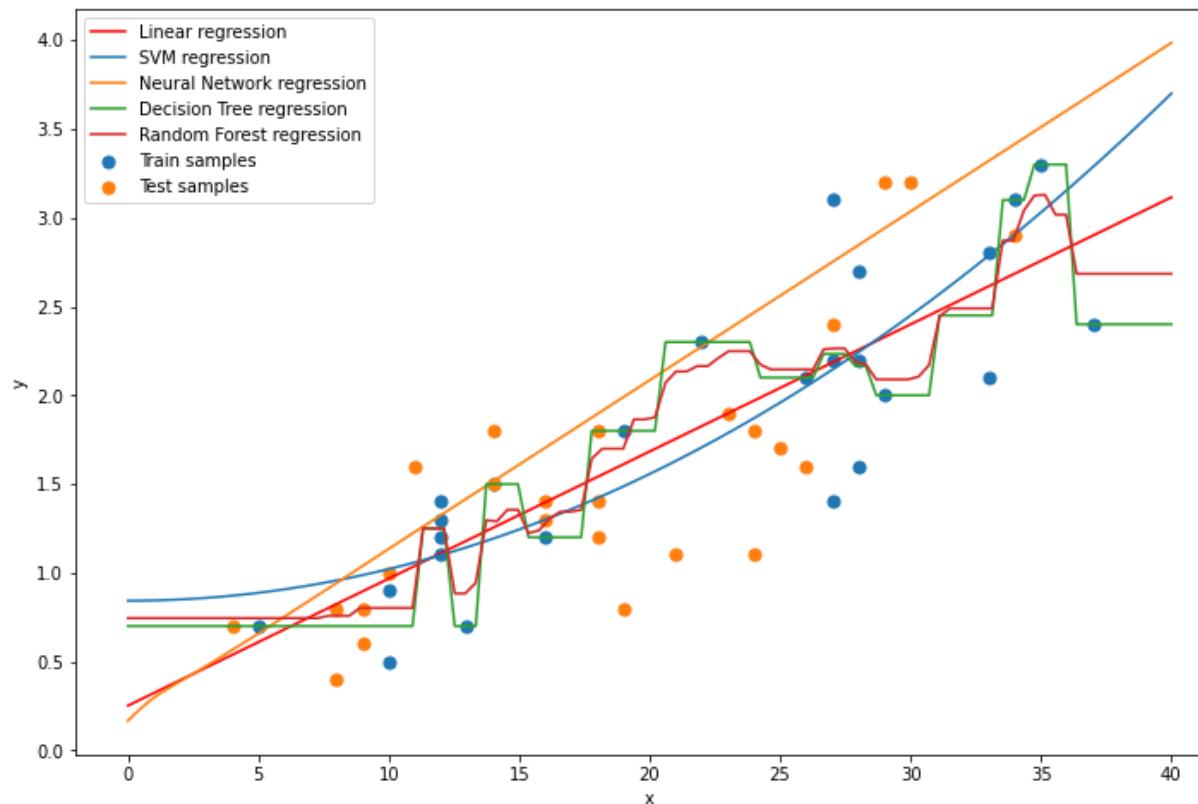
- Simpel = vaak beter
- x_1 en x_2 moeten onafhankelijk van elkaar zijn (collineariteit, komen we nog op terug)
- Meer variabelen vereist meer data

- Met de 3d printer dataset, voorspel de waarde van “tension_strength”, nu ook met andere variabele
- Gebruik de sklearn library
- Maak een scatterplot van de voorspellingen en de daadwerkelijke waarden.
`plt.scatter(y, y_pred)`
- Bereken de MSE en de R2

15 minuten de tijd

Andere regressie methodes

- Support Vector Machine Regress
- Decision Tree/Random Forest regressie
- Neural Network regressie



- Ridge regression (`sklearn.linear_model.Ridge`)

$$\min_{\beta} (y - \hat{y})^2 + \lambda \sum_{i=1}^p \beta_i^2$$

- Lasso regression (`sklearn.linear_model.Lasso`)

$$\min_{\beta} (y - \hat{y})^2 + \lambda \sum_{i=1}^p |\beta_i|$$

Regularization:

Het toevoegen van informatie aan het model om overfitting te voorkomen

- Met de 3d printer dataset, voorspel de waarde van “tension_strength”, nu met Ridge regression
- Bereken de MSE en de R²
- Visualiseer de β 's in een barplot van lineare regressie en ridge regressie:
`plt.barh(y = x_train.columns, width = linreg.coef_)`

10 minuten

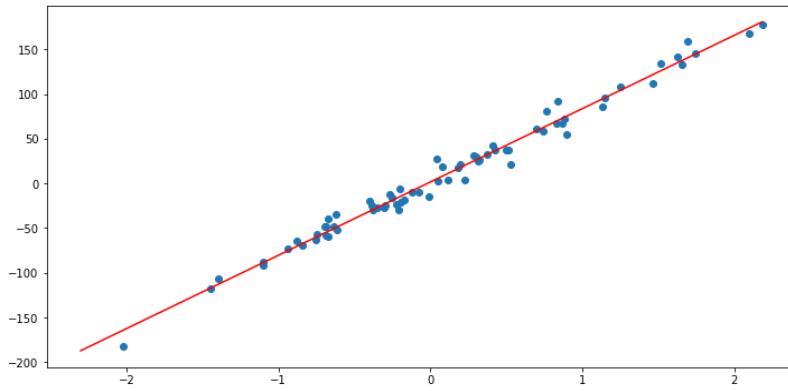
Lineair

- (Multi Variable) Linear regression
- Ridge Regression
- Lasso Regression



Non-lineair

- SVM Regression (SVR)
- Random Forest regression
- Neural Networks regression



METRICS

Performance measures

Accuracy

Hoeveel procent is correct geclassificeerd?

$$\text{Accuracy} = \frac{\text{aantal correct}}{\text{aantal voorspeld}}$$

Ground truth

		Predicted	
		A	B
A	A	True A	False B
	B	False A	True B

$$\text{Accuracy} = \frac{\text{True A} + \text{True B}}{\text{True A} + \text{False A} + \text{True B} + \text{False B}}$$

		Predicted	
		PLA	ABS
Ground truth	PLA	Correct PLA	Voorspeld ABS maar was PLA
	ABS	Voorspeld PLA maar was ABS	Correct ABS

Train-validatie-test split

Hele dataset

1
2
3
4
5
6
7
8
9
10

Train val test split

Verhouding 60/20/20

Trainset

3
5
2
8
6
10

Val

7
9

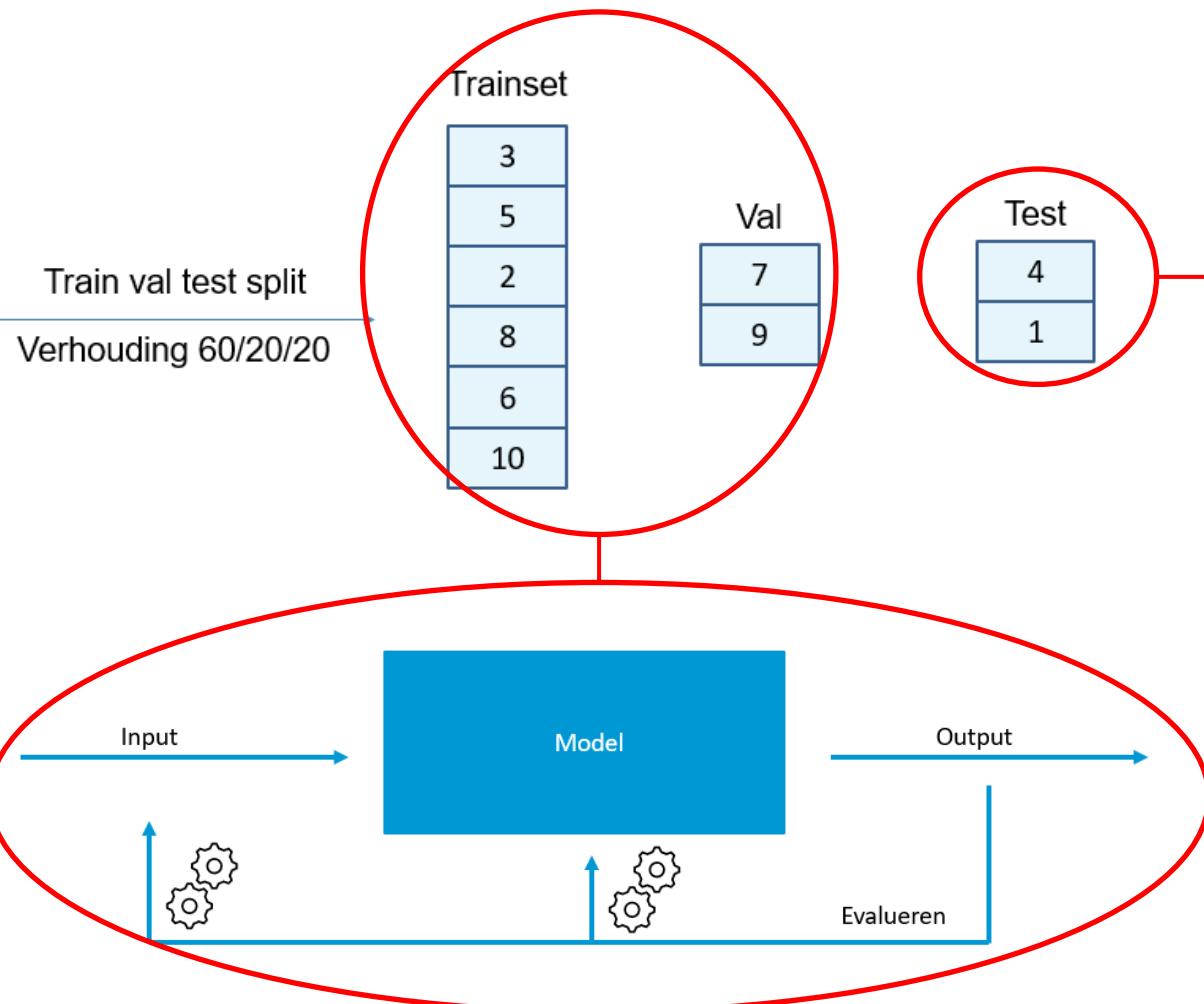
Test

4
1

Voorbeeld

Hele dataset

1
2
3
4
5
6
7
8
9
10



Betrouwbaarheid
99%

Oefening

Vergelijk alle modellen die je hebt gemaakt. Welk model is het beste om materiaal te voorspellen van de 3D-Printer dataset?

Welk model voorspelt het minste PLA terwijl dat eigenlijk ABS zou moeten zijn?

~15 minuten

Extra lesmateriaal

- Oefen met Pandas. Volg deze cursus (voor volgende week):
<https://www.kaggle.com/learn/pandas>
- Bias/variance tradeoff:
<https://www.youtube.com/watch?v=SjQyLhQIXSM>
- SVM kernel trick:
https://www.youtube.com/watch?v=3liCbRZPrZA&ab_channel=udiprod
- Decision Trees:
<https://scikit-learn.org/stable/modules/tree.html#classification>
- Random Forest:
https://www.stat.berkeley.edu/~breiman/RandomForests/cc_home.htm
- Hoe leert een machine (wiskundig):
https://www.youtube.com/watch?v=lHzwWFHWaw&t=5s&ab_channel=3Blue1Brown
- [Calling Bullshit](https://www.callingbullshit.org/tools/tools_misleading_axes.html)
(https://www.callingbullshit.org/tools/tools_misleading_axes.html): Hoe te misleiden met grafieken en data

Dagcasus (huiswerk)

Maak de opdrachten in de ‘dagcasus’ map.

Dag samenvatting

$$\mu = \frac{1}{n} \sum_{i=1}^n x_i$$

Hele dataset

